

# A (somehow) reasoned journey to GitOps

Consistency and Reproducibility  
in a Distributed System and Company



# TABLE OF CONTENTS



## Intro

The Operator as the basic Kubernetes pattern

## Helm

Kubernetes package management for reusability, consistency and reproducibility

## GitOps

Enforcing workflow in operations & standardizing on system state evolution

## Flux & beyond

Automating and standardizing CD pipelines for Kubernetes with the right mix of centralization & self-service



We are good people  
and create happiness

Positive mindset   Humility   Open-mindedness



We are  
the change

Courage   Agility   How yes



We win  
as a team

Diversity   Commitment   Collaborate



We are passionate  
and result-oriented

Ownership   Pragmatism   Hard work



We are transparent  
and respectful

Exemplarity   Trust   Integrity



We are customer  
and brand-centric

Impact   Loyalty   Engagement



We are passionate  
and result-oriented

Ownership

Pragmatism

Hard work

5

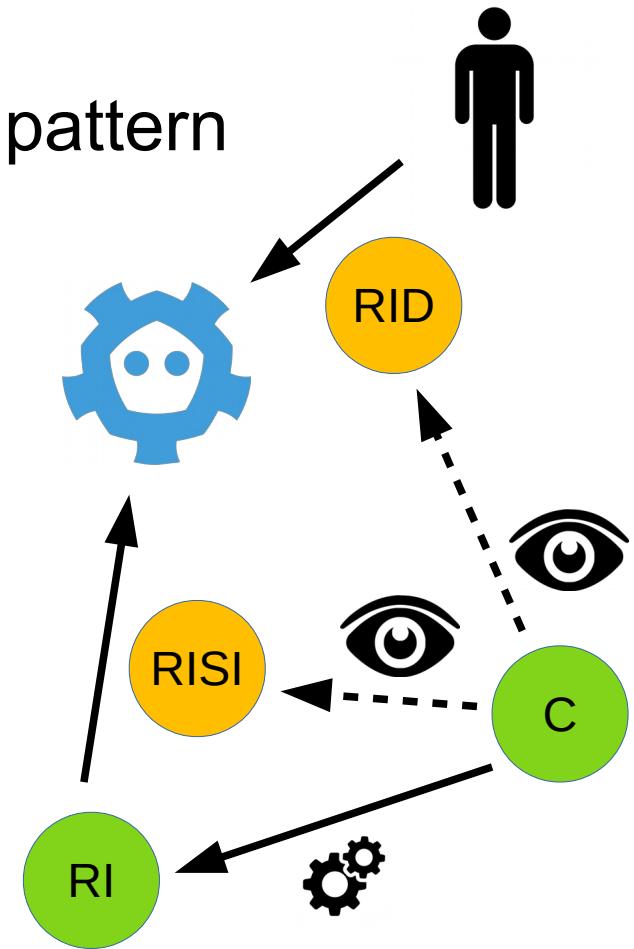


# Intro

The Operator as the basic Kubernetes pattern

# Operator: The basic Kubernetes pattern

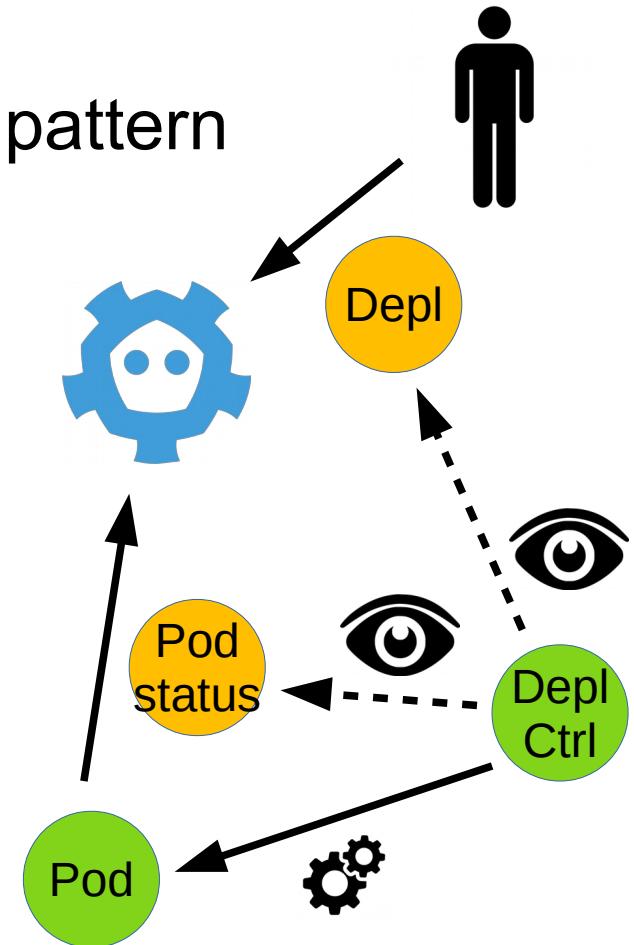
- Kube uses Etcd as a datastore to hold:
  - Resource Instance Definitions ( = desired state )
  - Resource Instance Status Info ( = actual state )
- Controllers
  - Watch Resource Instance Definitions (created, updated, deleted)
  - Create the Resource Instances as defined
  - Watch their current related Resource Instance Status Info
  - Strive to continuously match desired and actual state
    - In doing this, Controllers automate the operational know-how of handling each kind of Resource



# Operator: The basic Kubernetes pattern

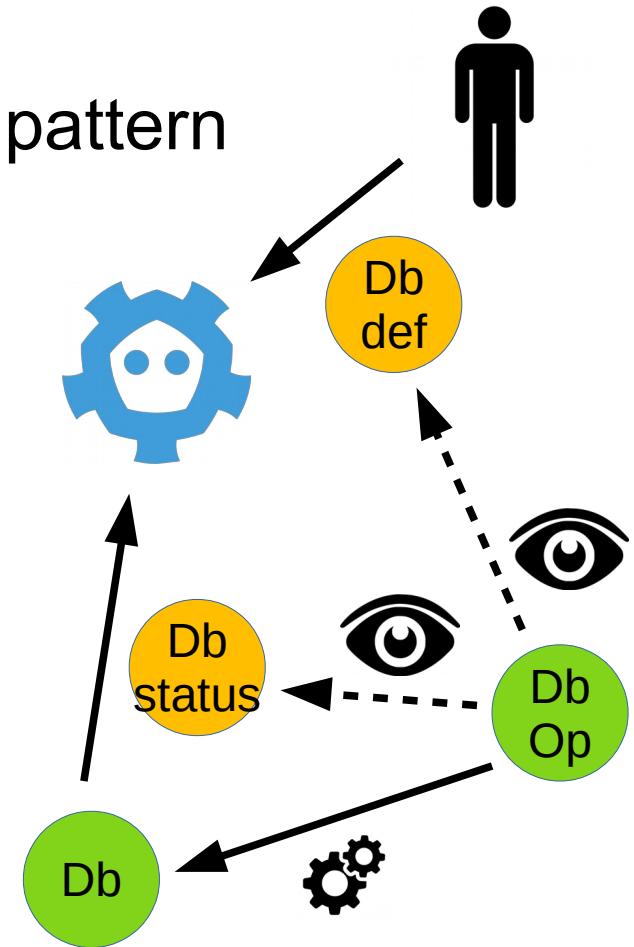
## Example: Deployment Controller

- Watches for Deployment definitions
- When a new one is detected, it creates a series of Pods (as required by def)
- Continuously watches Pod status to ensure all required Pods are running and healthy
- When Pods are killed unexpectedly, it creates them again
- When Deployments are changed to update the desired replica count, adds or removes Pods as needed



# Operator: The basic Kubernetes pattern

- **Controller** is a term used for code living in the K8s core
- But the same pattern can be applied for specialized applications that run in K8s like any other, but whose purpose is to manage and/or instantiate other applications: the **Operators**
- The pattern is similar: Users create **Custom Resource Definitions (CRD)** defined by the Operator itself, which expresses what they wish, and the Operator strives to make dreams (aka desired state) come true.
- The Operator handles the lifecycle of the **managed resources** by interacting with the **Kube API**, therefore **encapsulating the operational knowledge** of handling them, specially important when dealing with **stateful applications**.

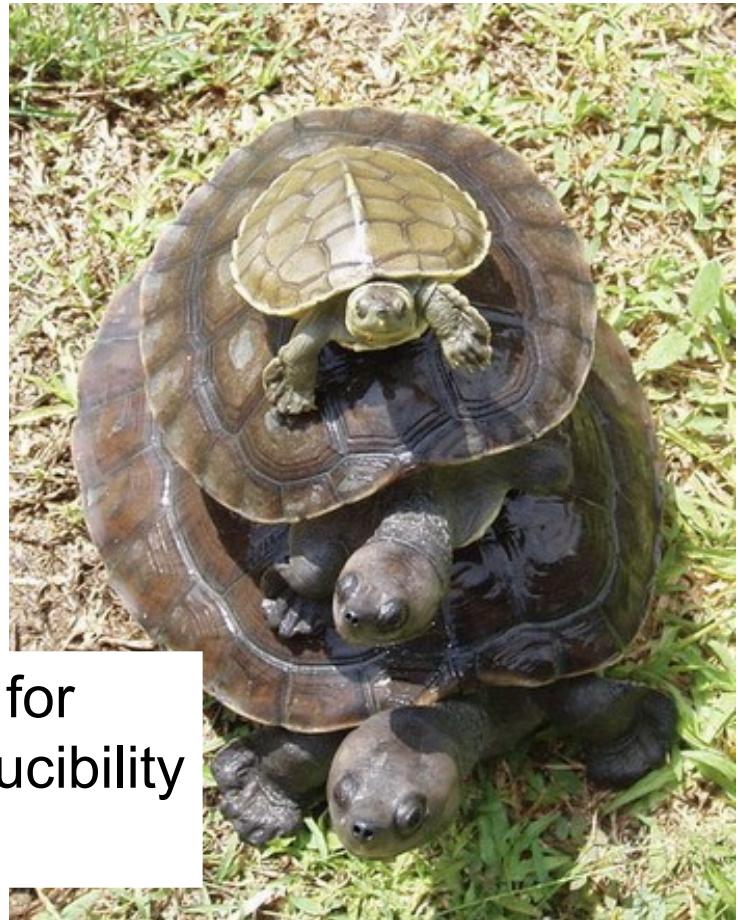


# DEMO #1

Basic Kubernetes & the Controller effect

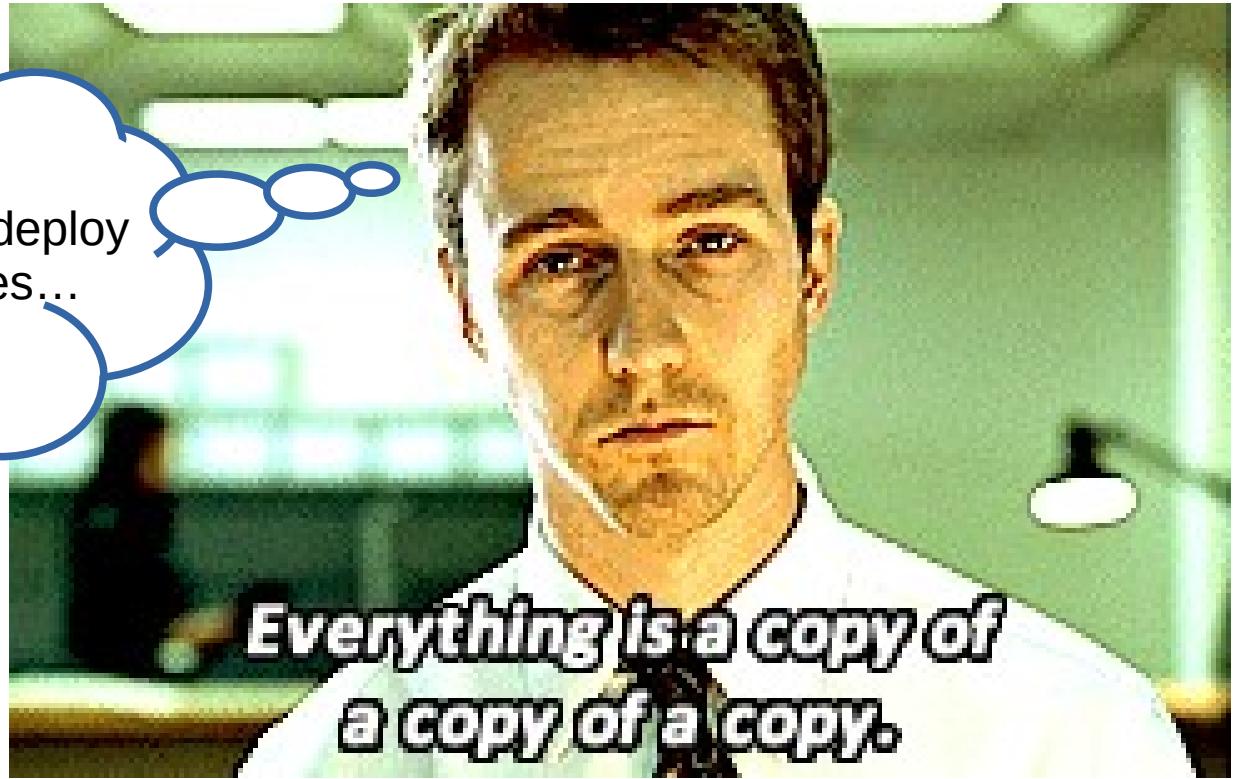


Kubernetes package management for  
reusability, consistency and reproducibility



# Helm

MyBoss asked me to deploy  
MySQL on Kubernetes...  
**Again.**



# Helm :: Purpose

**==> Reusable Kubernetes deployments**

- **Externalized configuration + Manifest templating**
  - Reusability of a deployment in many & different environments
- **Composability**
  - Creating a supersystem by reusing smaller systems
- **Package management**
  - Reusing public or privately published & versioned modules
- **Upgrade & rollback support**
  - Keeping track of published upgrades of a deployment



# Helm :: Concepts

- **Chart == Packaged (.tgz) set of Kube manifests, along with:**

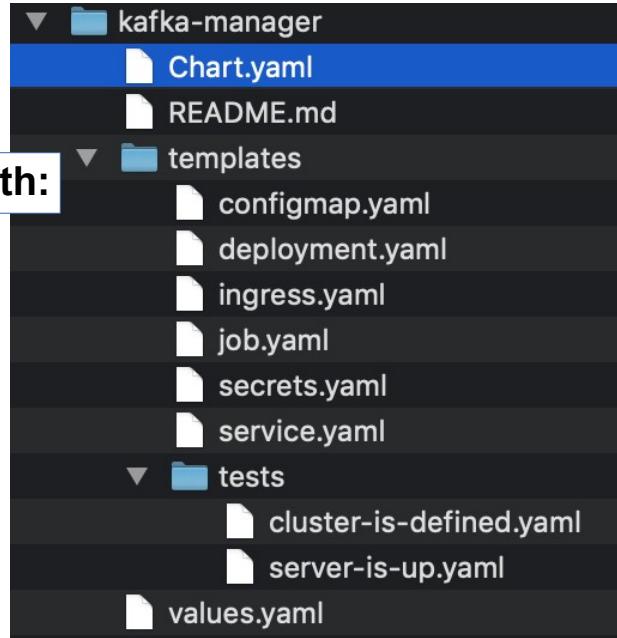
- Chart metadata
  - Default configuration values
  - Other support files (scripts, icons...)

- **Chart Configuration**

- Defaults (inside the Chart)
  - External configuration (using Helm CLI)
    - 0+ config value files (YAML)
    - 0+ overridden values

- **Helm Release = Chart + Configuration**

**== Chart Instance (Identified by a Release Name)**



```
helm install
--values global.yaml
--values environment.yaml
--set service.port=5432
stable/postgres
```



# Helm :: Concepts

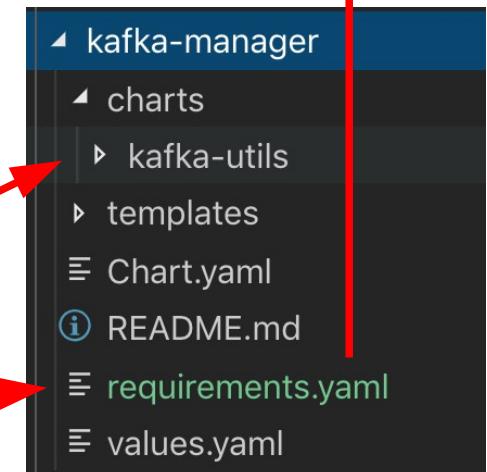
- **Chart Repository** == Glorified .tgz file server + /index.yaml with Chart metadata

```
dependencies:  
  - name: postgresql  
    version: 0.13.1  
    repository: https://kubernetes-charts.storage.googleapis.com/  
    condition: postgresql.enabled
```

- **Chart Dependencies**

== Other Charts, required by the main one, that provide additional content for the final Helm Release created

- Dependent charts can be:
  - Inline: Nested in the main chart
  - Referenced: In requirements.yaml
    - Fetched at build time from Chart Repositories



# Helm :: Concepts

- **Helm Release Revision**

== An update on a Helm Release deployed to K8s

- Required after a change in the Chart or the Config

Chart v1 + Config v1 → Release Revision #1

Chart v2 + Config v1 → Release Revision #2

Chart v3 + Config v2 → Release Revision #3

...

# Helm :: Concepts

- **HR Revision History**

- History kept inside K8s using ConfigMaps (do not touch!)
- Can go back and forth in History with Helm CLI
  - `helm install <releaseName> <chart+config>` → Revision #1
  - `helm upgrade <releaseName> <chart+config>` → Revision #2
  - `helm upgrade <releaseName> <chart+config>` → Revision #3
  - `helm rollback <releaseName> 2` → **Revision #4 (!!?)**

Back to v2 release, but creates a new revision!

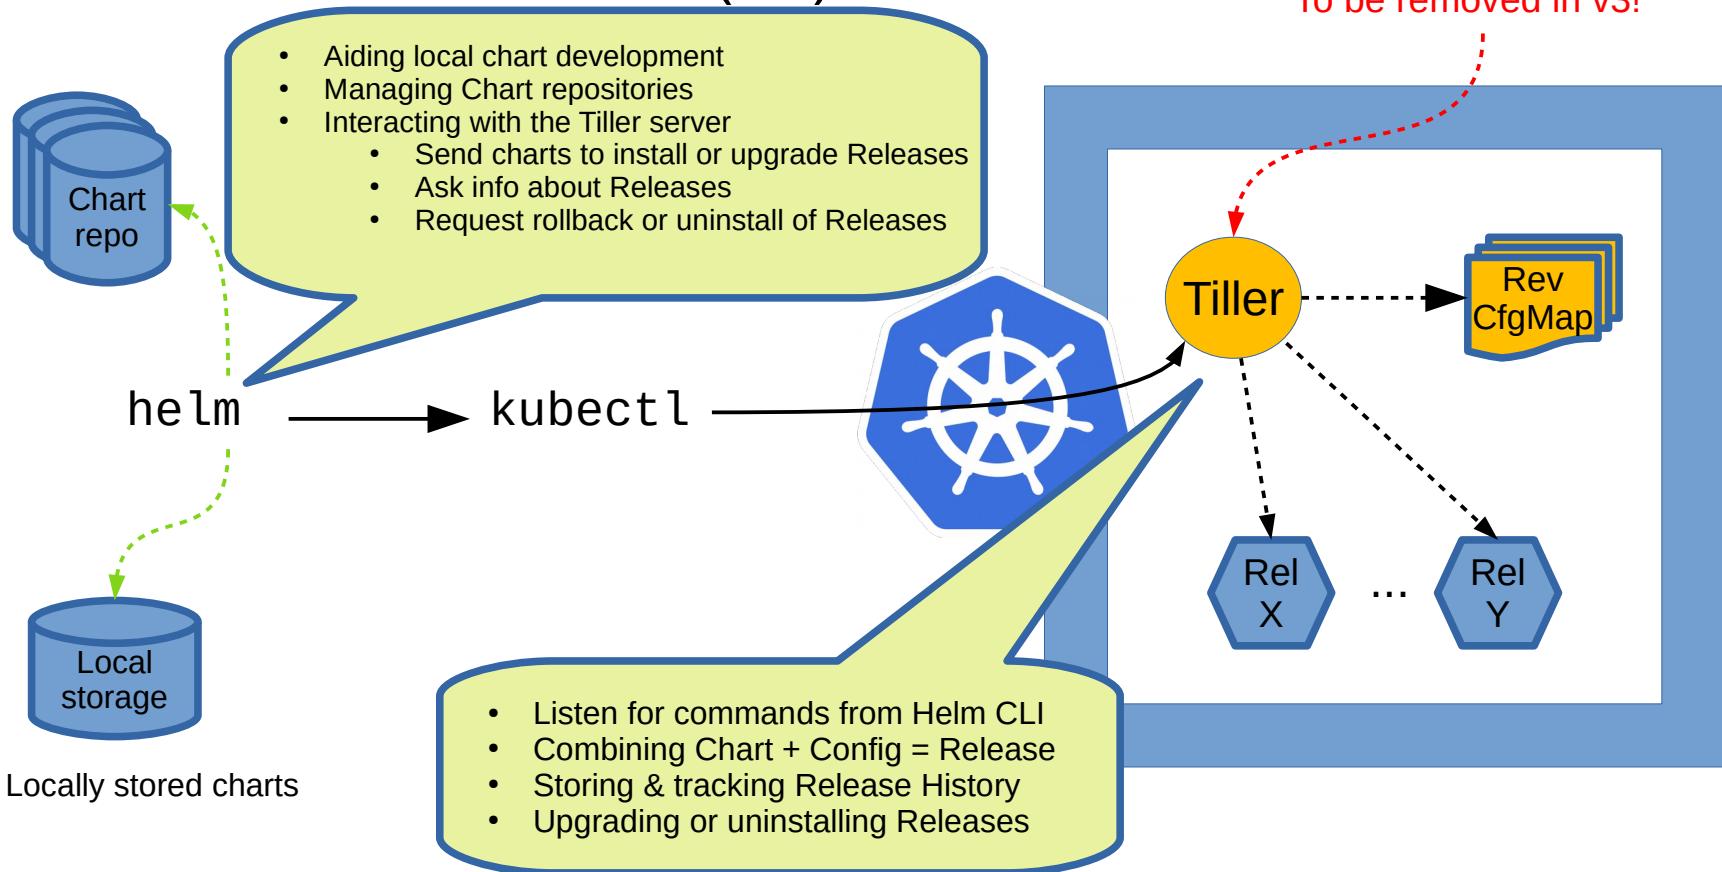
```
kubectl get configmaps -n kube-system
```

NAME	DATA	AGE
angry-bronco.v1	1	67s
angry-bronco.v2	1	27s
angry-bronco.v3	1	20s
angry-bronco.v4	1	10s

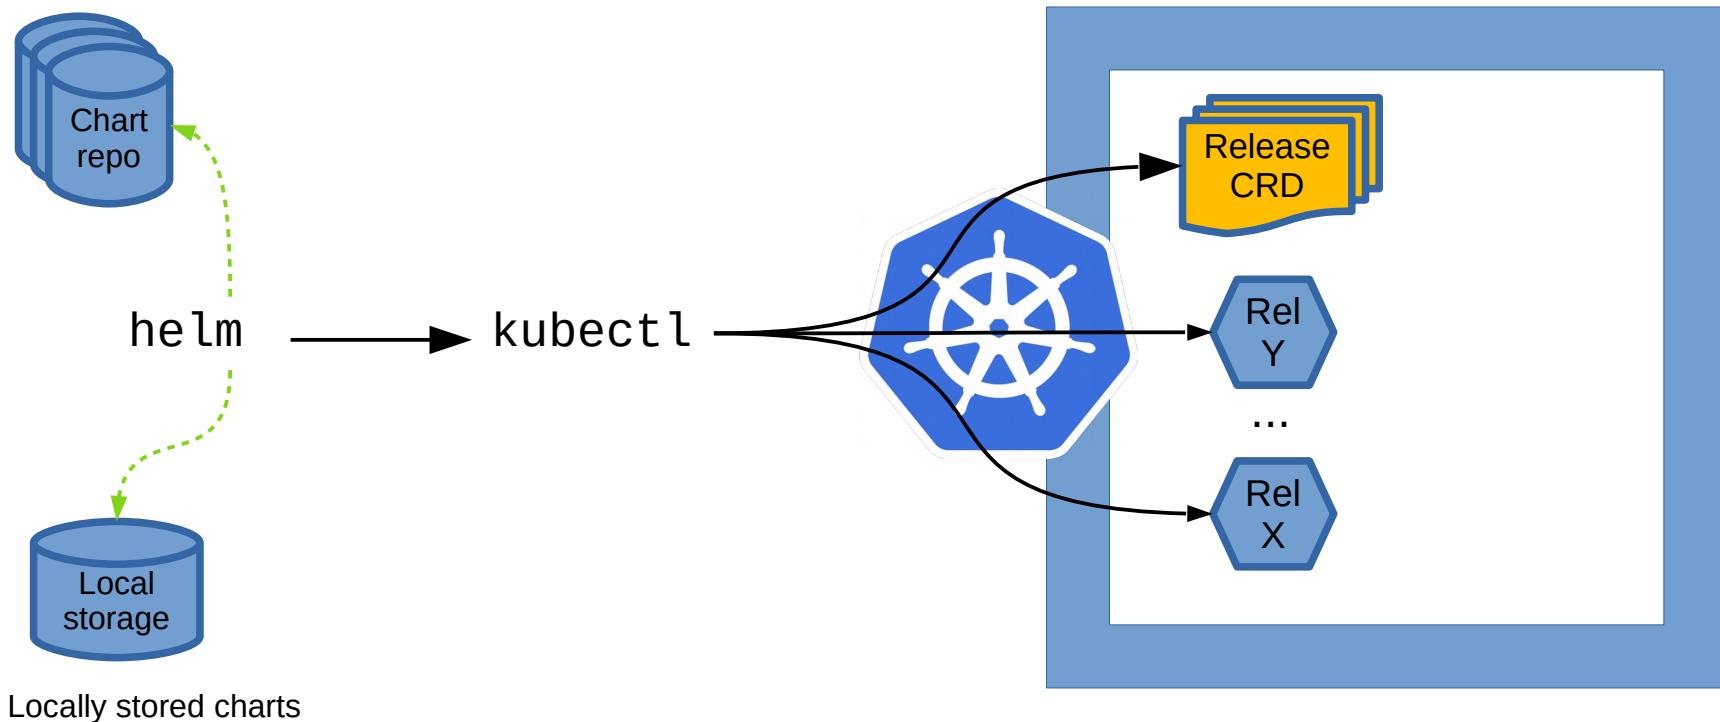
NAME	DATA	AGE
angry-bronco.v1	1	4m46s
angry-bronco.v2	1	4m6s
angry-bronco.v3	1	3m59s
angry-bronco.v4	1	3s



# Helm :: Architecture (v2)



# Helm :: Architecture (v3)

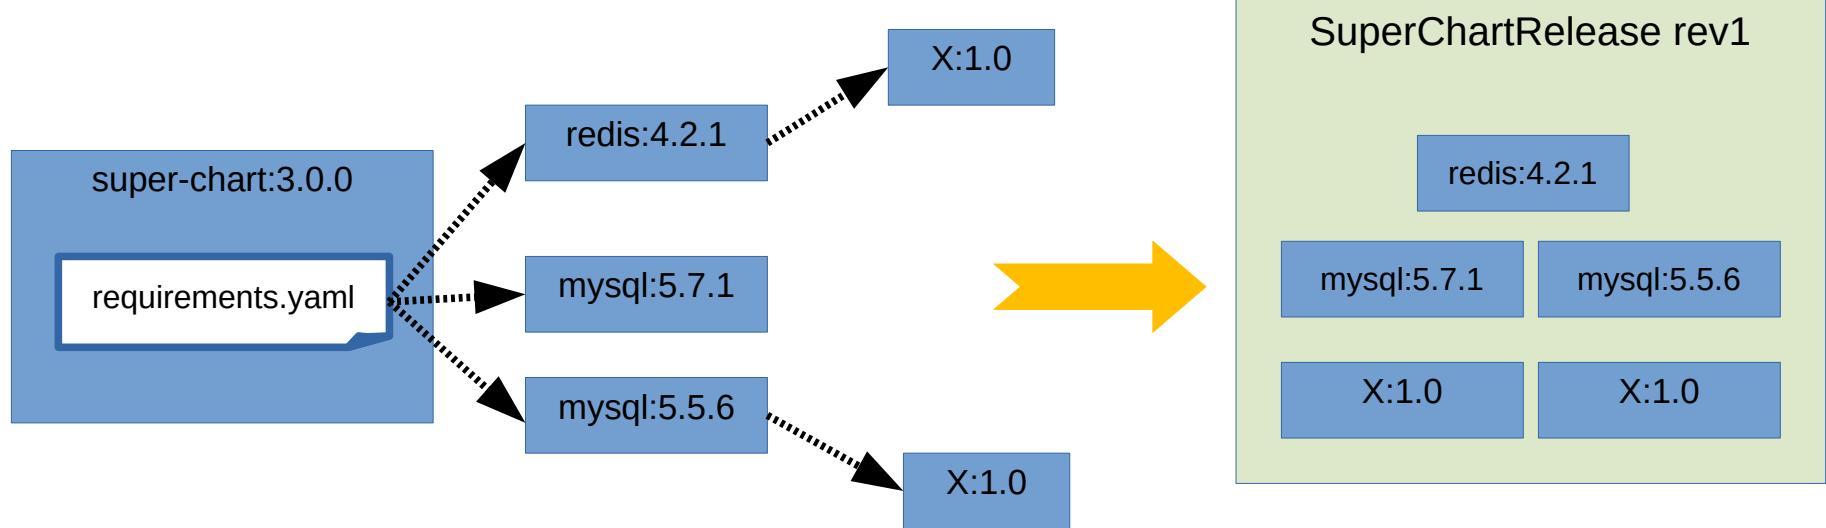


# DEMO #2

Helm Chart development workflow

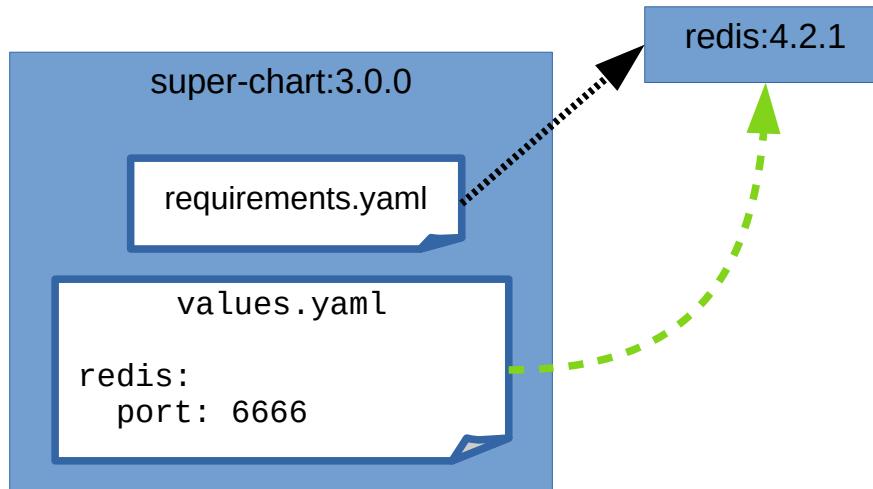
# Helm :: Feature Highlights

- **Package Dependency Management**
  - requirements.yaml == Non-shared build-time dependencies



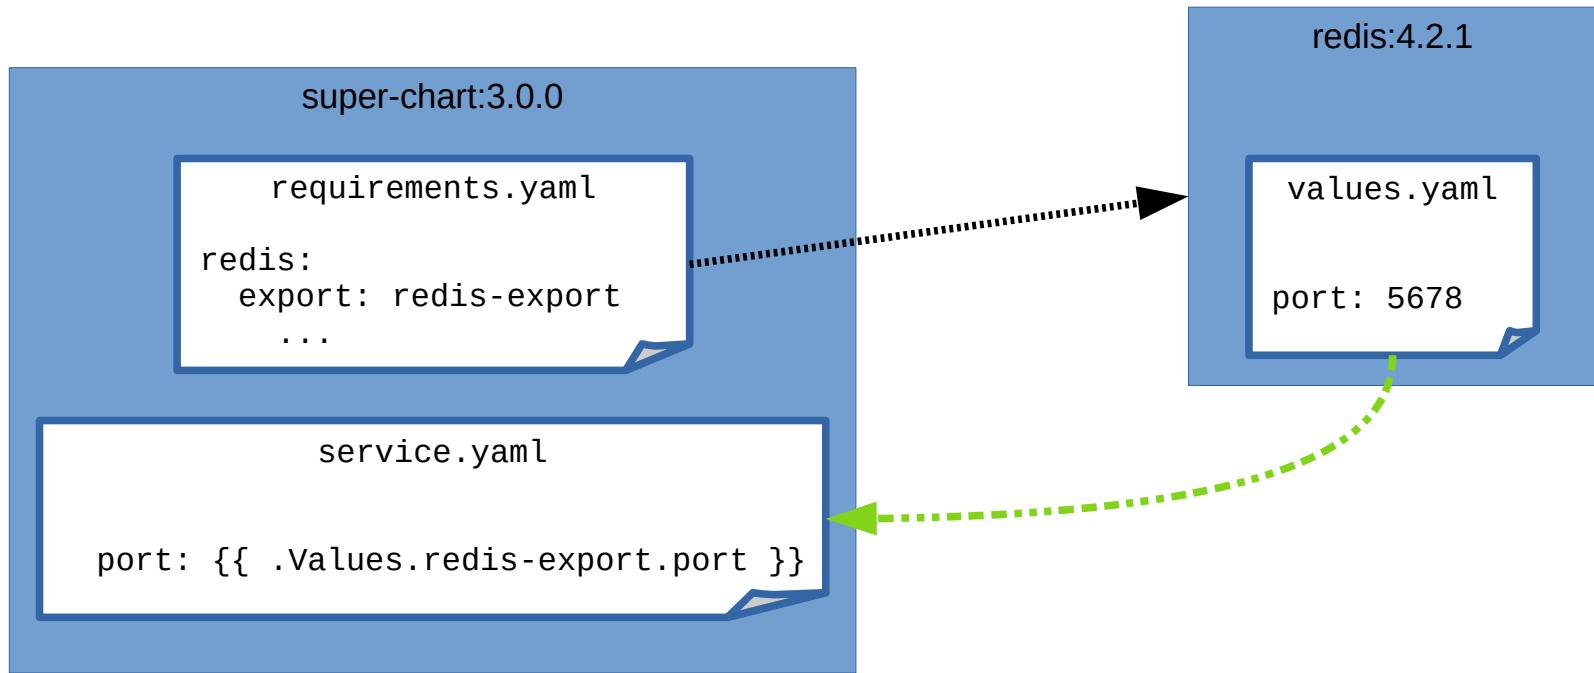
# Helm :: Feature Highlights

- Subchart config override



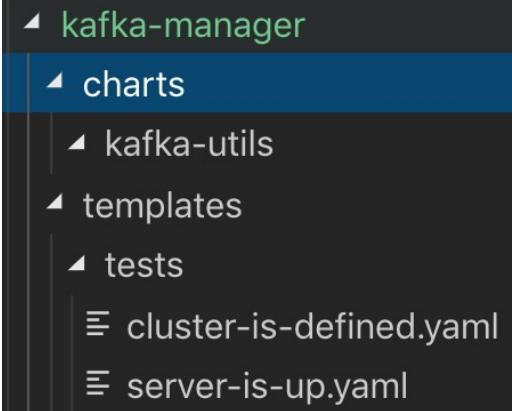
# Helm :: Feature Highlights

- Subchart config exports



# Helm :: Feature Highlights

- **Chart Tests (I)**
  - Kube Pod manifests, defined under: templates/tests/\*\*
  - Not deployed on Chart install
  - Stored with Release Revision data (Tiller ConfigMaps)
  - Deployed on command: `helm test <ReleaseName>`
  - **Test success == All Pod containers exit with (success | failure)**
    - Depends on per-test config



# Helm :: Feature Highlights

- **Chart Tests (II)**
  - **Recommended use: Smoke Tests**
    - Helps encapsulate knowledge about how to check if an application is working as expected after deployment
      - **Allows Dev staff to provide Ops staff with app-specific sanity checks they can run when in trouble**
    - Can serve as after-deployment tests in CD pipelines
      - E.g. Rolling back to the previous Release Revision on tests failure

# Helm :: Feature Highlights

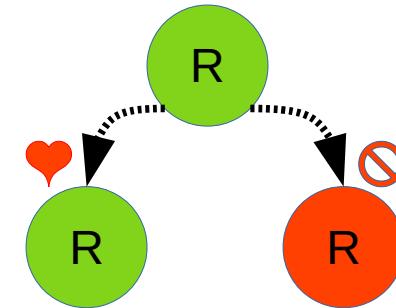
- **Chart Tests (III)**
  - **Drawbacks:**
    - **Cannot filter the tests to be executed with the “helm test” command**
      - Can be done in a home-made fashion, but it's hard
      - It's a problem for big Releases, composed of many Chart deps
        - **Big problem for UberChart pattern! (see later)**
    - **Lack of support for sidecar containers during tests**
      - All containers in the test Pod are considered tests
        - i.e. All of them must exit with success/failure according to test cfg
      - Can't have sidecar containers with ignored end status and auto killed
        - i.e. Must ensure sidecars fail if tests fail, and vice versa
          - Can be done in a home-made fashion, but painful!

# Helm :: Non Features

- **Tiller is not a Helm Release Operator**
  - Releases are not auto redeployed if Release resources are deleted manually
- **Cannot specify the deployment order of sub-Charts in a Release**
  - Applications deployed all at once
    - Must be able to wait for dependencies to be up → “Limbo” time!
- **Can't reference dependency Charts by Git repo URL**
  - Must be published in a Chart repo
    - Chart repos sometimes unnecessary for corporate Charts
    - But always forced to package Charts and maintain a Chart repo

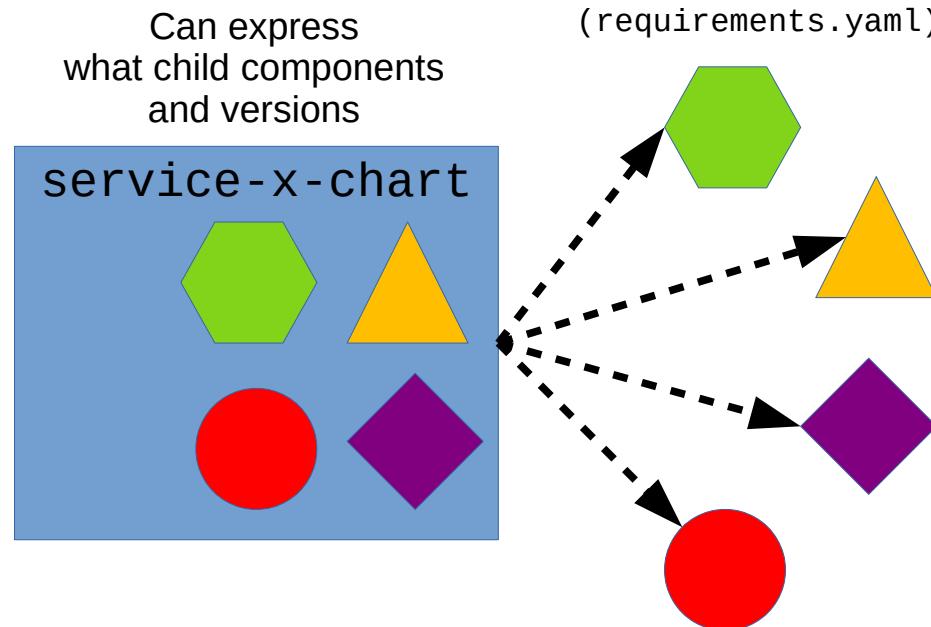
# Helm :: Non Features

- **No Release Dependency Management support**
  - Helm  $\approx$  Docker && Helm  $\neq$  DockerCompose
  - Chart dependencies in requirements.yaml are build-time deps → All part of a single Release
  - Release Dependencies = Runtime dependencies
    - Dependencies not included in the Release
    - But required to be there at run time for the Release to work properly
  - No support for deploying multiple interdependent Releases
    - Can be done using IaC tools (e.g. Ansible with “helm” module)
  - No support on Release deployment to check for depended-on Releases to be in place
    - Can be done in a custom-made fashion using Helm pre-install hooks
    - But must do it for every chart



# Helm :: Non Features

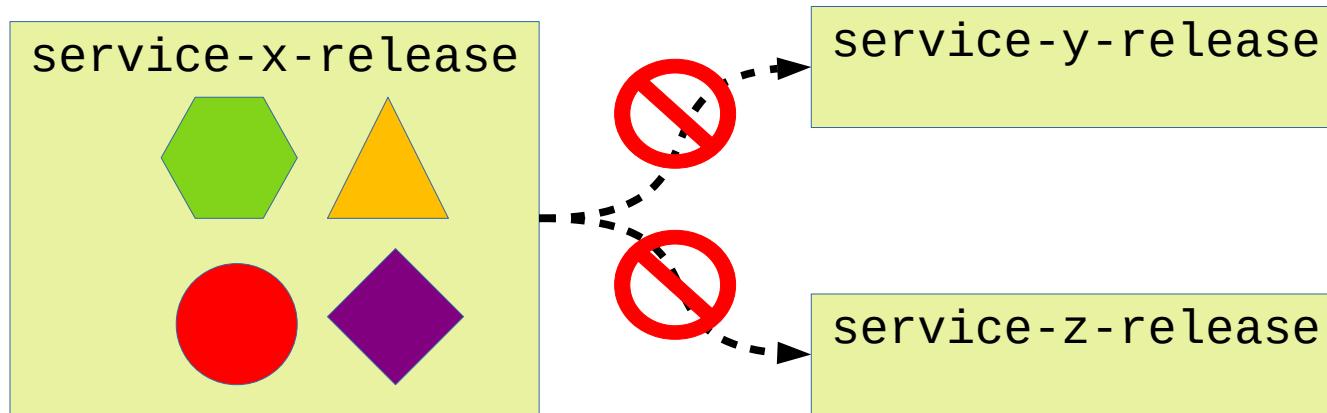
- **No Release Dependency Management support**



# Helm :: Non Features

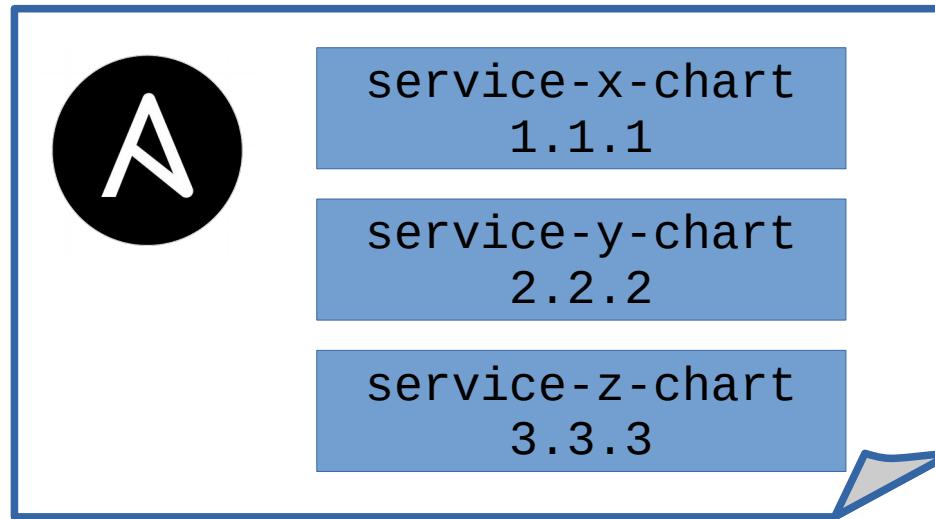
- **No Release Dependency Management support**

Charts cannot express what sibling releases and versions  
→ **Releases can be installed unlinked**



# Helm :: Non Features

- **No Release Dependency Management support**
  - Need a higher level tool to specify the right Chart versions to keep system consistency



# Helm :: Non Features



- **No Unit Test support**

- Unit Testing in Helm == Testing conditional Chart template generation
  - Input == Config values + Chart templates
  - Output == K8s manifests (YAML) ← Assert expectations!
- Specially useful for “library” Charts (i.e. designed for reuse)
  - Not so much for “system definition” Charts (i.e. specific system impl.)
- External tools available:
  - **HelmTest plugin**
    - Runs “helm template” command and asserts on YAML outputs
  - **Terratest**
    - Multipurpose Go library for infra testing
    - Includes test helpers to assert on Helm-generated templates

# Helm :: Non Features

- **No Integration Test support**

- Helm tests
  - Kinda similar to Int Tests... (runs checks on the deployed application)
  - But can only run tests where the Release has been deployed
    - **Test not in control of the deployment scenario → Need external tool for that**
- External tools available
  - **Terratest**
    - Multipurpose Go library for infra tests
    - Supports Helm Chart deployment & asserting on the deployed app
  - **KitchenCI**
    - Multipurpose test env generator & test executor
    - Can provision local VM with Minikube & deploy with Helm CLI
    - Supports several infra testing libraries to assert on the deployed app

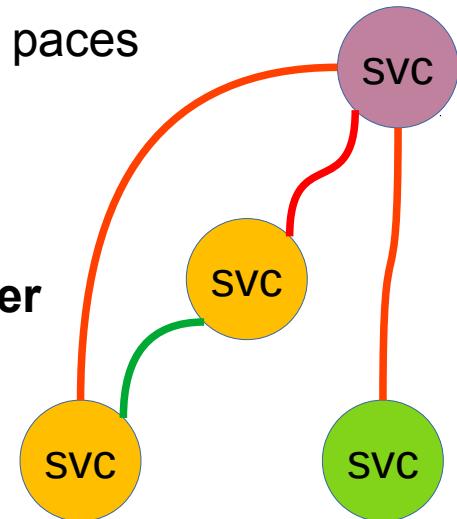


# Helm :: Alternatives

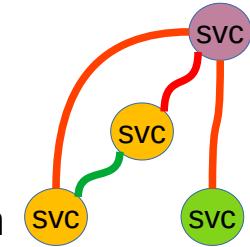
- **Kustomize**
  - Integrated into kubectl tool since K8s 1.14
  - Only for manifest templating
    - Other Helm features not supported
  - YAML-friendly syntax, cleaner than Helm's Go Template syntax
- **Helm + Kustomize**
  - Can use all other Helm features, but do templating with Kustomize
    - 1) Run “helm template” → Outputs the generated K8s manifests
    - 2) Run Kustomize over the generated manifests → Generates final manifests
    - 3) Send manifests to Kube with “kubectl apply”
  - Helps avoiding the more convoluted uses of Go Template, and still have Package Mgmt

# Distributed System Consistency

- **Distributed System**  
== A set of interdependent, network connected Services
- **Services:**
  - Are usually developed by different teams, at different paces
  - Need to be able to work correctly with each other
    - Need to keep deployed always the right versions
      - **Every system change must result in another set of compatible versions**
      - **We need to specify somewhere that versions set**



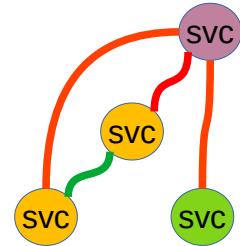
# Distributed System Consistency



- Besides that, **we should always be able to reconstruct our system from scratch** to recover from disasters.
- For that you need:
  - **A declarative specification of your current system state**
    - What applications, and which versions
    - What configuration
  - **A tool to build the system from the declarative spec**
    - **Can we do that with Helm?**
    - **Problem: Helm only deploys individual releases**

# Distributed System Consistency

- Option #1: Single Super Chart (to rule them all)



Deep tree (= monorepo )

```
capsule-inc-global/  
  charts/  
    capsule-inc-hhrr/  
      charts/  
        ...  
    capsule-inc-research/  
        ...  
    capsule-inc-accounting/
```

Shallow wrapper

```
capsule-inc-global/  
  requirements.yaml  
    --> capsule-inc-hhrr  
    --> capsule-inc-research  
    --> capsule-inc-accounting
```



Single Super Helm Release !

# DEMO #3-1

One Super Chart to rule them all

# Distributed System Consistency

- Option #1: Single Super Chart (to rule them all)

- Pro:

- No extra tool needed, just Helm to build the full system

- Con:

- Service charts not independently manageable

- Some chart metadata becomes corrupted...

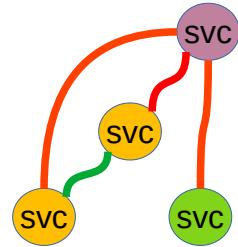
- Release needs to be deleted...

- Full system takedown!!

- Independent teams wanting to do their own thing...

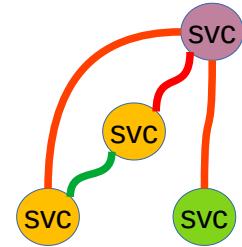
- Unable to smoke-test a single malfunctioning sub Chart

- “helm test” command cannot filter tests!



# Distributed System Consistency

- Option #1: Single Super Chart (to rule them all)
  - Con:
    - No built-in way to test if Service charts have all their runtime dependencies (i.e. other Service charts) satisfied
      - Chart metadata doesn't support expressing runtime dependencies
      - Helm doesn't implement checking runtime deps
      - Can be done, but in a custom way

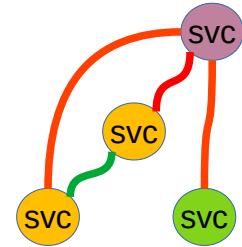


# Distributed System Consistency

- Option #2:

## Orchestrated deployment of multiple Helm Releases

- Full System = 2+ Helm Releases
  - Usually 1 Release = 1 Logical Service (funct cohesion)
    - Logical Service = 1+ closely integrated applications
- **Requires: Release orchestration tool**
  - Simple shell script
  - Any IaC tool
    - e.g.: Ansible + “helm” Ansible module
    - **Helmfile !**

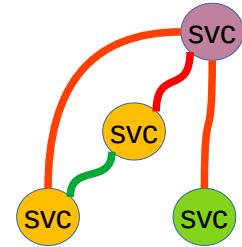


# Distributed System Consistency

- Option #2:

## Orchestrated deployment of multiple Helm Releases

- Con:
  - Need an extra tool
  - Still no support for runtime dependency management
- Pro:
  - Independently manageable Helm Releases !



# Distributed System Consistency

- Option #2:

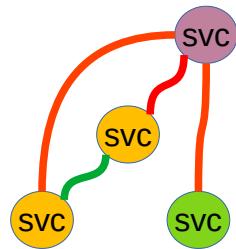
## Orchestrated deployment of multiple Helm Releases

- Example tool: **Helmfile**

- **Features:**

- Declarative Helm CLI automation
    - Go Template support (like in charts! → familiarity)
      - Allows conditional Release installation & config
    - Environment definition support → Grouping of config values per env
    - Infinitely nestable (Helmfiles can include other Helmfiles)
    - Configurable Release deployment order
      - Releases are deployed by their order in the Helmfile

```
releases:  
- name: kafka-manager  
  chart: ./kafka-manager  
  namespace: transverse  
  labels:  
    release-name: kafka-manager  
  values:  
    - ./values-global.yaml
```



# DEMO #3-2

Simple Helmfile

# DEMO #3-3

Complex Helmfile

# Distributed System Consistency

- Option #2:

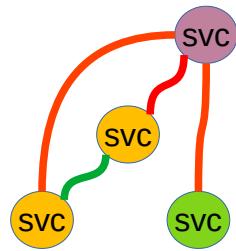
## Orchestrated deployment of multiple Helm Releases

- Example tool: **Helmfile**

- **Drawbacks:**

- An extra tool to learn/install/use
    - No support for Helm v3 yet
    - Beta! (but used in Prod by many companies)
    - Can't reference Charts directly by Git repo URL
      - Just delegates to Helm behind the scenes, and Helm doesn't support it.
    - Configurable Release deployment order
      - Releases are deployed by their order in the Helmfile

```
releases:  
  - name: kafka-manager  
    chart: ./kafka-manager  
    namespace: transverse  
    labels:  
      release-name: kafka-manager  
    values:  
      - ./values-global.yaml
```



# Current Status

- Managing 3 tools for app deployment

helmfile

Release orchestration

helm

Release installation

kubectl

K8s manifest execution

... can we do better?

# Current Status

- Managing 3 tools for app deployment

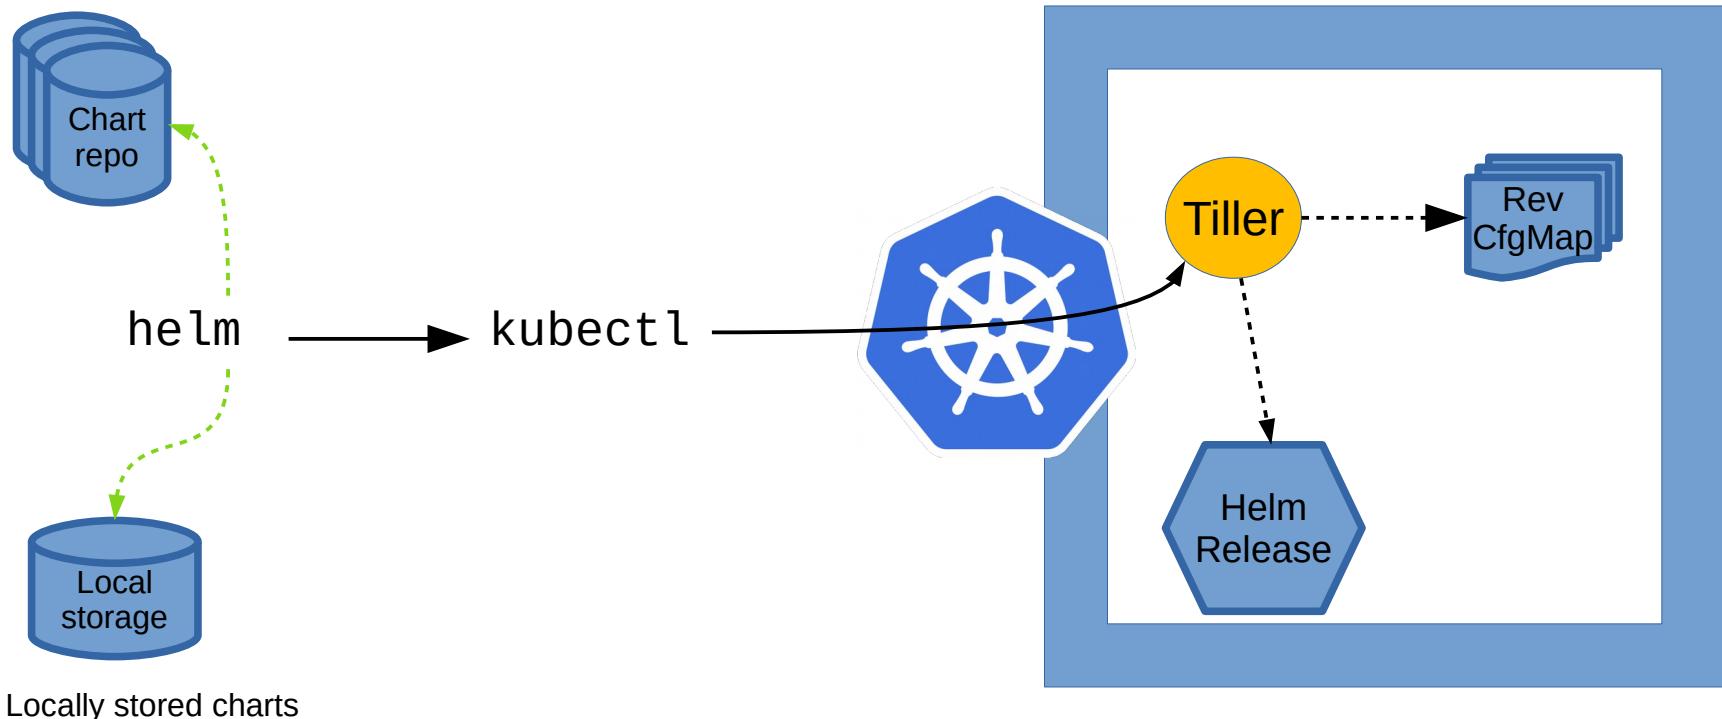
helmfile	Release orchestration
helm	Release installation
kubectl	K8s manifest execution

... can we do better?

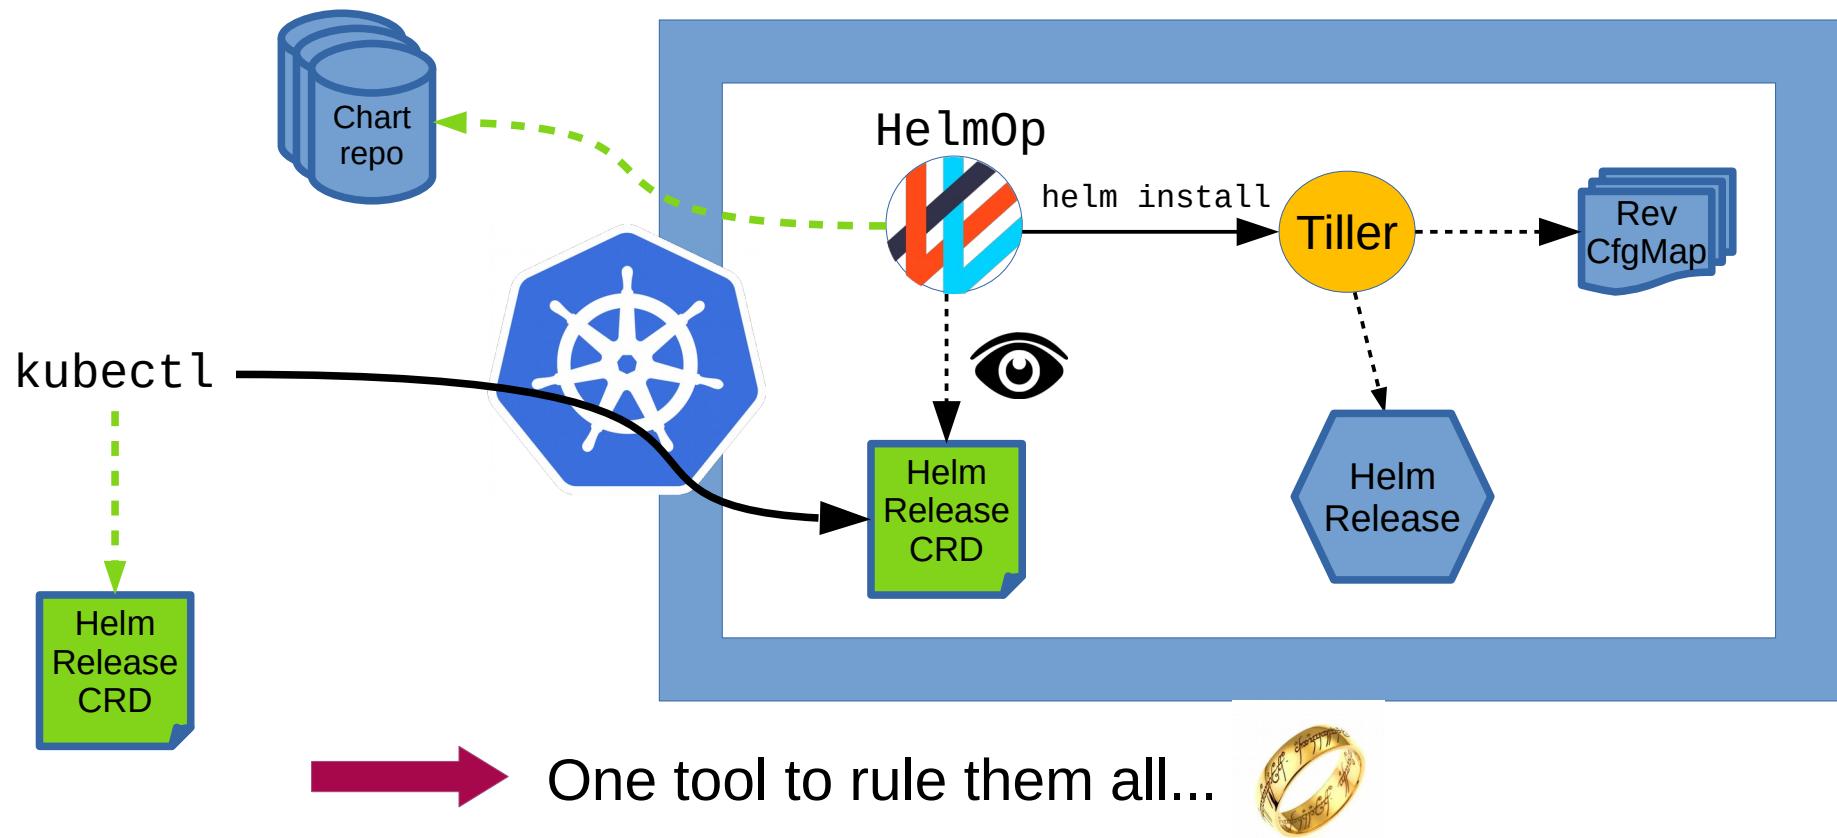


Yes, by needing less tools

# Before: Standard Helm



# Enter Weaveworks' Helm Operator



# Helm Operator explained

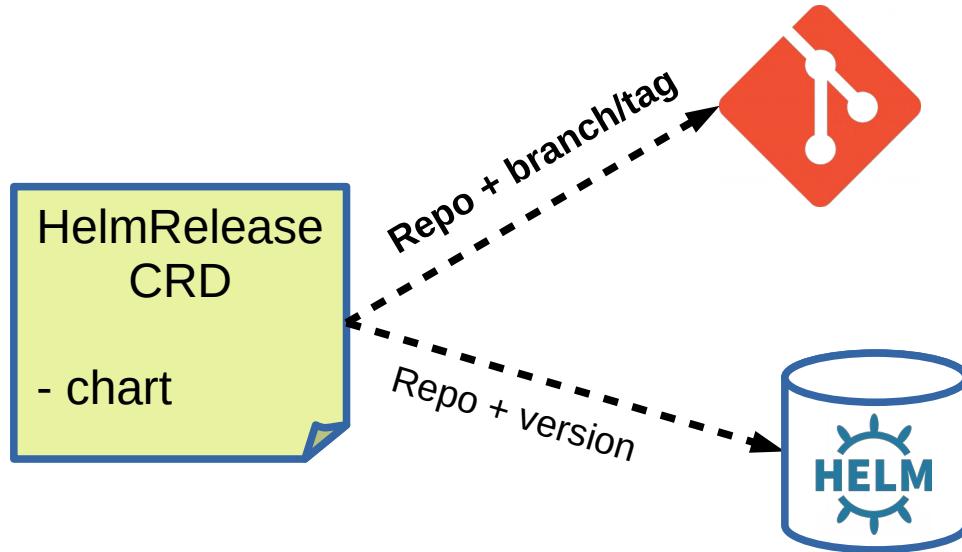
- **Operator application**
  - User creates K8s CRD's of type HelmRelease (defined by HO)
  - HO watches for HR objects and interacts with Tiller to install or upgrade Releases
- **Can run inside (usual) or outside Kube (reuse and/or security hardening)**
- **Centralizes:**
  - Installation of the Helm CLI tool
  - Access to Helm chart repos and their credentials
  - Auditing of deployments (who = Helm Release CRD creator)



**Simplified user setup and authorization**

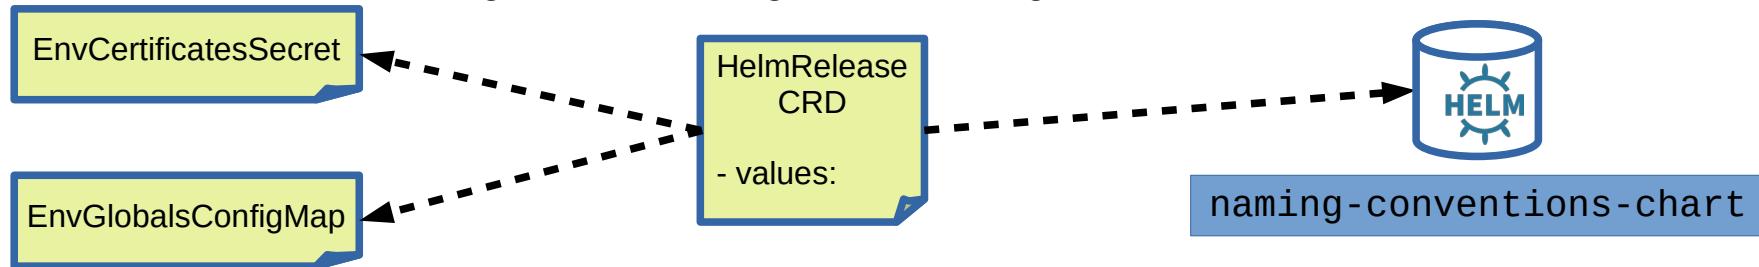
# Helm Operator explained

- HelmRelease CRD allows referencing Helm Charts by:
  - Helm chart repo reference
  - Git repo URL → **Missing feature in Helm!**
    - No need to package and upload to chart repo!



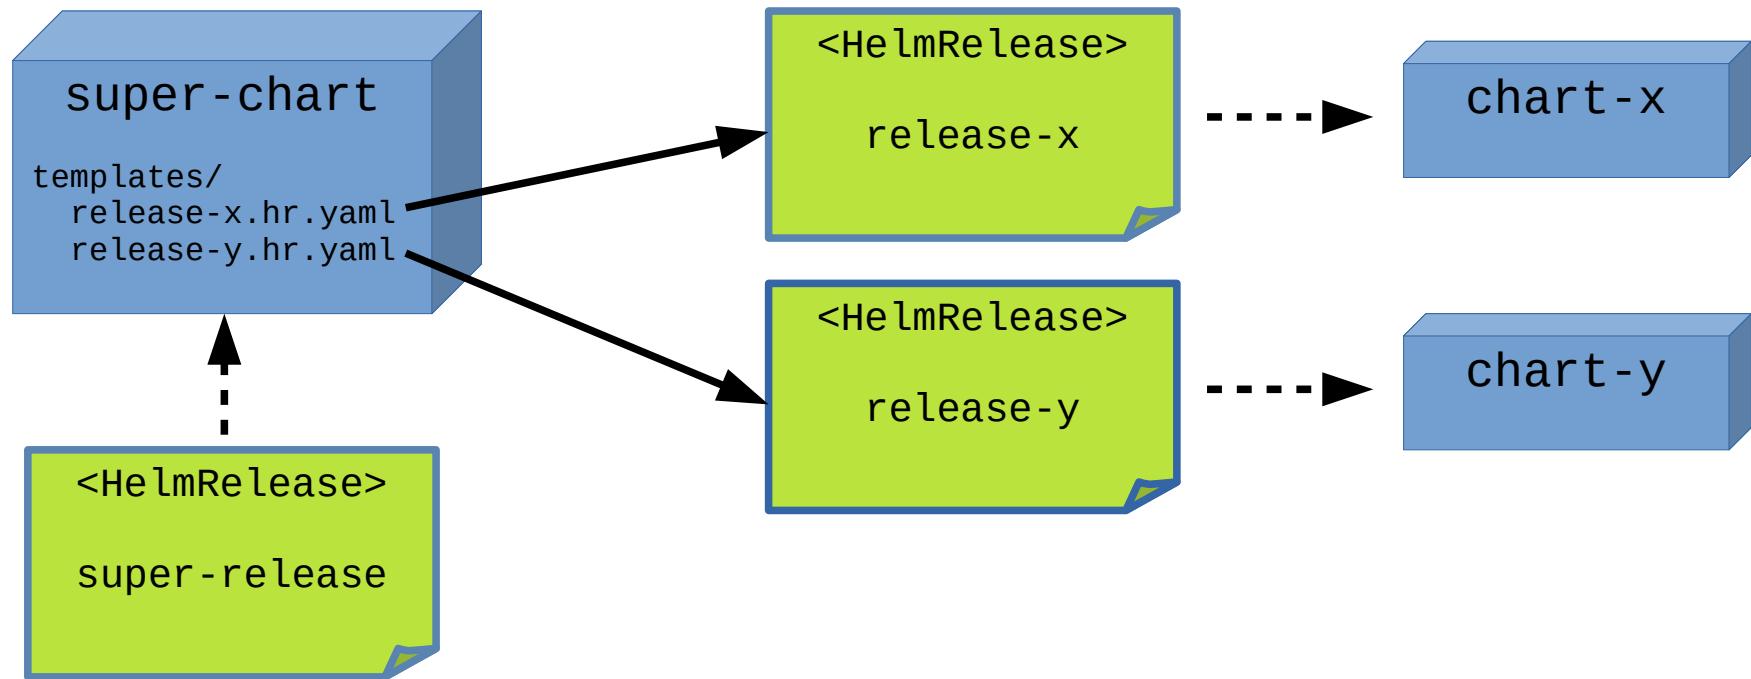
# Helm Operator explained

- Much higher composability for a chart's config values
  - Config value sources (can use many)
    - ConfigMaps, Secrets
    - External file (URL)
    - Another chart
  - Example uses:
    - Deploy a single ConfigMap for all environment-specific configs, and have all charts use it as config value source
    - Sharing common config values among all the charts



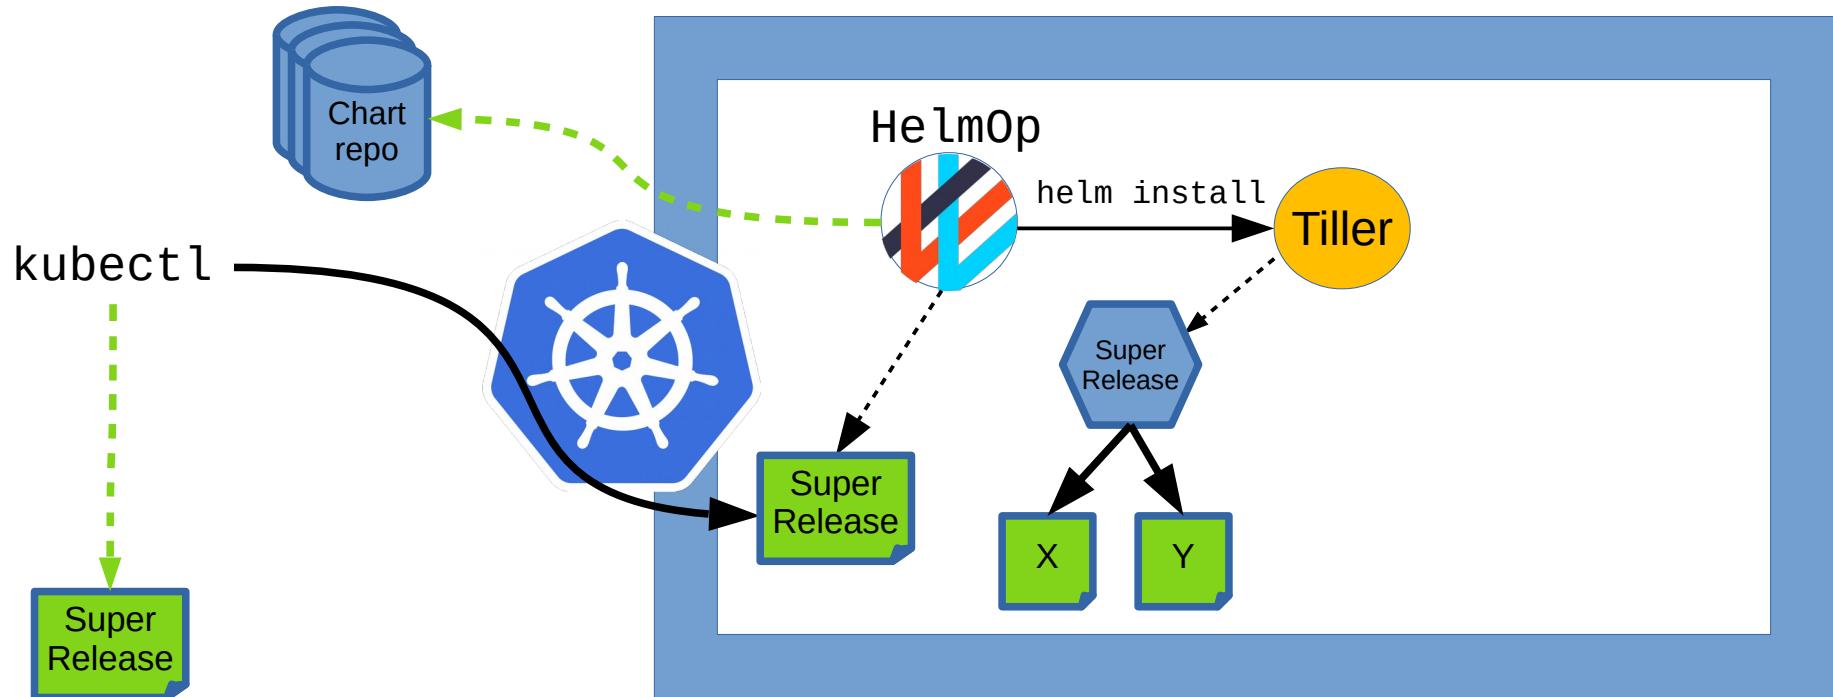
# Helm Operator explained

## Helm Release Orchestration revisited



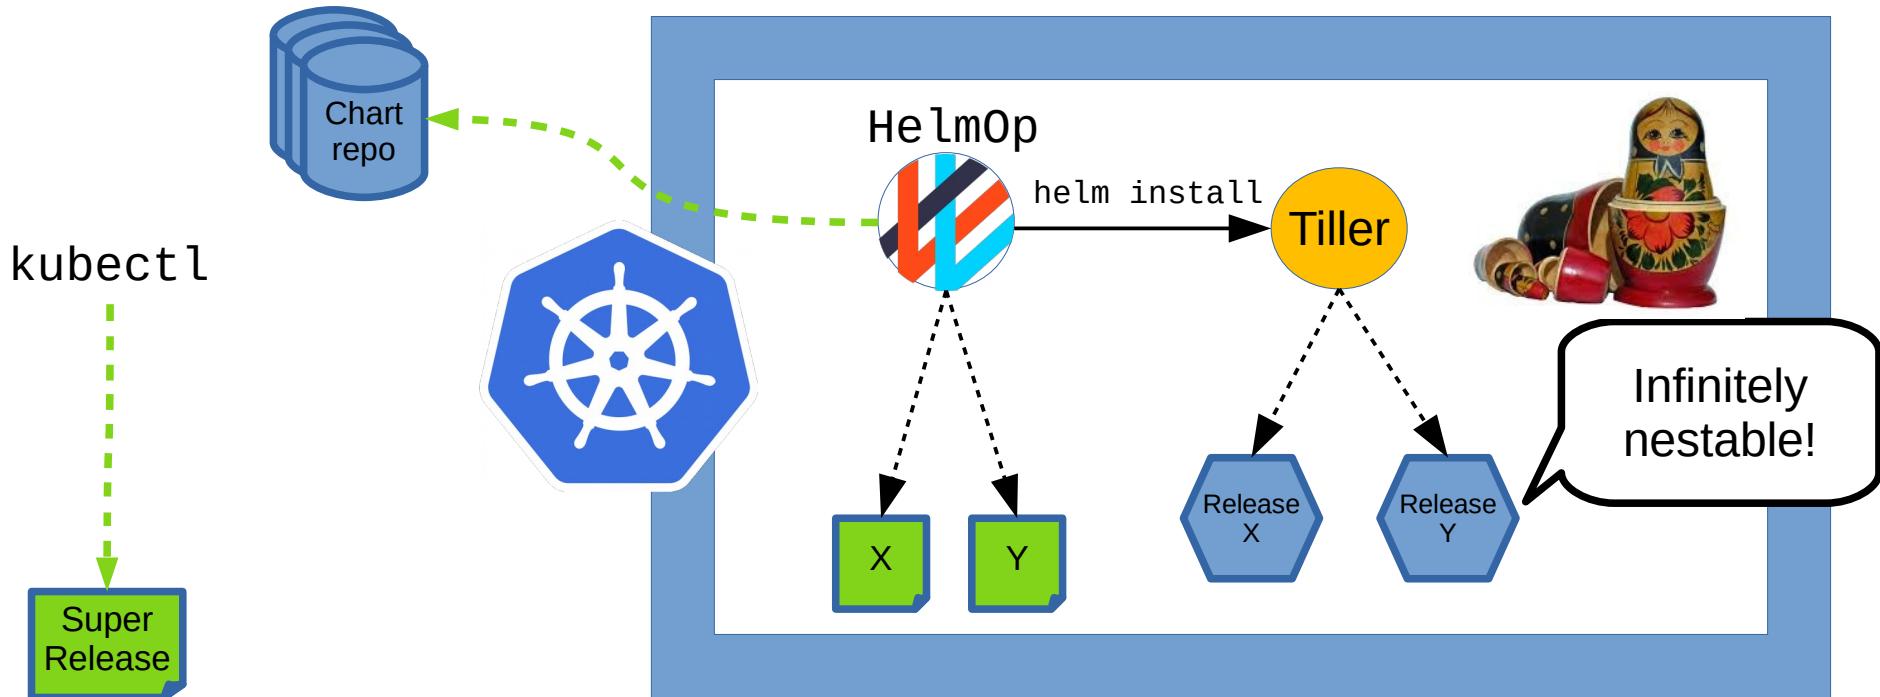
# Helm Operator explained

## Helm Release Orchestration revisited



# Helm Operator explained

## Helm Release Orchestration revisited



# DEMO #4

## Helm Operator

# Helm Operator explained

## Drawbacks

- **Still delegating finally to Helm/Tiller to perform the deployment**
  - Still carrying Helm's limitations
    - Can't specify HelmRelease deployment order
    - Can't filter executed tests in a single HelmRelease
- **The chained sequence of HelmRelease deployments makes “limbo time” longer**
  - “Limbo time”
    - = The time that a deployed app might have to wait to its runtime dependencies (other apps)
  - During “limbo time” apps might get restarted because of missing runtime deps
    - Introduces “noise” in logs during deployment
- **We no longer have a way to trigger ALL the tests (no link between HelmReleases)**
  - Need to loop all the HelmReleases and test

# Quick status recap...

- **Deploying Helm charts just using kubectl <== Sweet!**
  - No need to install Helm / Helmfile, just kubectl
  - No need to have locally configured access to chart repos, chart git repositories, etc. All centralized in the Helm Operator deployment.
- **Can do deployments at many levels of granularity**
  - Extremely coarse grained: Single superchart, single Helm release
  - Extremely fine grained = 1 Helm Release per application
  - Recommended:
    - Superchart as the global system definition
      - Contains 1 HelmRelease per Logical Service
      - Contains global config values shared by all LS's

# **Can we do better?**

# Yes, we can!



- **One single tool (`kubectl`) is still one too much**
  - Ops still need to configure `kubectl` locally
- **Auditability**
  - We have no centralized log about who deployed what and when
- **Need careful access control for Ops**
  - Because Ops must use `kubectl`, permissions must be carefully assigned to avoid security issues.

# Yes, we can!



- **Lack of a consistent workflow**
  - **System consistency**
    - We have no process to ensure that teams don't step over each other in the deployments (e.g. breaking runtime dependencies between deployed HelmReleases)
  - **System change promotion between environments**
    - We have no process to ensure an orderly progression of system changes from one environment to the other

# Yes, we can!



- **Lack of a consistent workflow (II)**
  - **System divergence control**
    - We have no way to know when the current system state is diverging from the desired state
      - e.g. Ops making system changes in troubleshooting scenarios, but forgetting to translate those changes into Git
        - Relying on Ops' discipline & good practices → Danger!

# Yes, we can!



- Lack of a consistent workflow (II)
    - System divergence control
      - We have no way to control the current system state is in sync with the desired state
        - e.g. we can detect system changes in various operating scenarios, but we have no way to translate those changes
- Killing the CD Pipeline!  
Relying on Ops' discipline & good practices → Danger!

# Yes, we can!



- Lack of a consistent workflow (II)
  - System divergence control
    - We have no way to control the system changes in the current system state is not aligned with the desired state
      - e.g. system changes in testing scenarios, but translating those changes relying on Ops' good practices → Data

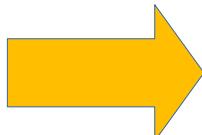
CD Pipeline!



# Is a CD pipeline enough?

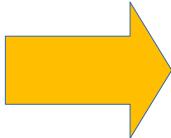
- One single tool (`kubectl`) is still one too much
  - Ops still need to configure `kubectl` locally
- Need careful access control for Ops
  - Because Ops must use `kubectl`, permissions must be carefully assigned to avoid security issues.



- 
- Yes, deployments centralized by the CD pipeline
    - Less need for `kubectl` & less privileges needed
  - However, if Ops still need `kubectl` for fine tuning or troubleshooting apps or Kube resources, we are stuck

# Is a CD pipeline enough?

- **Auditability**
  - ~~We have no centralized log about who deployed what and when~~
    - Yes, now the CD pipeline tooling provides a tracing mechanism for deployments
    - However:
      - If Ops still use `kubectl` for tweaking or troubleshooting
        - We have no log about those actions
      - Sometimes CD pipelines are just about individual applications, not global system definition
        - No single place to track global system evolution
          - “When did we have service X working with service Y?”



# Is a CD pipeline enough?

- Lack of a consistent workflow

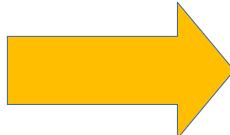
- System consistency



- We have no process to ensure that teams don't step over each other in the deployments (e.g. breaking runtime dependencies between deployed Helm Releases)

Yes, we can achieve that

- As long as CD pipelines perform integration or smoke tests after deployment, and rollback in case of failure
  - If they don't rollback, temptation to leave it broken is too strong

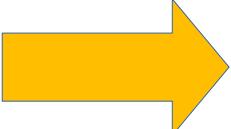


# Is a CD pipeline enough?

- Lack of a consistent workflow
  - System change promotion between environments
    - We have no process to ensure an orderly progression of system changes from one environment to the other



**Yes, we can achieve that**

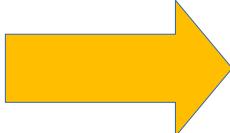
- 
- As long as we check for overall system consistency after deployment at each environment, and refuse promotion on failure
  - Example:
    - CI pipeline for X performs integration tests of X against Y with version v2, and CD deploys to INT & STAGING
    - Works fine in INT env, but STAGING env still has v1 of Y
- Integration tests pass, INT env works, but STAGING breaks

# Is a CD pipeline enough?

- Lack of a consistent workflow
  - System divergence control 
    - We have no way to know when the current system state is diverging from the desired state

Yes, we can achieve that, but not by default

- It's something very dependent on the deployment environment
  - Not provided by CD tool, unless it's specialized in env X
  - It's very usual to not have this, and after CD deployments the environments becoming the source of truth
    - “Would you be able to redeploy in Mars tomorrow?”



# Other CD pipeline considerations

- **Usually very ad-hoc**
    - Well-known mature CD tooling (Concourse, Jenkins, Spinnaker...)
    - But every org needing their own deployment pipeline
      - Different target runtime environments (K8s, AWS, Nomad...)
      - Different approval processes
- **Little chance for reuse across organizations**
- **Deployment process tightly coupled to the CD tool**
  - Hooking into the process usually means hooking into the tool

# Other CD pipeline considerations

- Sometimes CD pipelines are just about individual applications
  - 1 application → 1 CD pipeline
    - No coordinated effort to keep system consistency
      - Nothing prevents us from deploying incompatible services! ← Specially if integration tests are lacking...
- Usually CD tooling allows us diverging the system away from the System Definition
  - E.g. Your typical button to rerun an old deployment job
    - Environments become the real source of truth!
    - We stop being able to reproduce the system from scratch



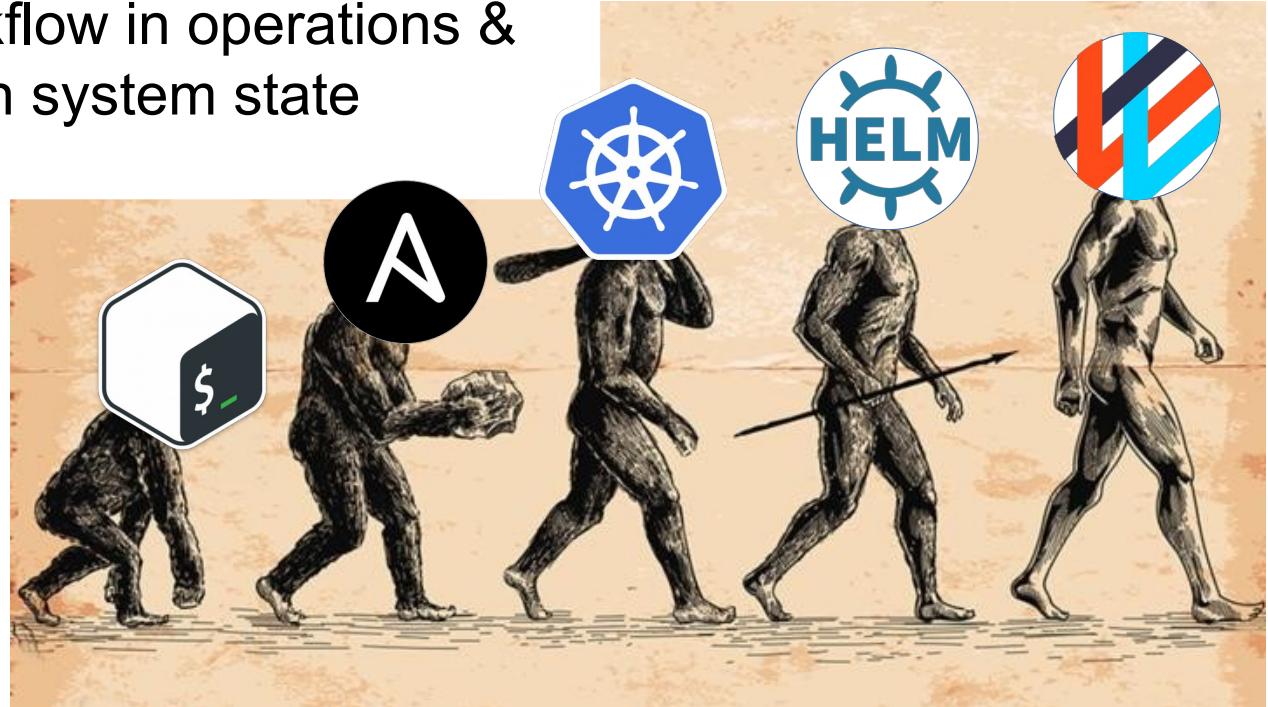
# The GitOps proposal

- Have a declarative Global System Definition in Git
- Allow system changes only through GSD changes
  - No direct operations on the System by Ops
    - That includes manual deployments, too
  - No rerunning old deployment jobs either!
- Leverage Git's transactional capabilities to enforce the proper workflow between teams for System changes
- Continuously monitor System divergence and automatically roll it back to latest accepted GSD



# GitOps

Enforcing workflow in operations &  
standardizing on system state  
evolution



# GitOps History

- 2014 – Weaveworks
  - SaaS CD solutions for K8s
- 2016 – Flux OSS project
- 2017 – GitOps article →
- Currently
  - GitOps = Pattern
  - Codefresh: GitOps 4 Terraform

The screenshot shows a blog post from the Weaveworks website. At the top left is the Weaveworks logo, which consists of a stylized 'W' icon followed by the word 'weaveworks'. On the far right is a three-line menu icon. Below the header, the date 'AUGUST 07, 2017' is displayed. Underneath the date are three blue links: 'Gitops', 'Kubernetes', and 'Product features'. The main title of the post is 'GitOps - Operations by Pull Request', centered above a horizontal line. Below the line, there is a paragraph of text: 'At Weaveworks, developers are responsible for operating our Weave Cloud SaaS. "GitOps" is our name for how we use developer tooling to drive operations.' A small decorative flower icon is located at the bottom right corner of the page.

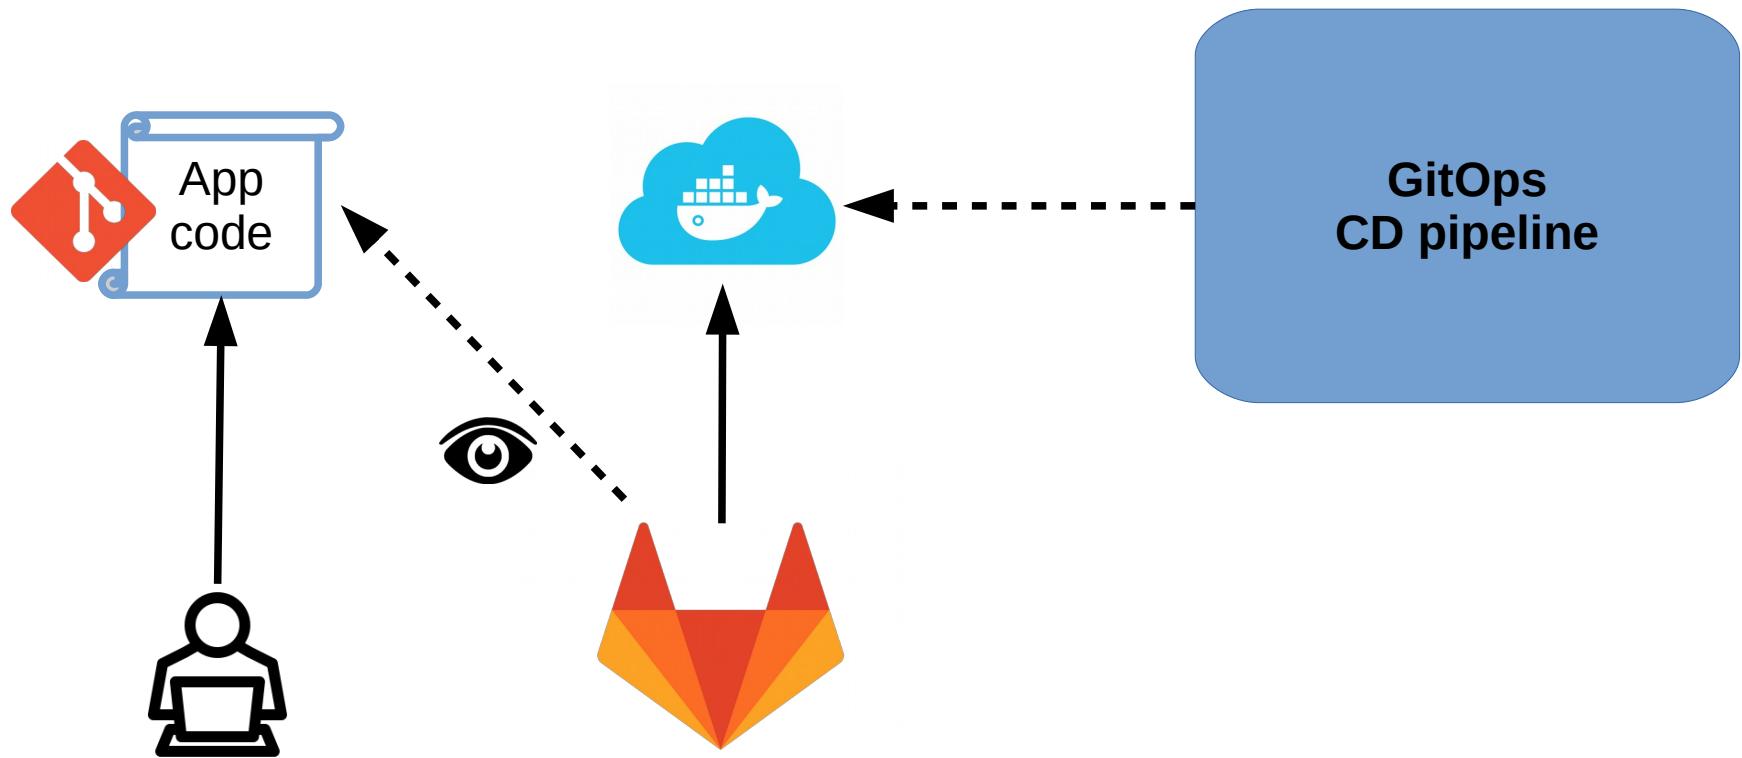
AUGUST 07, 2017

Gitops | Kubernetes | Product features

## GitOps - Operations by Pull Request

At Weaveworks, developers are responsible for operating our Weave Cloud SaaS. "GitOps" is our name for how we use developer tooling to drive operations.

# The GitOps proposal

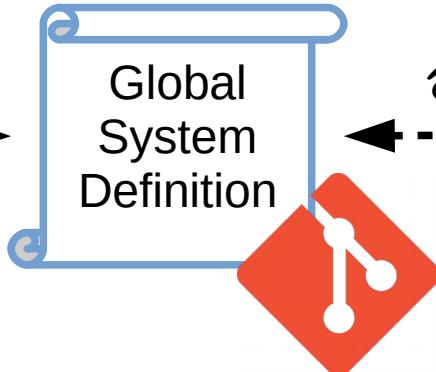


# The GitOps proposal



Corporate  
workflow  
automation

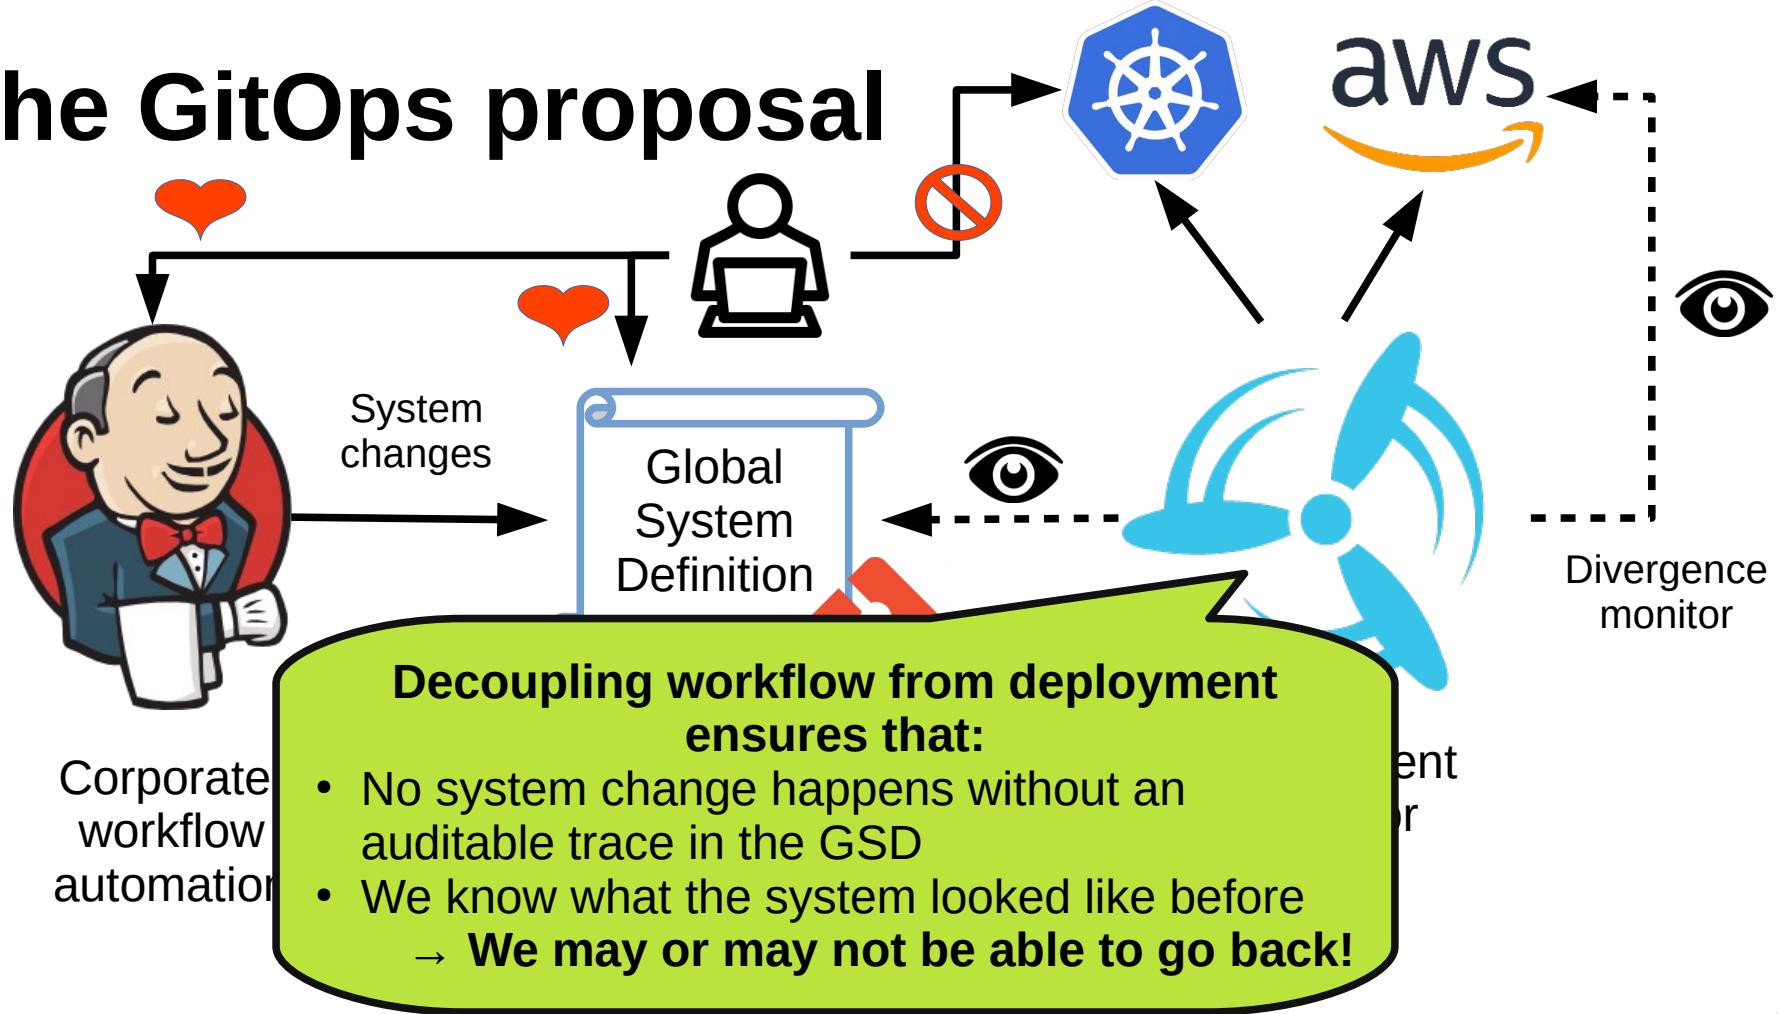
System  
changes



Deployment  
operator

Divergence  
monitor

# The GitOps proposal

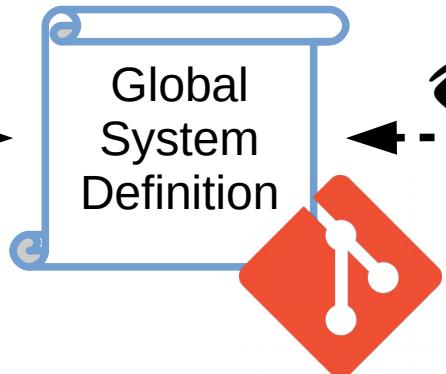


# The GitOps proposal



Corporate workflow automation

System changes



Divergence monitor

Deployment  
operator

Alerts about system divergence.  
Divergence means that right now we  
wouldn't be able to recreate the system.

Corporate workflow automation

- Workflow Engine & Sys log
- Keeps the **GSD** change log
- Keeping the **current state** of all the running workflows
- Provides **transactional support** for the workflows
- Multi purpose **state machine**
- Can easily **plug-in** new tools into the workflows

# proposal

Global System Definition



## Not just Git!

- Can be just Git...
- But for complex workflows it's much more convenient to have MR/PR, approvals...
- So use GitHub, GitLab or equivalent!



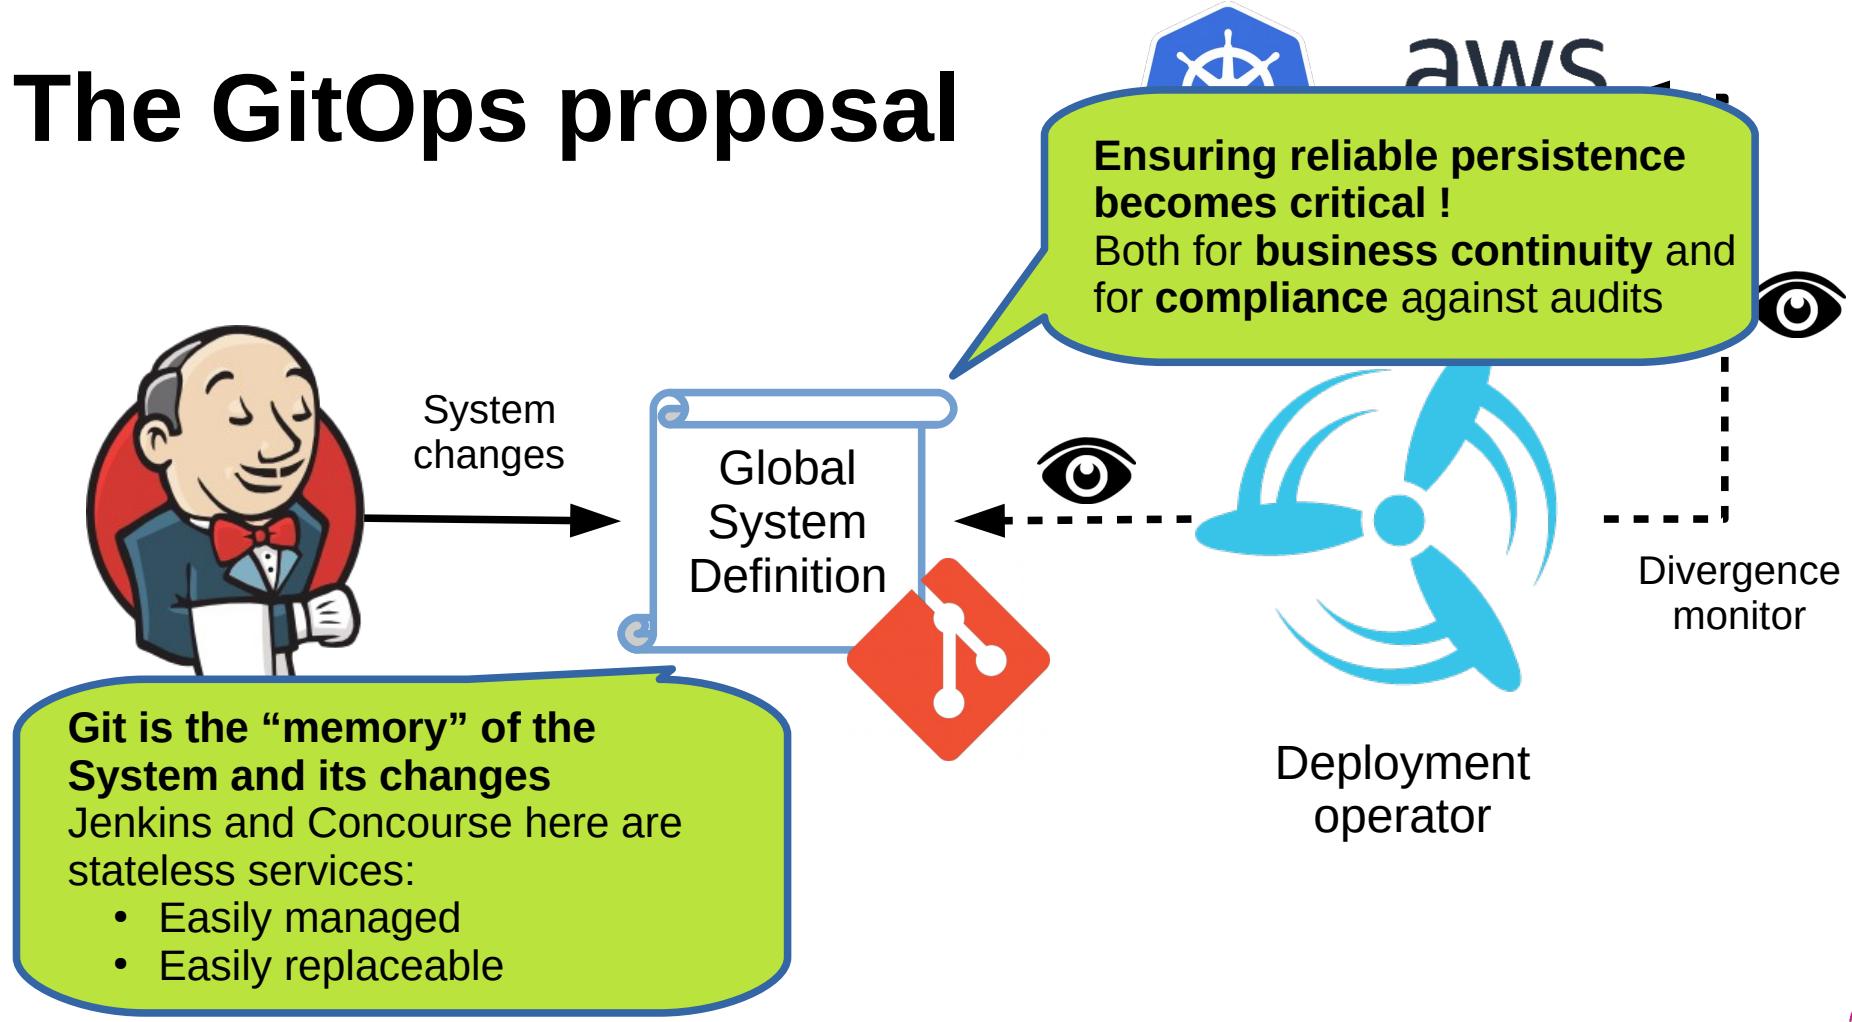
aws

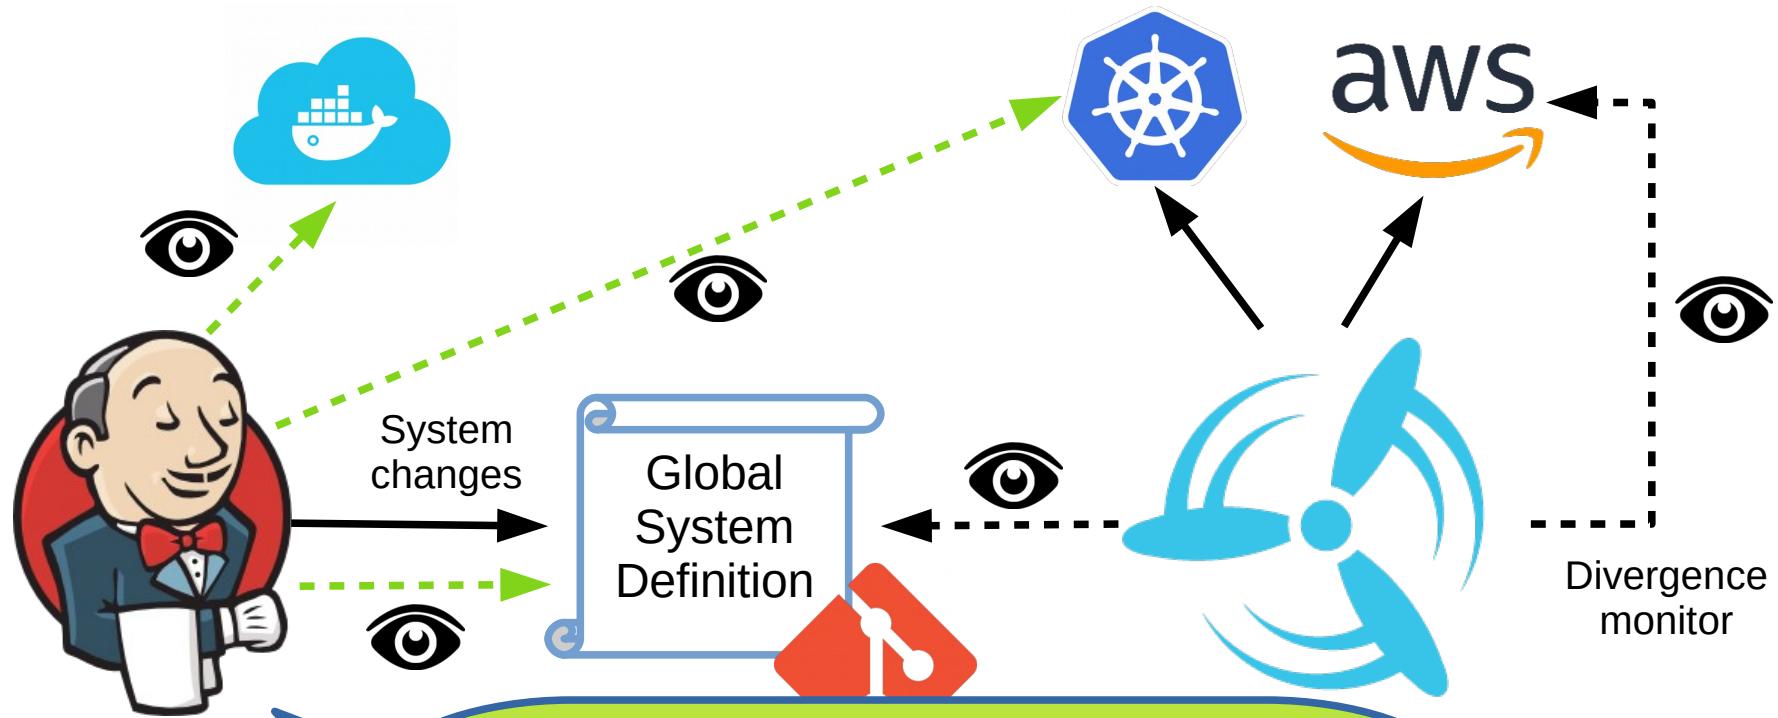


see



# The GitOps proposal



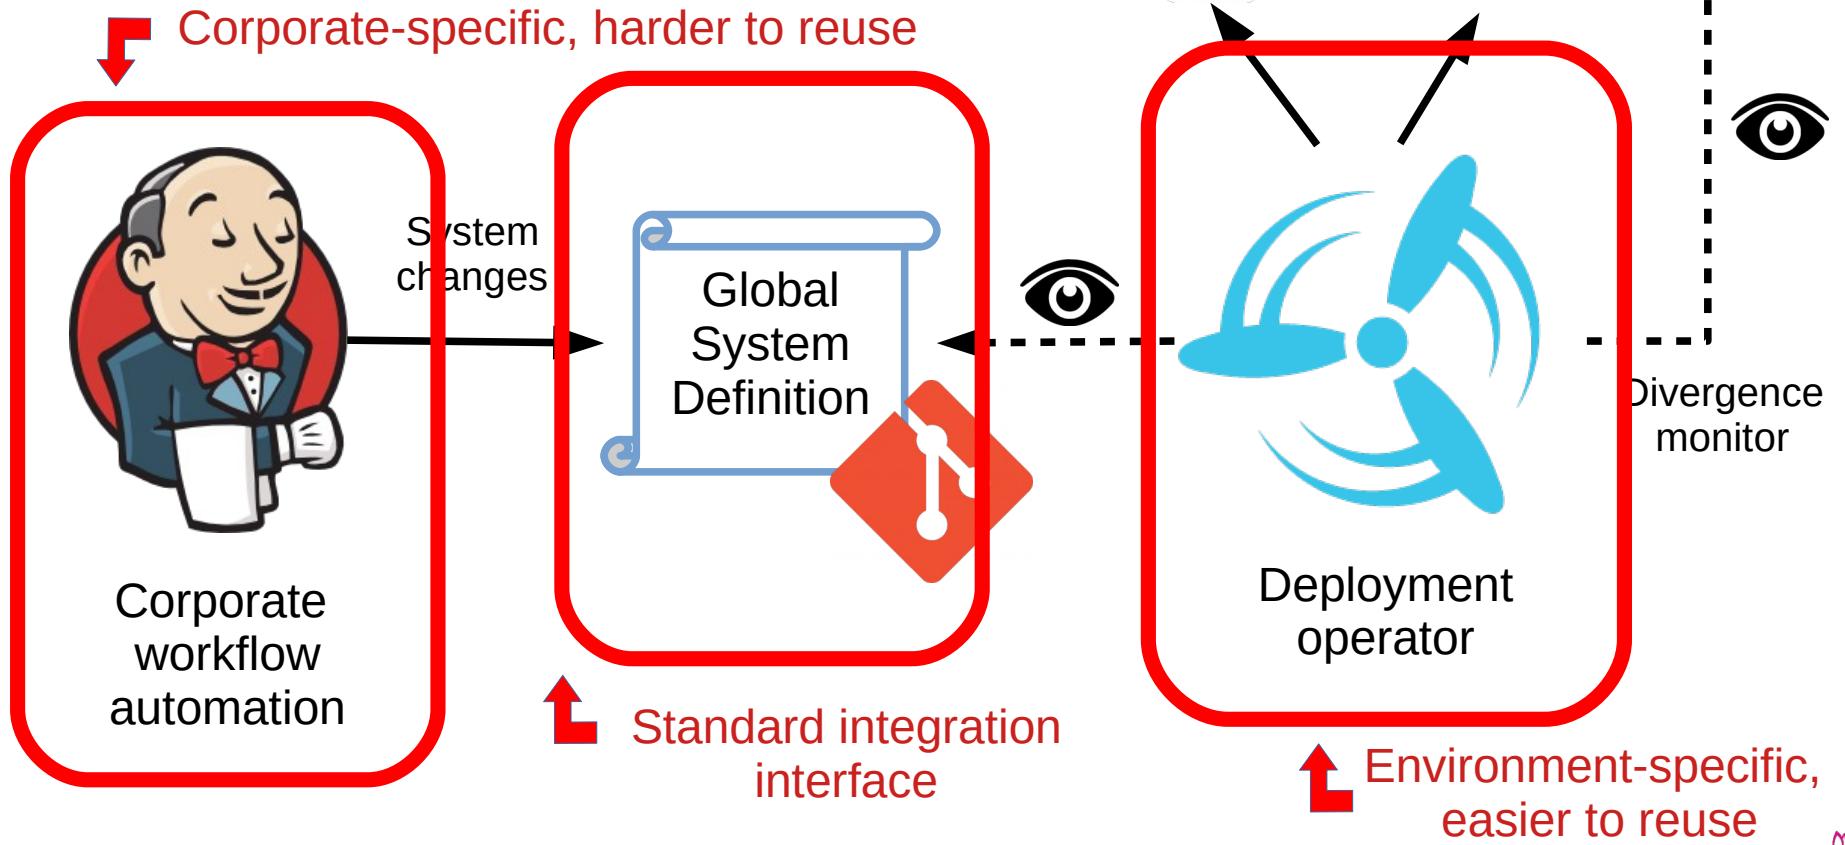


Corporate workflow automation

### Workflow enabler

- Watches external events
- Triggers workflows
  - Auto merges, MR + required approvals...
- Handles workflow timeouts

# The GitOps proposal

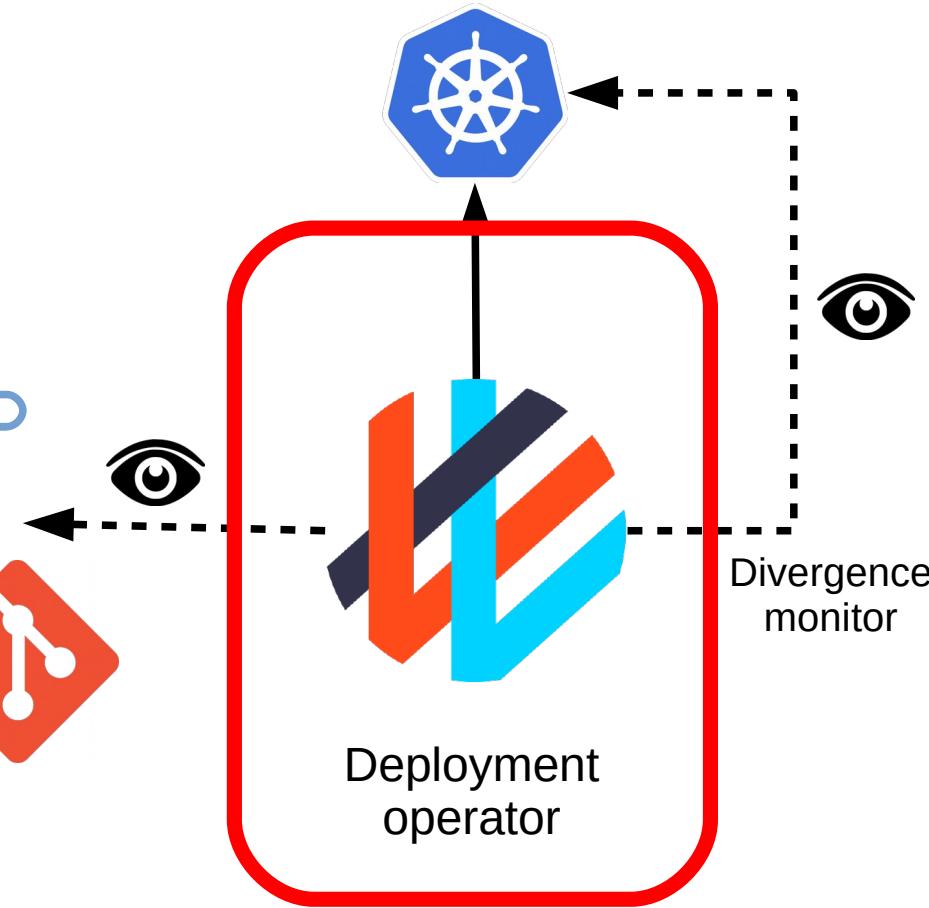
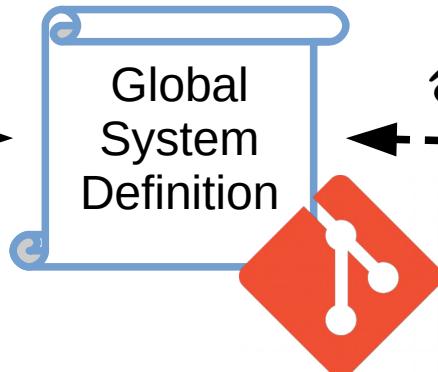


# The Flux proposal



Corporate  
workflow  
automation

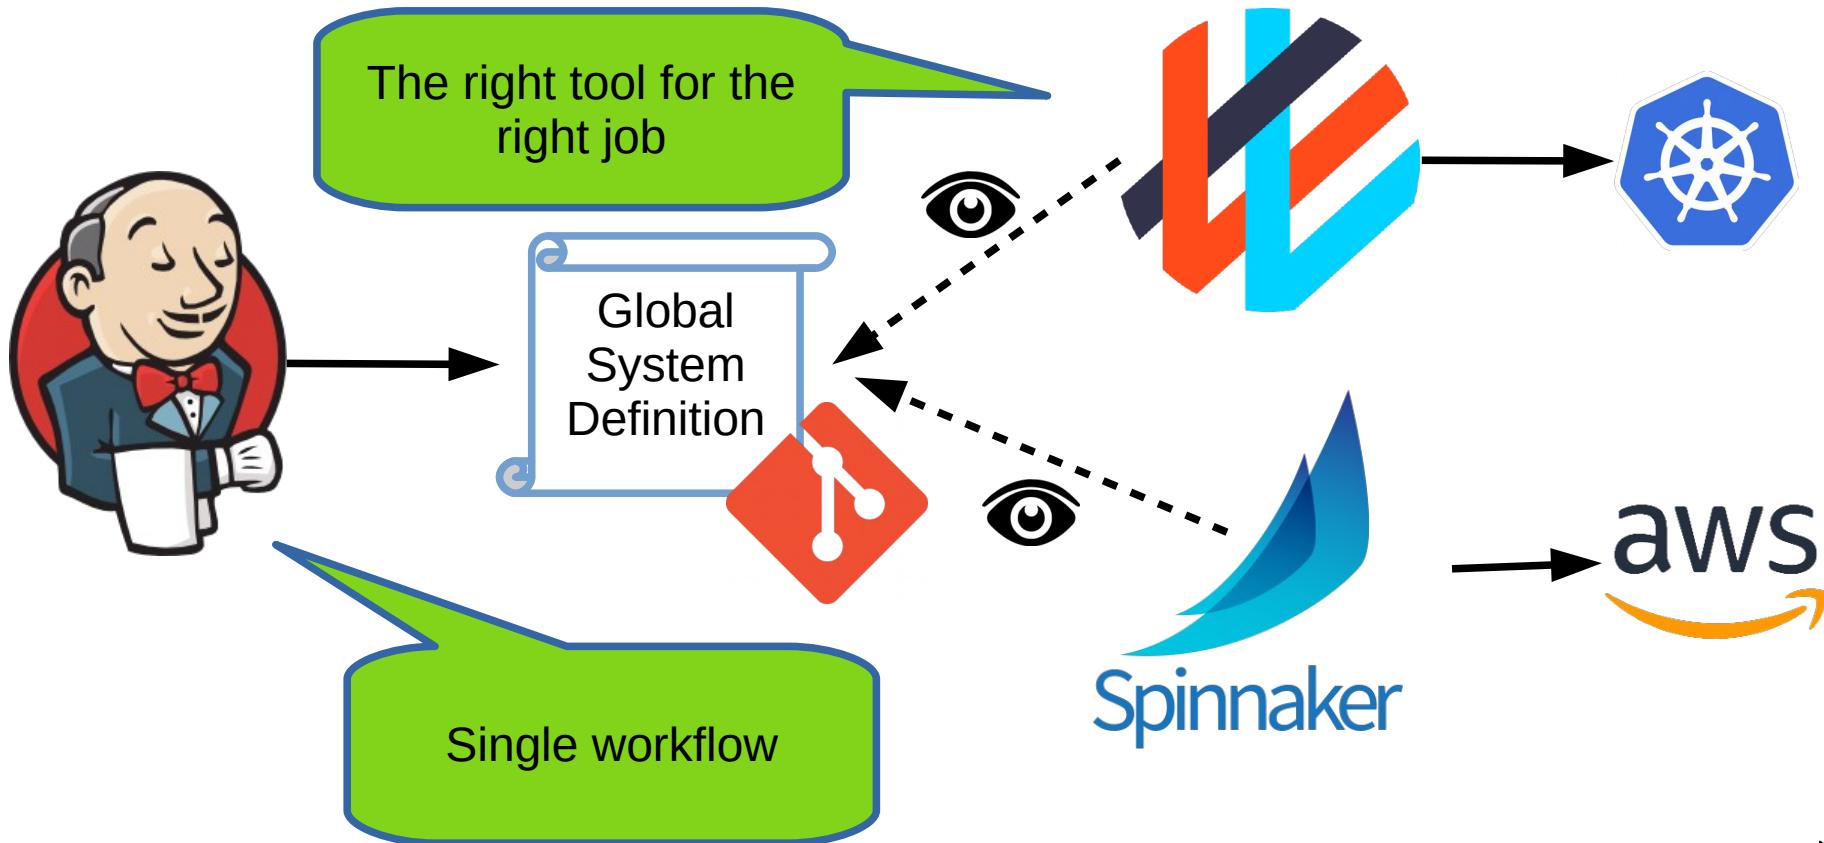
System  
changes



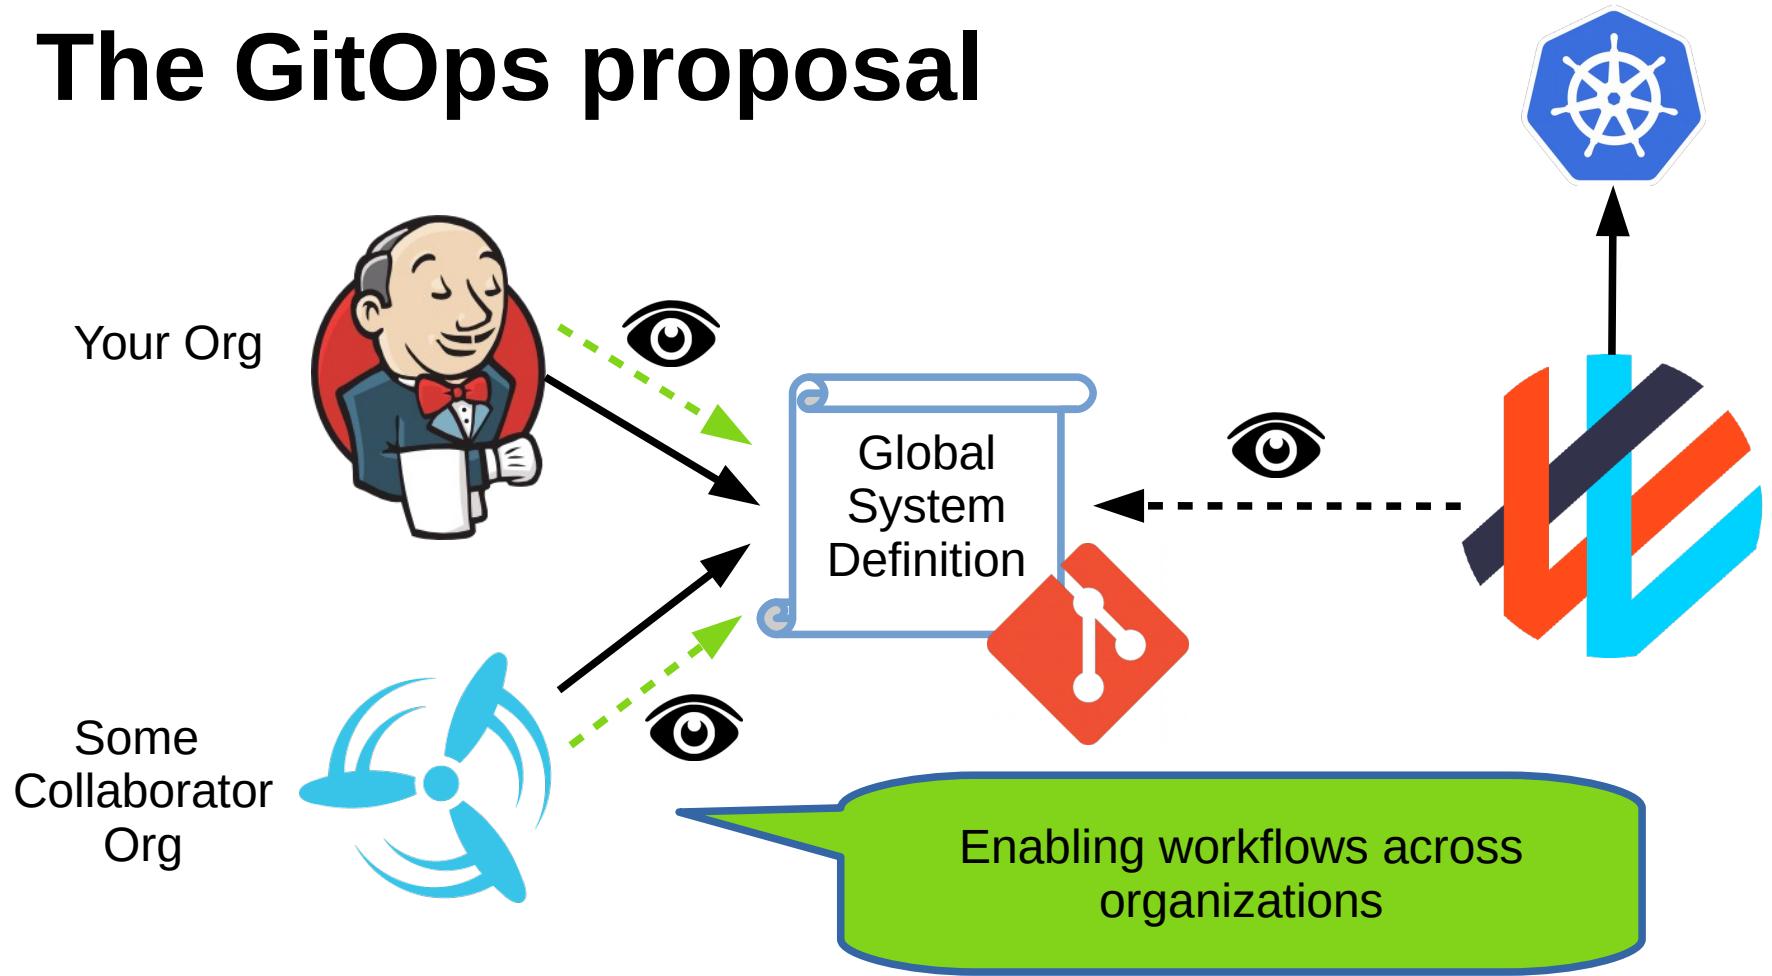
Kubernetes-specific, easier to reuse.  
Kubernetes-specific concerns.



# The GitOps proposal

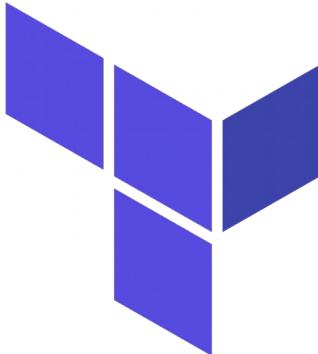


# The GitOps proposal



# GitOps practices

- **Global System Definition in Git**
  - Expressed in declarative form using any appropriate tool

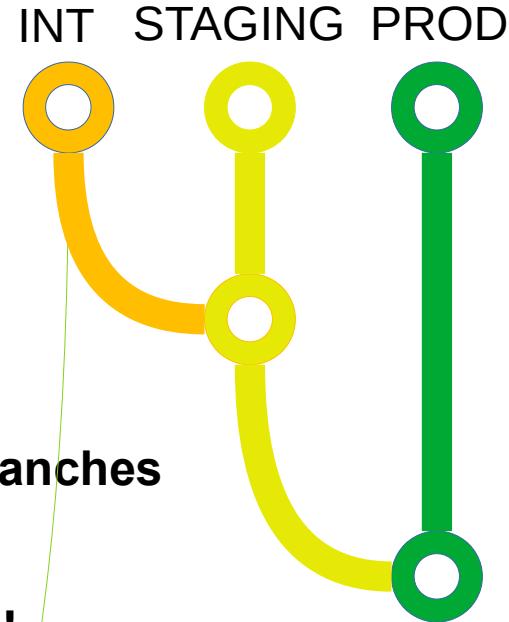


# GitOps practices

- **Git as the system's Single Source of Truth**
  - All the config in Git:
    - **Application config**
    - **Infrastructure config**
      - K8s templates...
    - **Observability config**
      - Grafana dashboards definitions...
    - **Controllability config**
      - Automated operation playbooks...
        - Extra value if integrated with Obs tools (actionable dashboards)

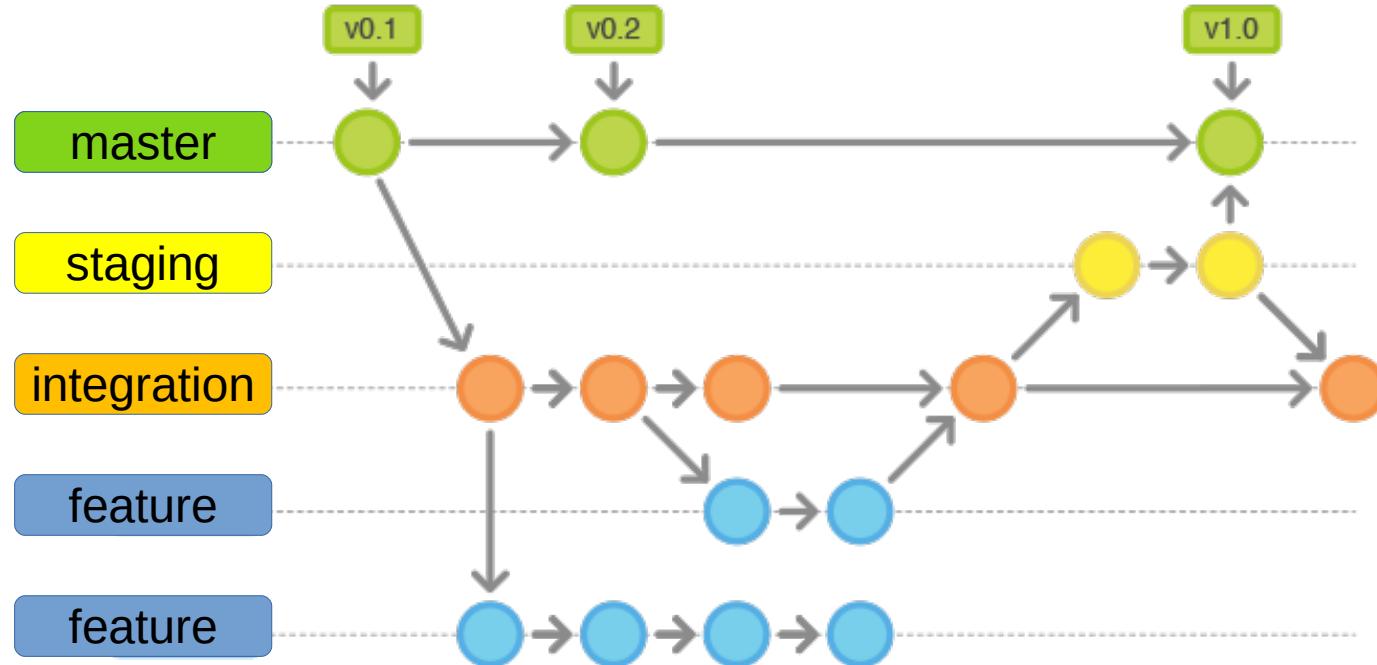
# GitOps practices

- **Branch by Environment**
  - Production env config in **master** branch
  - Other branches for previous envs
  - **Promotion between envs = PR + Merging env branches**
    - e.g. Merging INTEGRATION → STAGING
    - **No direct commits allowed on env branches!**
    - **Brings Git workflow capabilities to system evolution**
      - Prevents “race conditions” during deployments
      - Forces teams to integrate system upgrades
        - Prevents accidental regressions



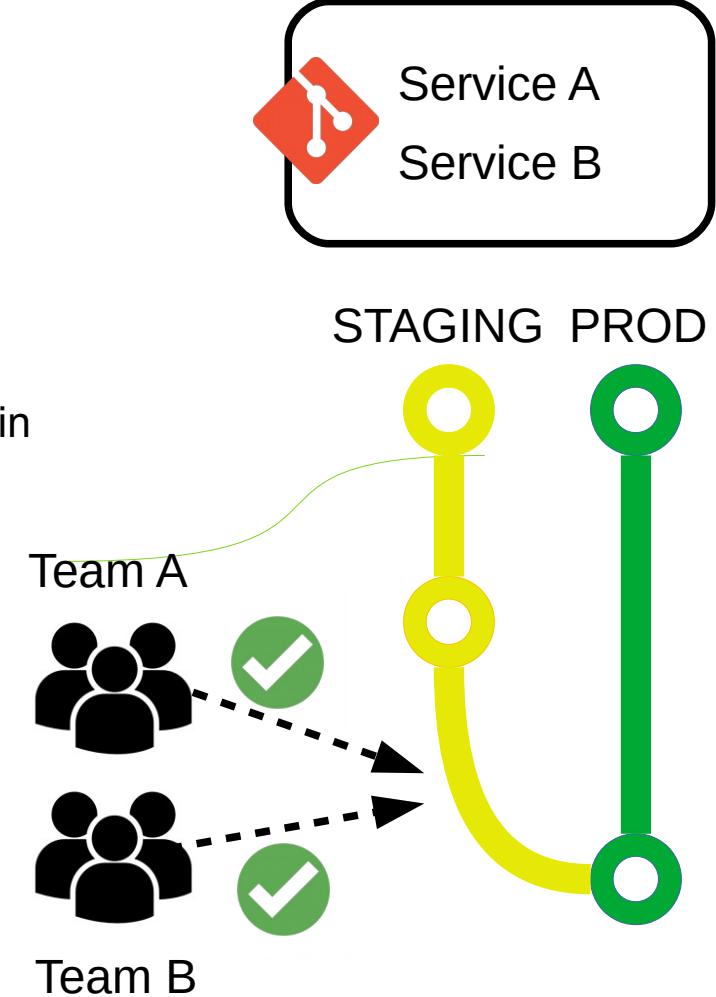
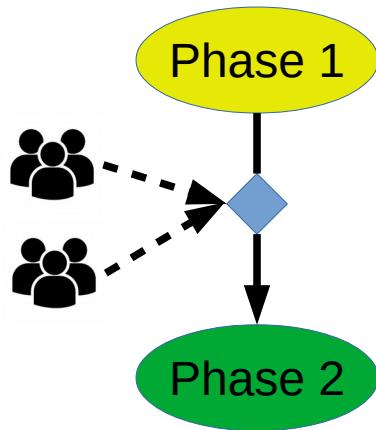
# GitOps practices

Git Flow applied to system evolution



# GitOps practices

- **Branch by Environment**
  - Workflow modeling using Git
    - **Branches**
      - Can be used to model **stages** in the workflow
    - **Merge Requests**
      - Can be used to model **transitions** between stages
    - **MR approvals**
      - Can be used to model **conditions or inputs** for transitions



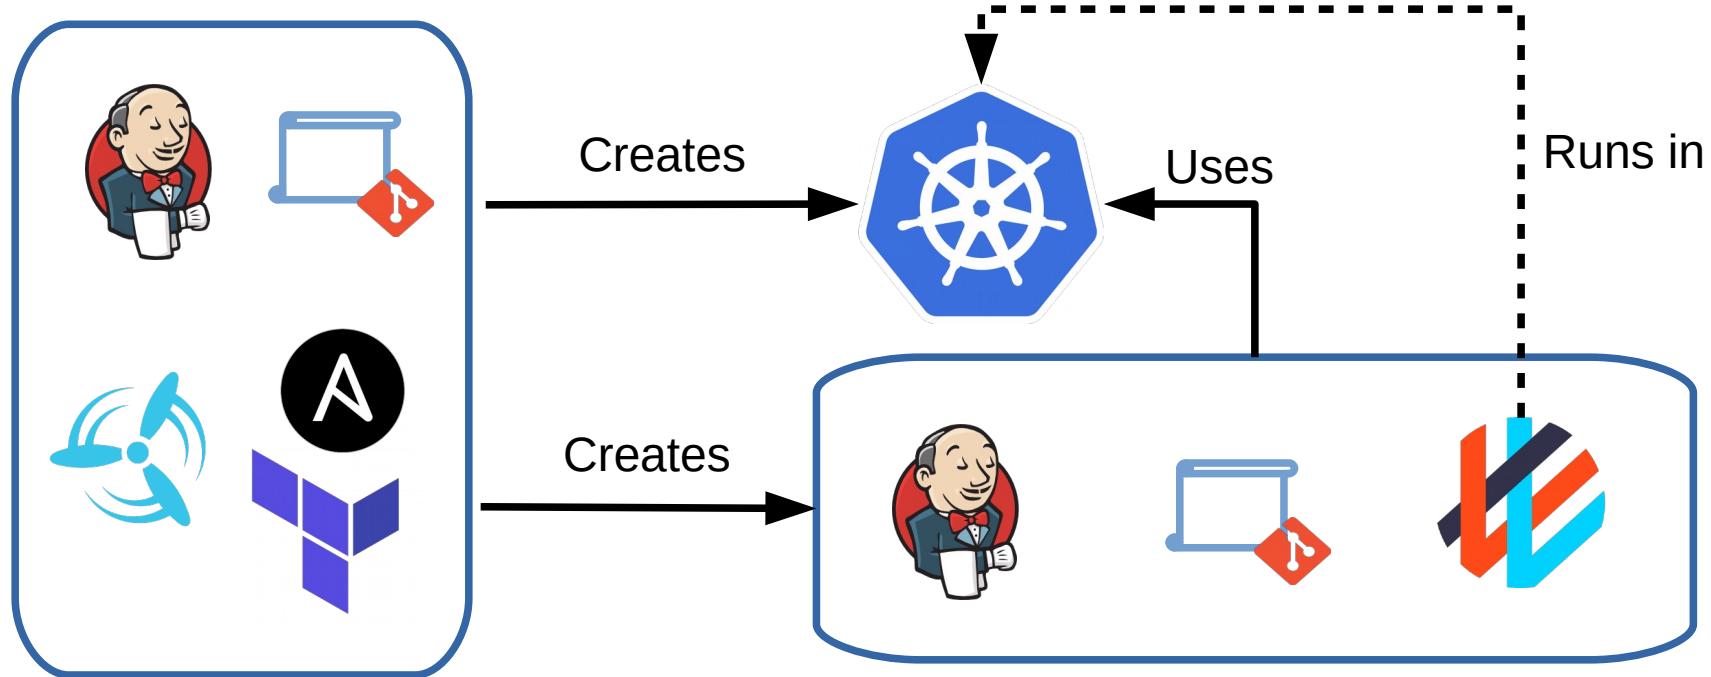
# GitOps practices

- **Divergence control**
  - Git == Single Source of Truth
    - **Divergence == Danger**
      - We can't reproduce the system anymore
  - Use **agents** in the target environments to watch for divergence
    - **Alerting**
    - **Automated convergence (if applicable/possible)**
      - i.e. Rolling back to last accepted GSD state

# GitOps practices

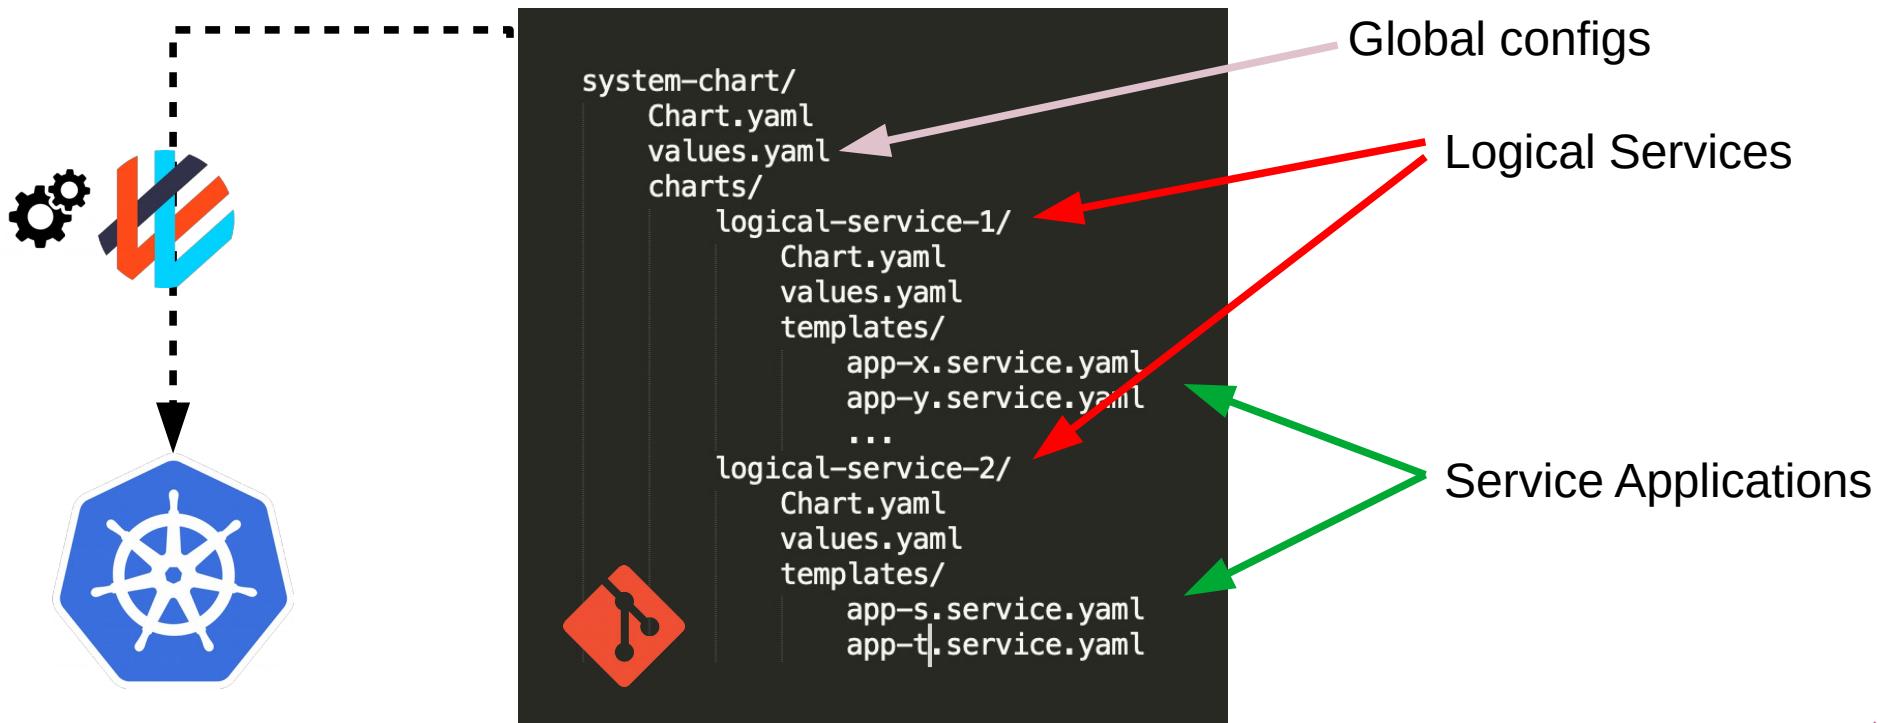
- **GitOps 4 GitOps**

- Your environments & GitOps platform also needs consistency!



# GitOps 4 Kube

## GitOps example code structure: Deep Tree System Chart



# GitOps 4 Kube

## GitOps example code structure: Deep Tree System Chart

```
system-chart/
  Chart.yaml
  values.yaml
  charts/
    logical-service-1/
      Chart.yaml
      values.yaml
      templates/
        app-x.service.yaml
        app-y.service.yaml
        ...
    logical-service-2/
      Chart.yaml
      values.yaml
      templates/
        app-s.service.yaml
        app-t.service.yaml
```



- **Pro:**

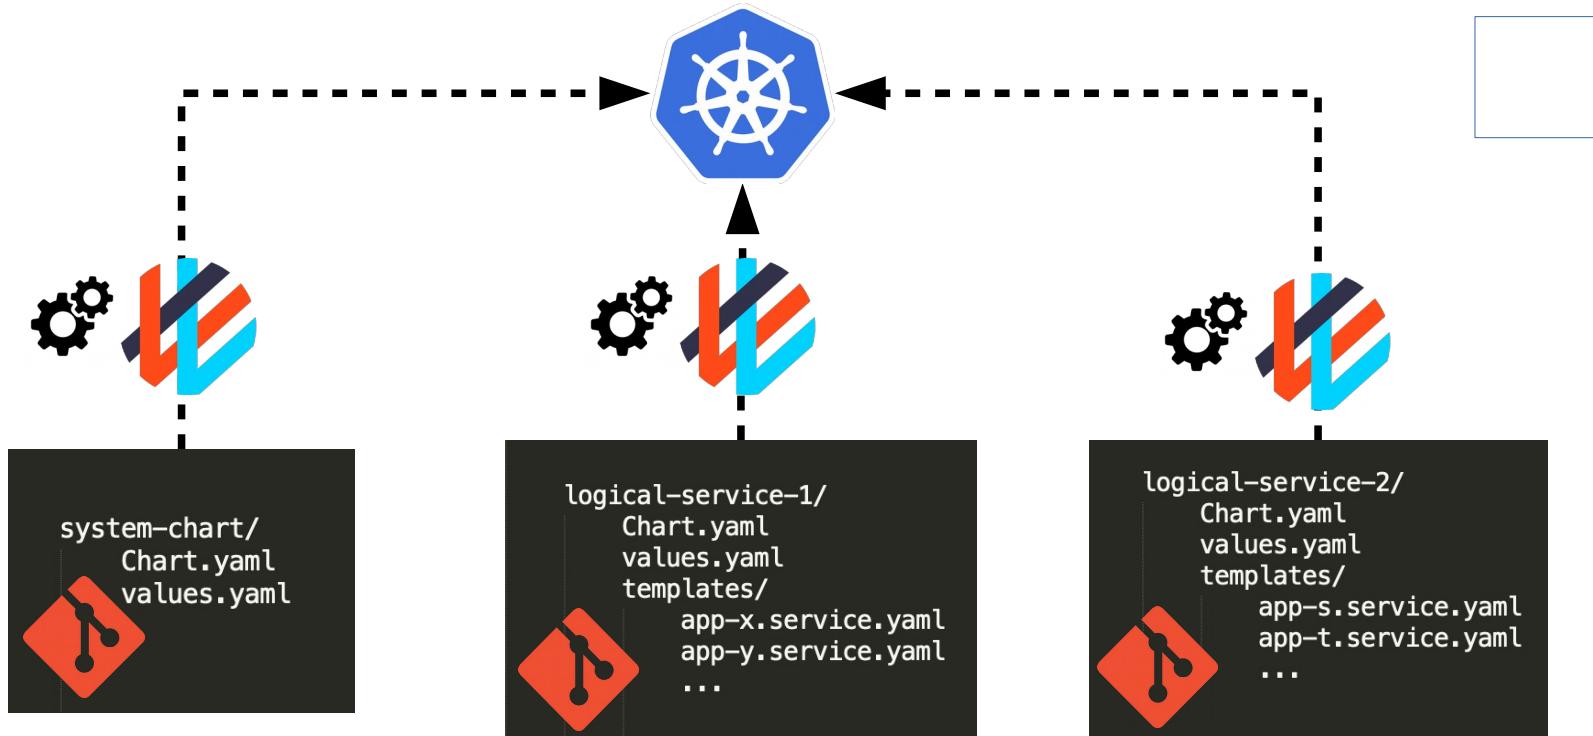
- Single Git repo
- Single GitOps pipeline

- **Con:**

- Can get cumbersome if Merge Requests require several manual approvals to merge into Production branch
  - e.g. LS-1 deploying changes all the time, and every MR requiring also the approval of LS-2
  - e.g. Doc changes in LS-1 require MR

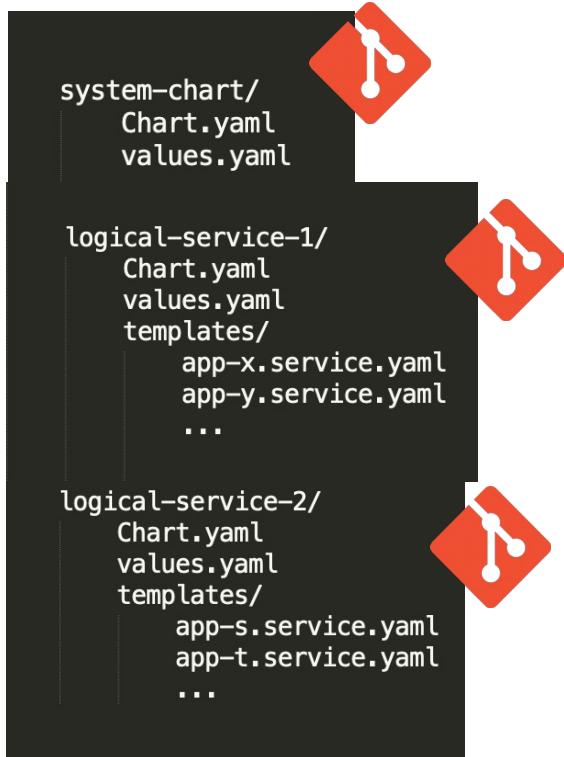
# GitOps 4 Kube

GitOps example code structure: Pipeline per Logical Service



# GitOps 4 Kube

## GitOps example code structure: Pipeline per Logical Service



- **Pro:**

- Separate Git repos for every chart
  - e.g. doc changes in LS chart don't require MR in SystemChart
- Simpler MR approval in each Git repo
  - As the Git repo belongs to the team owning each Logical Service, approving the MR only requires their own approval

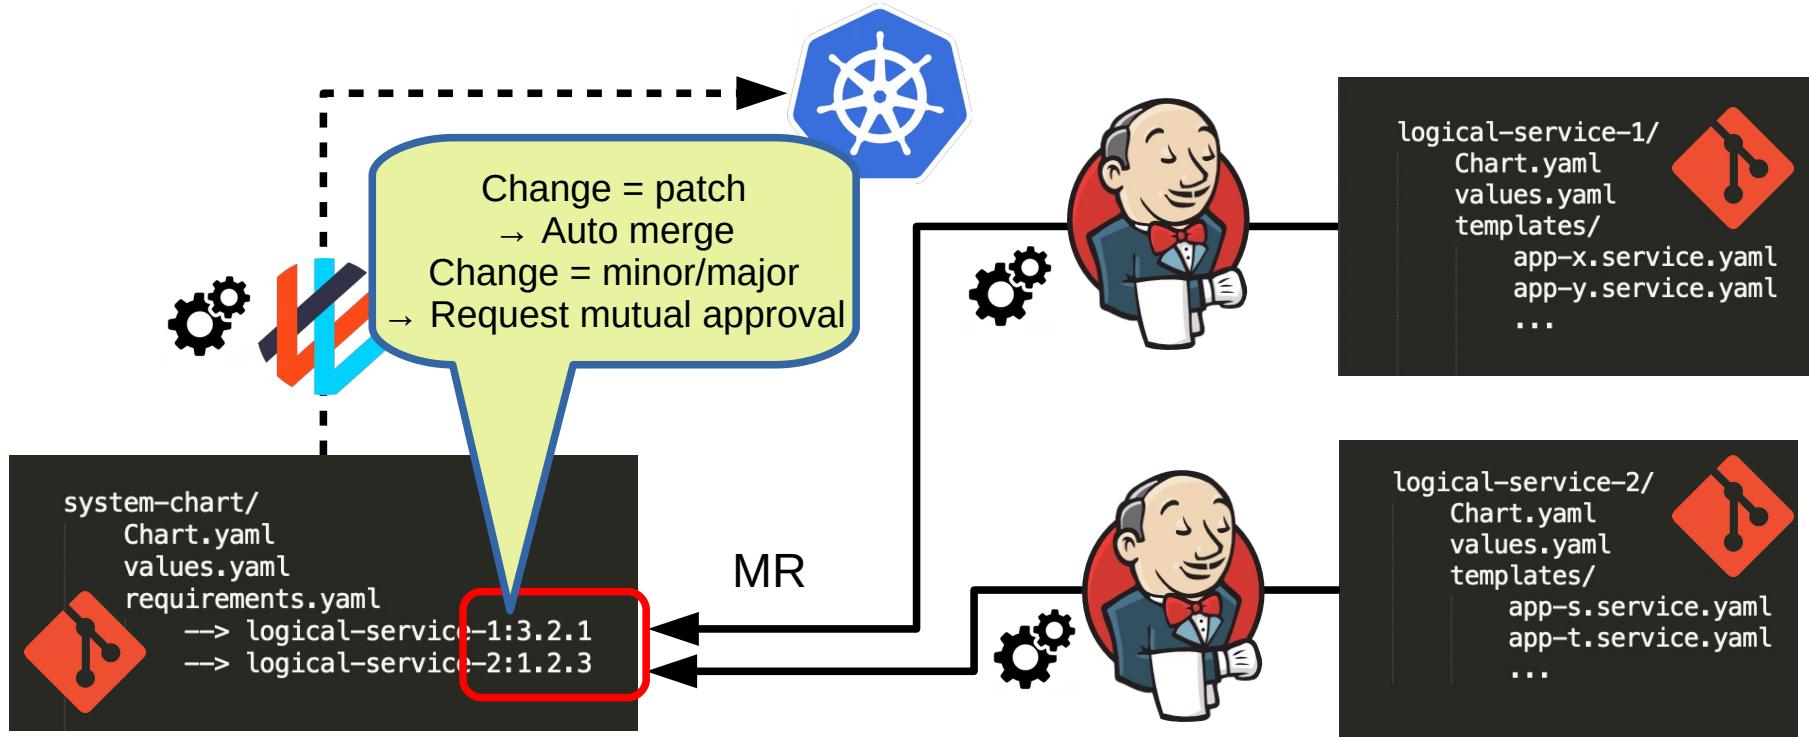
- **Con:**

- If the LS depend on each other, incompatible versions of both LS could end up deployed in Production



# GitOps 4 Kube

## GitOps example code structure: Chained pipelines



# GitOps 4 Kube

## GitOps example code structure: Chained pipelines

```
logical-service-1/
Chart.yaml
values.yaml
templates/
  app-x.service.yaml
  app-y.service.yaml
  ...
  
```



```
logical-service-2/
Chart.yaml
values.yaml
templates/
  app-s.service.yaml
  app-t.service.yaml
  ...
  
```



```
system-chart/
Chart.yaml
values.yaml
requirements.yaml
  --> logical-service-1:3.2.1
  --> logical-service-2:1.2.3
  
```



- **Pro:**

- Separate Git repos for every chart
    - Doc changes don't require a MR in SystemChart

- **Con:**

- More complex setup, CI pipelines must create a MR to SystemChart to bump up the versions of the LS charts
    - Alternatively, when bumping just patch versions the MR could be automatically merged



# GitOps benefits

<https://www.weave.works/technologies/gitops/>

## 1) Increased Productivity

- Faster deployments



CD benefit

## 2) Enhanced Developer Experience

## 3) Improved Stability

## 4) Higher Reliability

## 5) Consistency and Standardization

## 6) Stronger Security Guarantees

# GitOps benefits

<https://www.weave.works/technologies/gitops/>

- 1) Increased Productivity**
- 2) Enhanced Developer Experience**
  - Devs & Ops only use Git  CD benefit
- 3) Improved Stability**
- 4) Higher Reliability**
- 5) Consistency and Standardization**
- 6) Stronger Security Guarantees**

# GitOps benefits

<https://www.weave.works/technologies/gitops/>

- 1) Increased Productivity**
- 2) Enhanced Developer Experience**
- 3) Improved Stability** 
  - Full system change log in Git → Easier troubleshooting  
(both for deployments & operations)
- 4) Higher Reliability**
- 5) Consistency and Standardization**
- 6) Stronger Security Guarantees**

# GitOps benefits

<https://www.weave.works/technologies/gitops/benefits>

- 1) Increased Productivity
- 2) Enhanced Developer Experience

- 3) Improved Stability



- 4) Higher Reliability



– Full system change history in Git

→ Can recreate system. **Can rollback accurately.**

- 5) Consistency and Standardization

- 6) Stronger Security Guarantees



© Nestlé S.A.



# GitOps benefits

<https://www.weave.works/technologies/gitops/>

## 1) Increased Productivity

## 2) Enhanced Developer Experience

## 3) Improved Stability



Creating the right workflow for your organization doesn't have to be easy, though

## 4) Higher Reliability



## 5) Consistency and Standardization



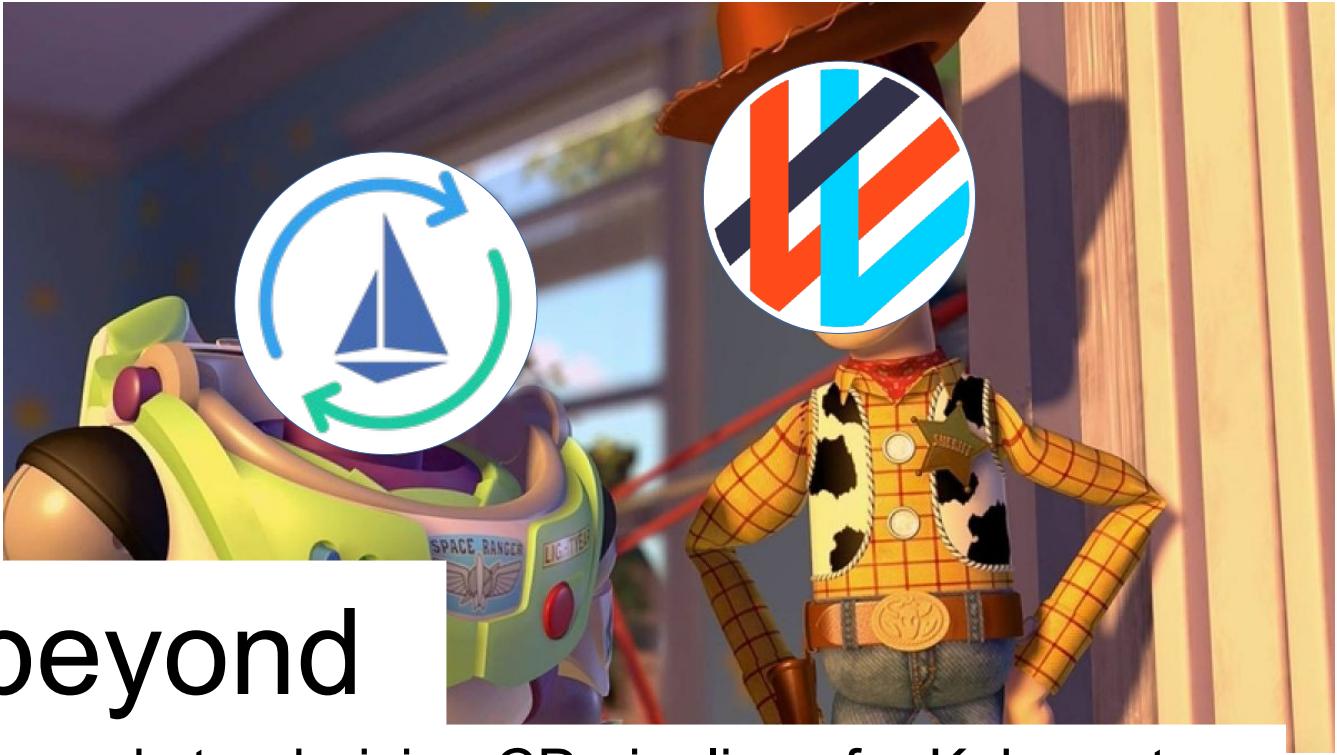
- Consistent workflow for system changes by using Git
- Git as the standard workflow tool, easy to hook-in external tools

## 6) Stronger Security Guarantees

# GitOps benefits

<https://www.weave.works/technologies/gitops/>

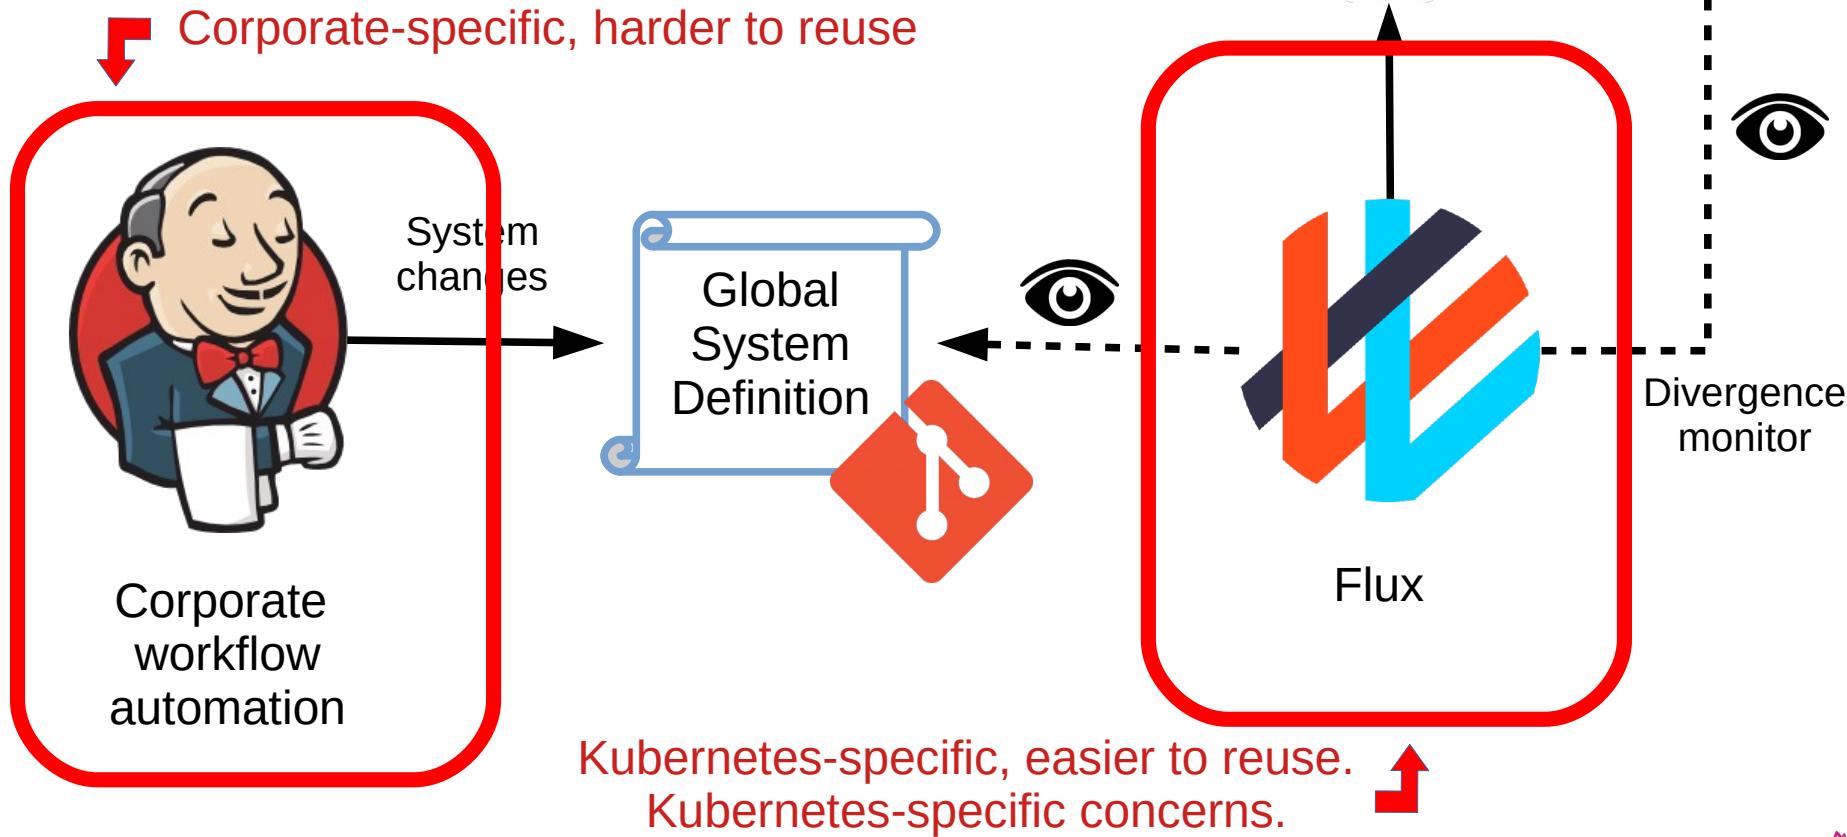
- 1) Increased Productivity**
- 2) Enhanced Developer Experience**
- 3) Improved Stability** 
- 4) Higher Reliability** 
- 5) Consistency and Standardization** 
- 6) Stronger Security Guarantees** 
  - Sys changes only through Git (requires auth & authoriz)
  - Auditability of system changes



# Flux & beyond

Automating and standardizing CD pipelines for Kubernetes  
with the right mix of centralization & self-service

# The Flux proposal



# Weaveworks Flux

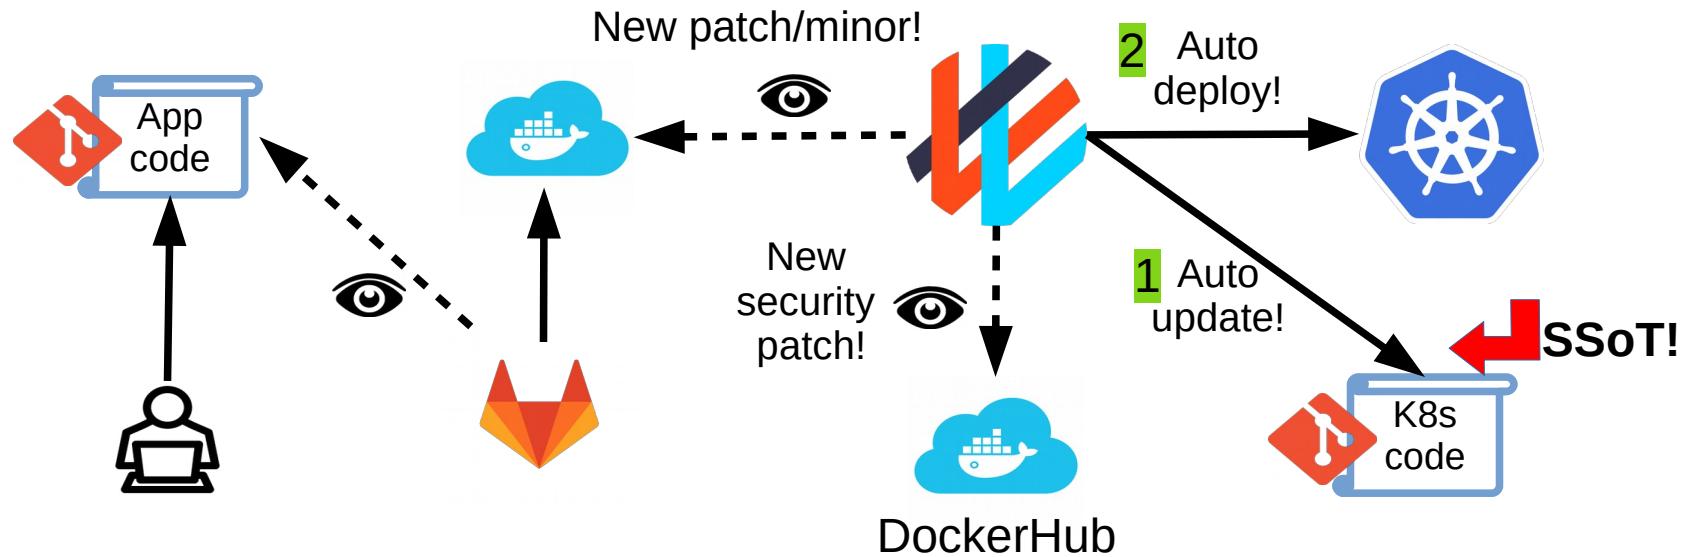


- **Implements a GitOps CD pipeline for Kube & Helm**
- **Flux = Helm Operator + Deployment Operator**
- **Deployment Operator**
  - Watches 1 Git repo+branch for changes
    - Note: Maybe in the future Flux will support **1 instance → \* repo+env**
    - **Need 1 DO instance per repo & environment**
  - Can deploy:
    - Kube workloads: CronJob, DaemonSet, Deployment, StatefulSet
    - Helm workloads: HelmRelease (CRD)

# Weaveworks Flux



- Deployment Operator
  - Can auto scan container registries for updated images of installed workloads & redeploy them



# Weaveworks Flux



- Deployment Operator
  - Can auto scan container registries for updated images of installed workloads & redeploy them
    - USEFUL !!!
      - Reduces Kube routine operation efforts (less MRs)
      - Improves system security → Auto installs latest patches
      - Simplifies workflows by skipping MR approvals in some cases
        - e.g. auto install app patch & minor upgrades, MR for majors
      - Can be fine tuned
        - e.g. auto install only images tagged as \* - release
    - Git SSoT: After auto deploy, the workload is modified & committed
    - Drawback: We loose some control about system changes (but tolerable)

# DEMO #5-1

Automated Kube releases with Flux

# DEMO #5-2

GitOps workflows with Flux

# DEMO #5-3

Controlling workloads with Flux

# A Secret problem



# What about Secrets?

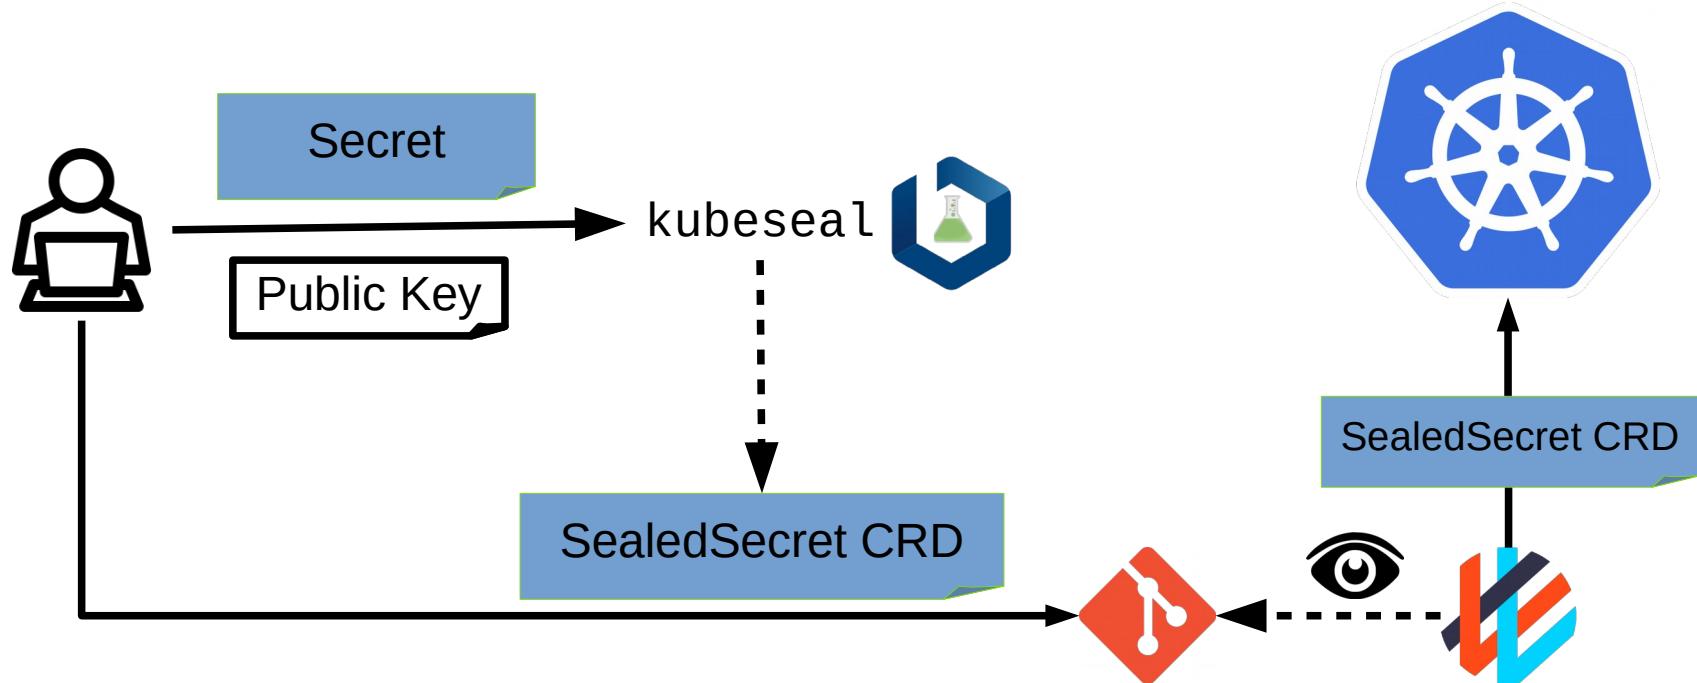
- **To keep the GitOps benefits, Secrets need also to be stored in Git**
  - They can be stored externally (e.g. using Vault)
    - But then we loose the “reliable rollback” capability
      - e.g. Upgrade workload, change secret in Vault, rollback  
→ Cannot recover old secret
        - New secret may lack some key/values of the old
      - Or we need to complicate the workflow to synch with Vault
  - **If we want to keep them in Git, however, they can't be in plain text**

# What about Secrets?

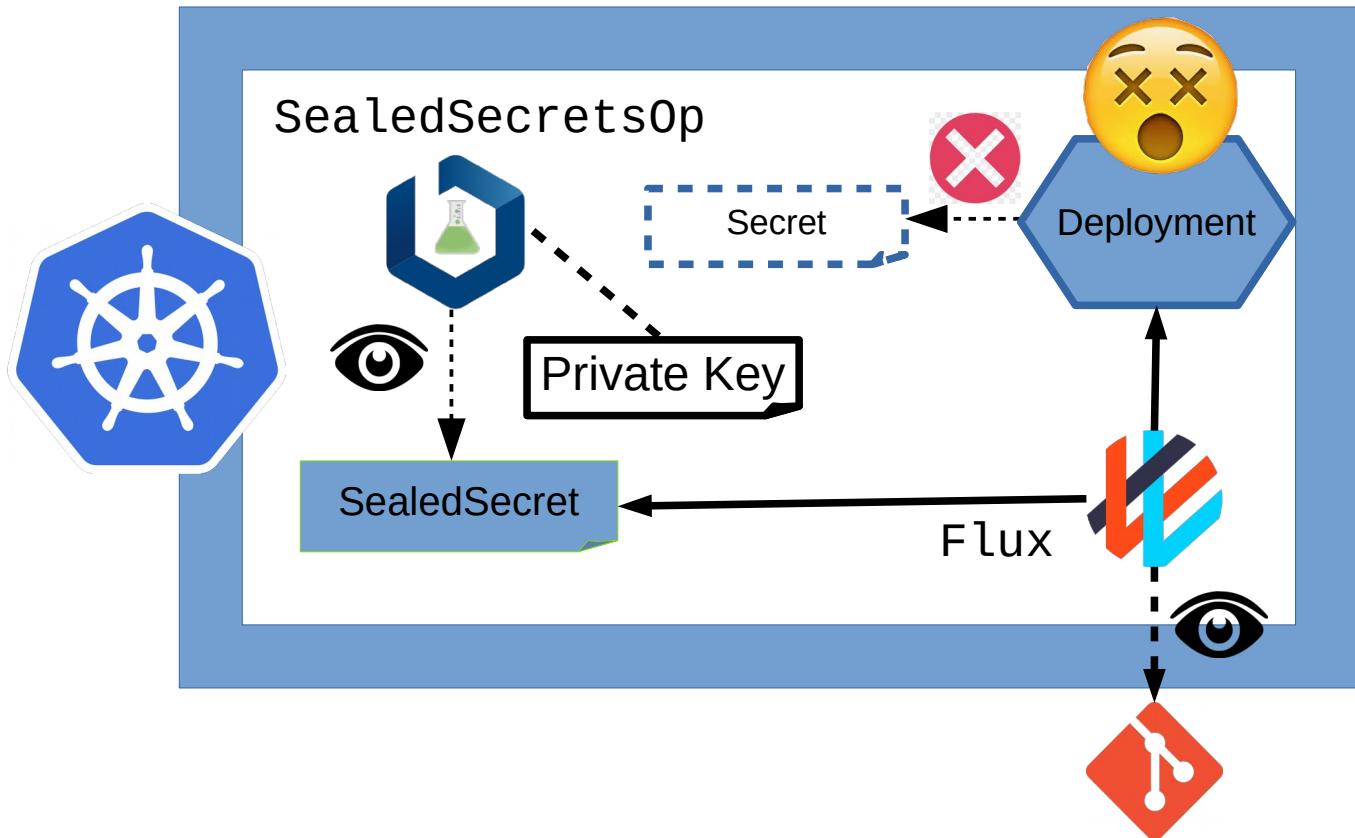
- **Option #1 – Bitnami's SealedSecrets Operator**
  - See: <https://github.com/bitnami-labs/sealed-secrets/>
  - An early solution by Bitnami to solve the Secret problem

# What about Secrets?

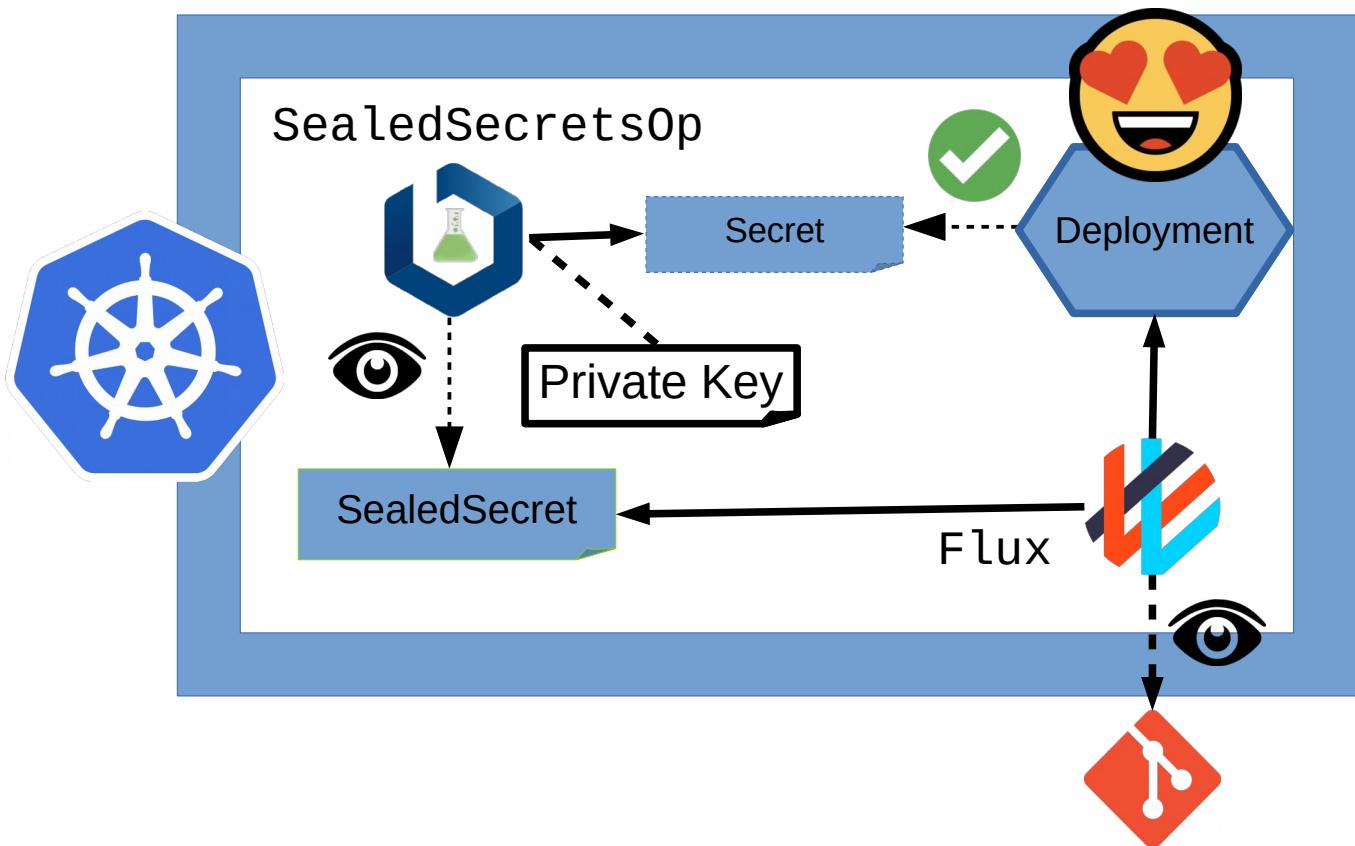
- Option #1 – Bitnami's SealedSecrets Operator



# Option #1 – Bitnami's SealedSecrets Operator



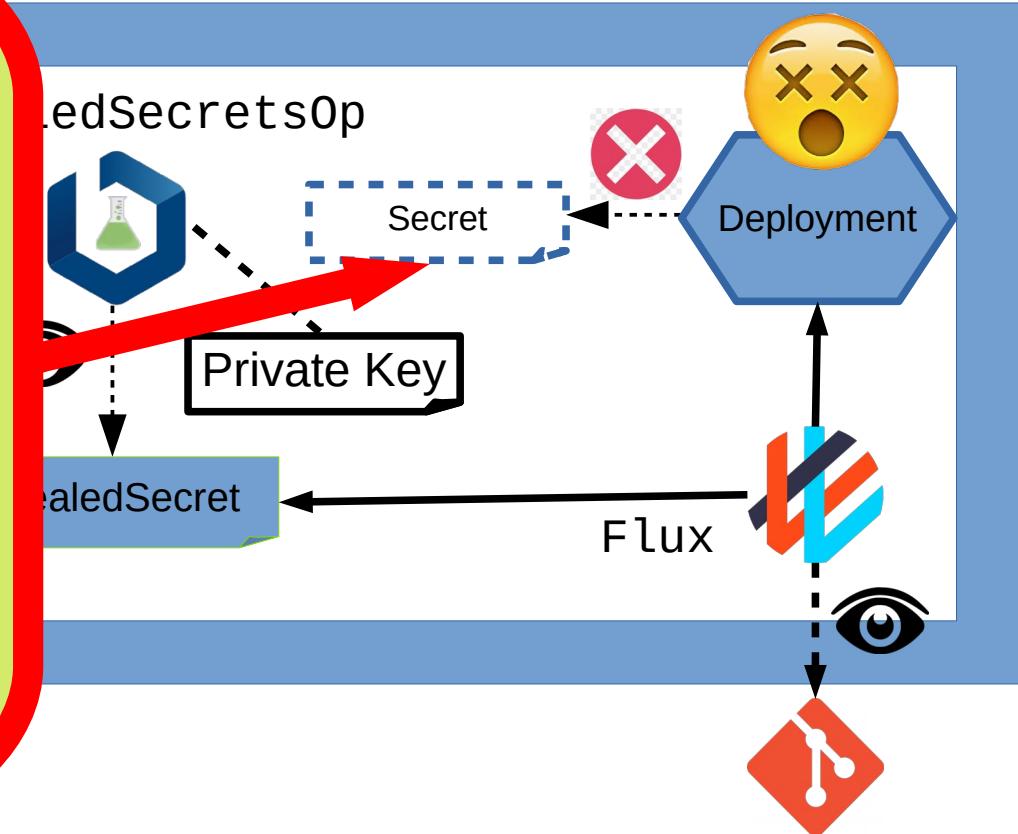
# Option #1 – Bitnami's SealedSecrets Operator



# Option #1 – Bitnami's SealedSecrets Operator

## Drawbacks

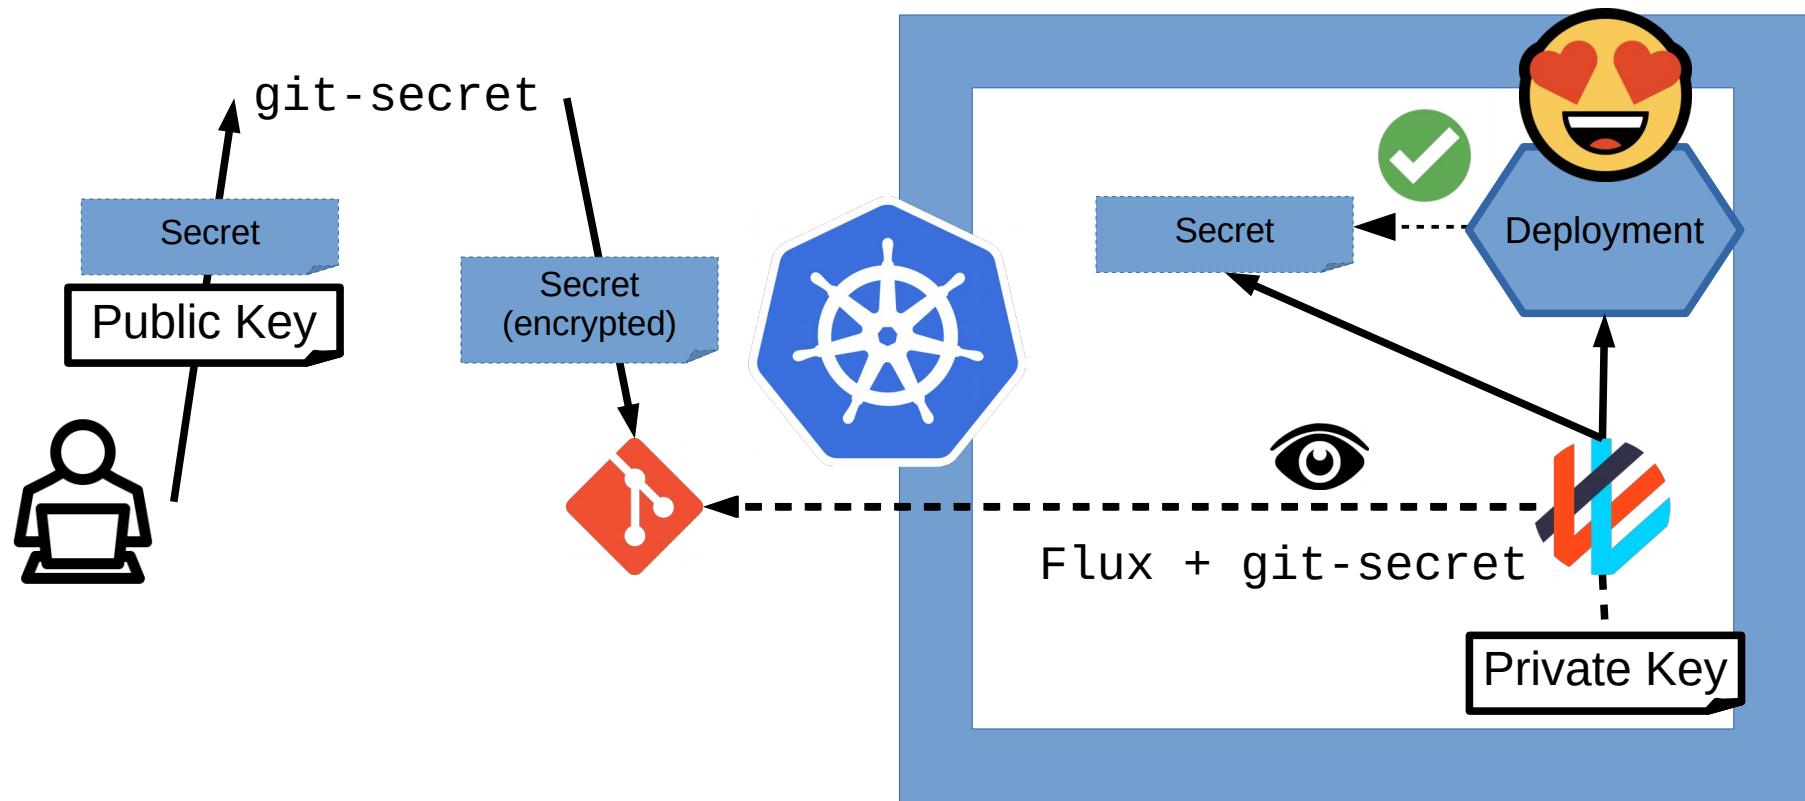
- SSOp: Another moving part of the system...
- SSOp & Flux disconnected from each other
  - Depending on config, Flux may rollback the deployment because of the missing Secret!
- SSOp: Another Operator...
  - Increases “limbo time”



# What about Secrets?

- **Option #2 – Flux + git-secret**
  - See:
    - <https://github.com/fluxcd/flux/pull/2159> ← **Very recent feature!**
    - <https://git-secret.io/> ← Transparent enc/dec for Git
  - Manifest files with Secrets stored in encrypted form in Git repo
  - Flux decrypts Secrets with git-secret before installing them in K8s

## Option #2 – Flux + git-secret



# DEMO #5-4

Flux + git-secret

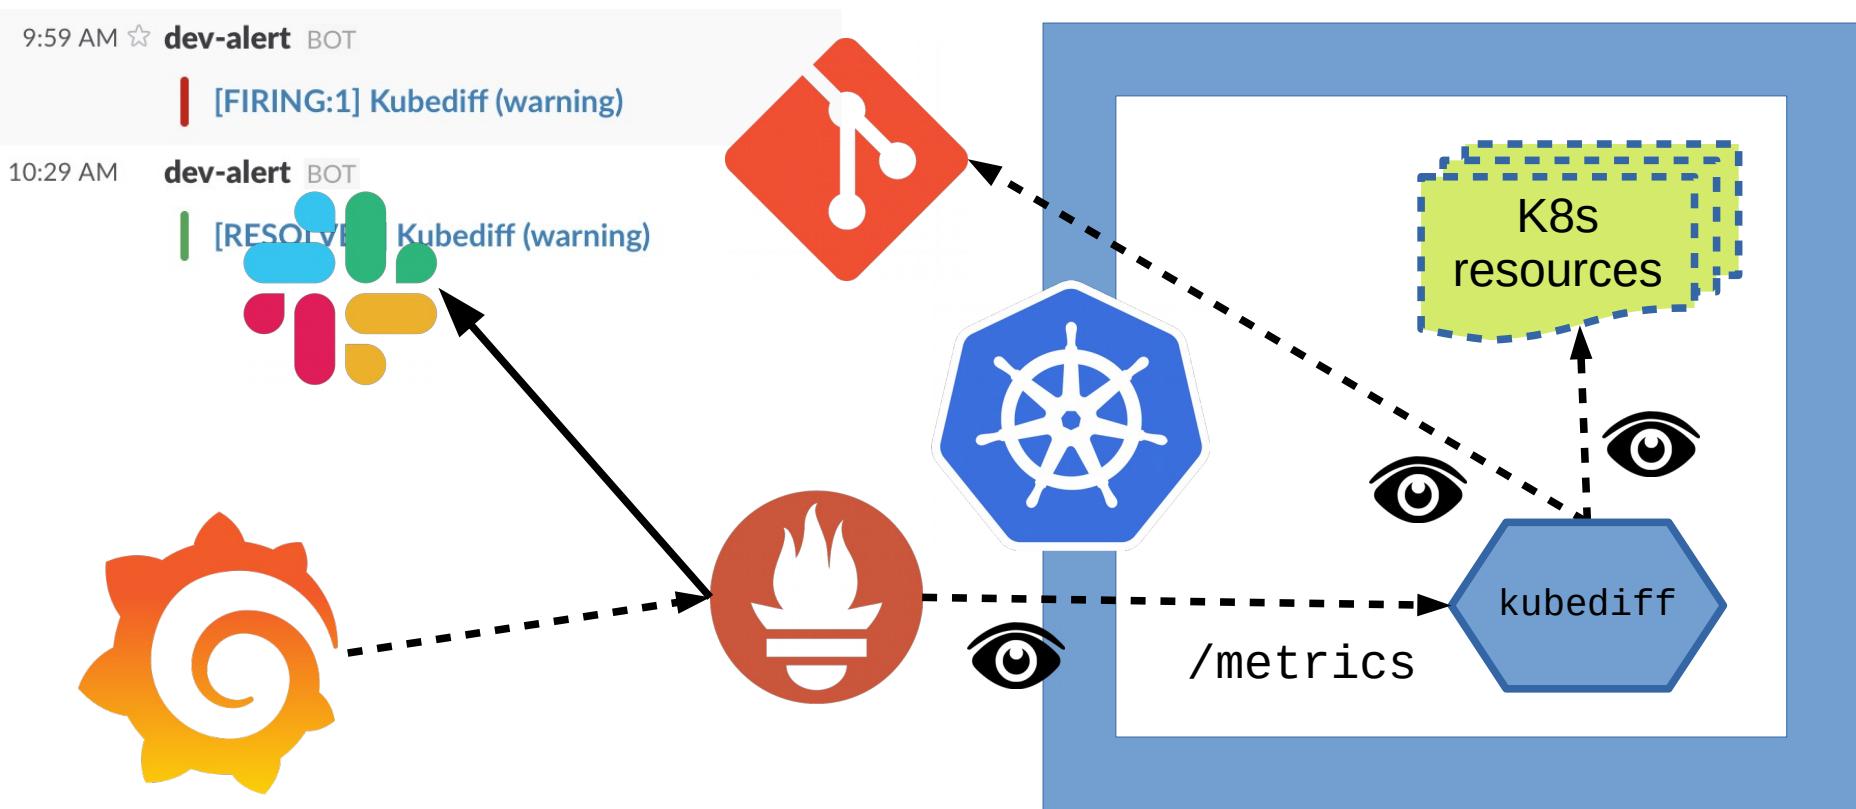
# A Secret “no problemo”



# What about System divergence?

- Weaveworks works currently with three main techs
- For each one they use a matching tool for system divergence control:
  - **K8s** → **kubediff**: <https://github.com/weaveworks/kubediff>
  - **Terraform** → **terradiff**
    - Internal tool, wrapper for standard TF CLI “plan” command
  - **Ansible** → **ansiblediff**
    - Internal tool, wrapper for standard Ansible “check mode”

# Kubediff



# Current Status

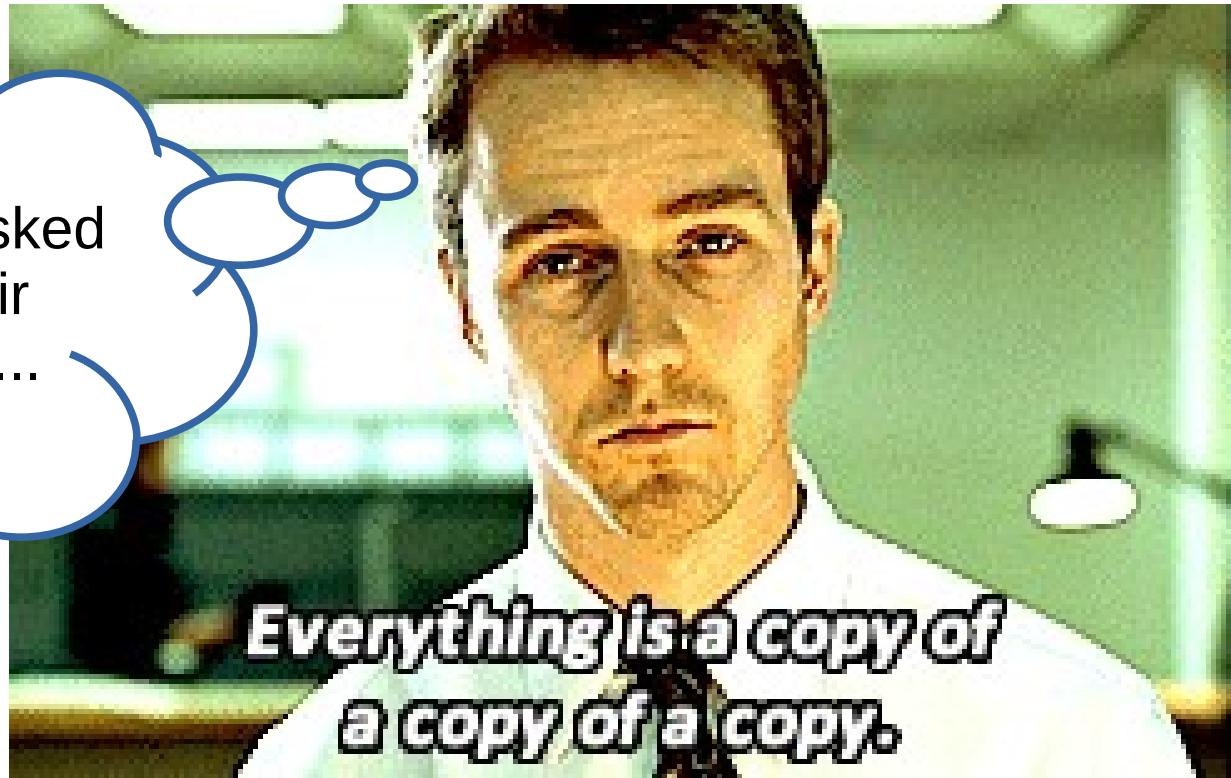
- **kubectl-less operations**
  - Offloading of Helm Release installation to CD (Flux)
- **Guaranteed reproducibility by using Git as SSoT**
  - Solved the problem of Secrets in Git
- **Separated corp workflow from env deployment (CD reuse)**



Can we do better?

# New first world problems...

Team Kolontái asked  
me to update their  
Flux configuration...  
**Again.**



Yes, we can!



# Flux friction points

- **Flux instance per Git repo + env (branch)**
  - Note: Future = 1 inst → \* repo+env
- **If Ops handling the full Flux installation:**
  - Devs requiring new or updated Flux configurations
    - **Ops overload**
- **To avoid Ops overload**
  - Devs given full control of Flux installation
    - Devs with access to sec creds & config options they shouldn't have
      - **Not DevSecOps'ey !!**

# Flux separation of concerns by role

## Ops (aka “the Platform team”) : Setting up Flux defaults

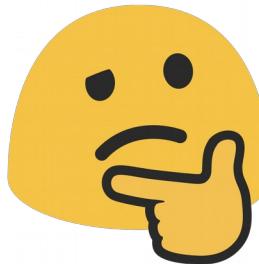
- Container image repo management
  - Allowed corporate/external repos
  - Access credentials
  - Image/Chart white/black listing & caching
  - ...
- Deployment target setup
  - Kube clusters + namespaces
  - Credentials
  - ...
- System Def Git repo management
  - Git repo whitelisting
  - Access credentials
  - ...

# Flux separation of concerns by role

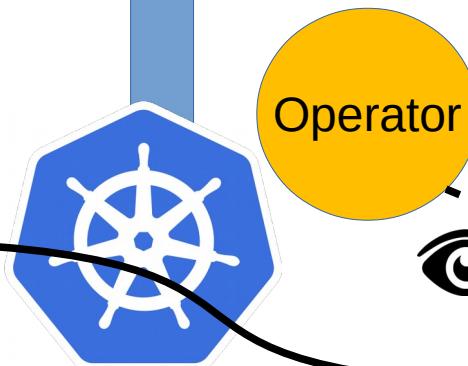
## Devs : Configuring specific Flux use cases

- **Specifying what to deploy**
  - i.e. The source Git repo K8s configs
- **Specifying where to deploy**
  - i.e. Identifying the target K8s cluster by a logical ID
    - And the target namespace,  
(if using 1 cluster → \* environments)

# Wouldn't it be great if...



Dev Kolontái



Operator

NS Kolontái

FluxConfig

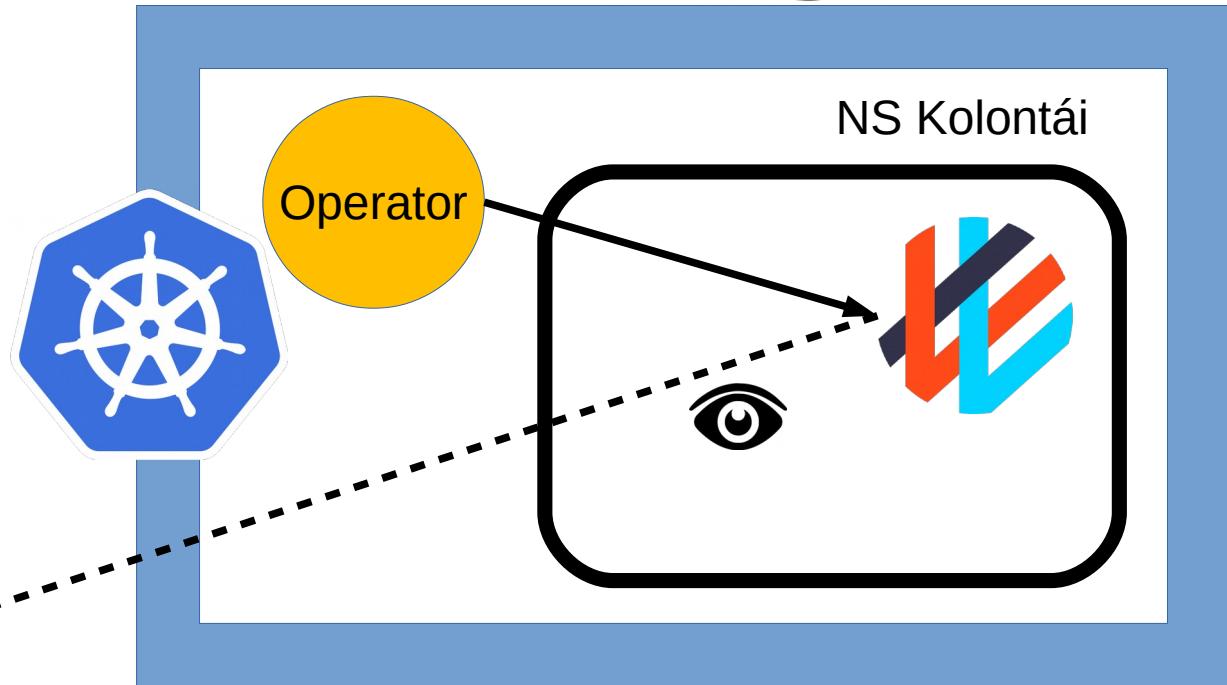
references



# Wouldn't it be great if...



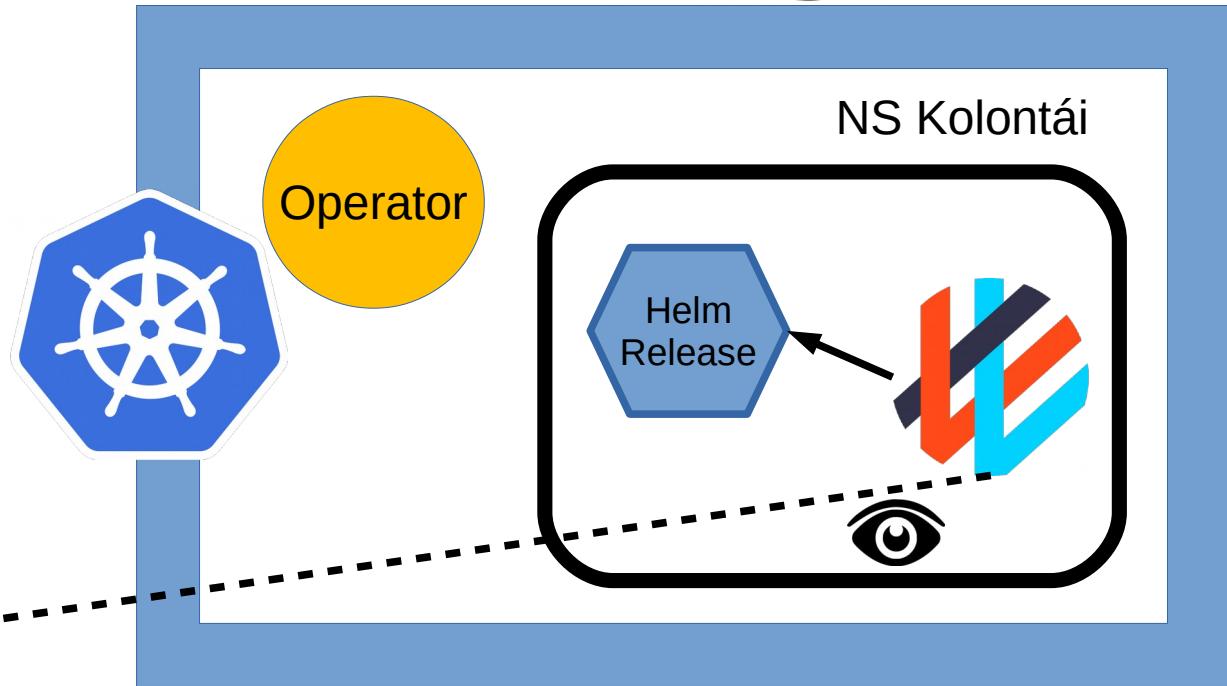
Dev Kolontái



# Wouldn't it be great if...



Dev Kolontái



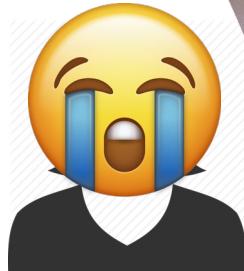
# Quick quiz:

Looks DevOps'ey,  
doesn't it?



# Wouldn't it be great if.

Dev Kolontái



YOU KNOW NOTHING



JON SNOW



NS Kolontái



quickmeme.com



# Wouldn't it be great if...

Dev Kolontái



Not committed  
anywhere

→ SNOWFLAKE!

NS Kolontái

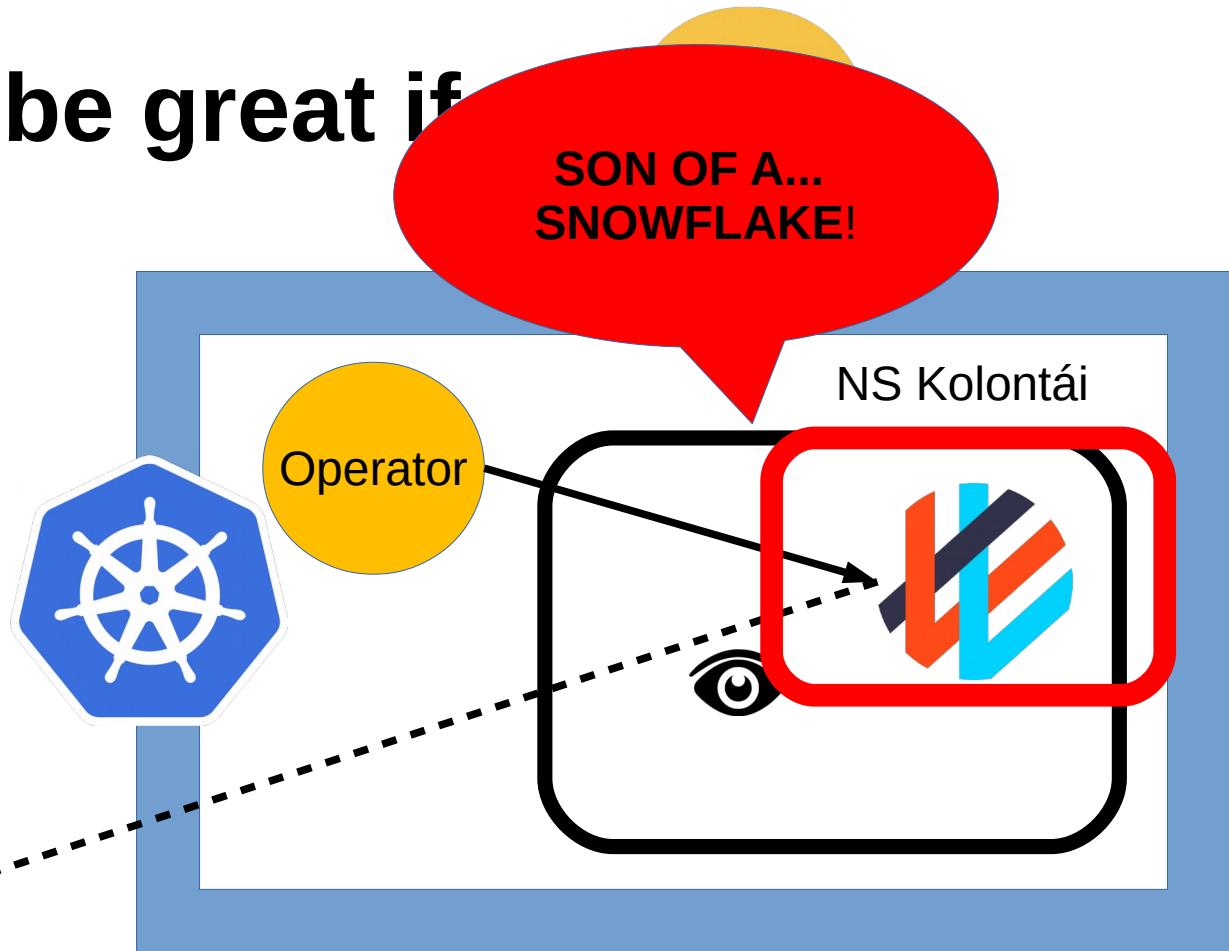


references

FluxConfig

# Wouldn't it be great if

Dev Kolontái



# Wouldn't it be great

DevOps = **IaC** + CD + SRE

**IaC** = Reproducible systems

Manually created stuff

**==**

Irreproducible stuff

SON OF A  
SON OF A  
SNOWFLAKE!

Operator

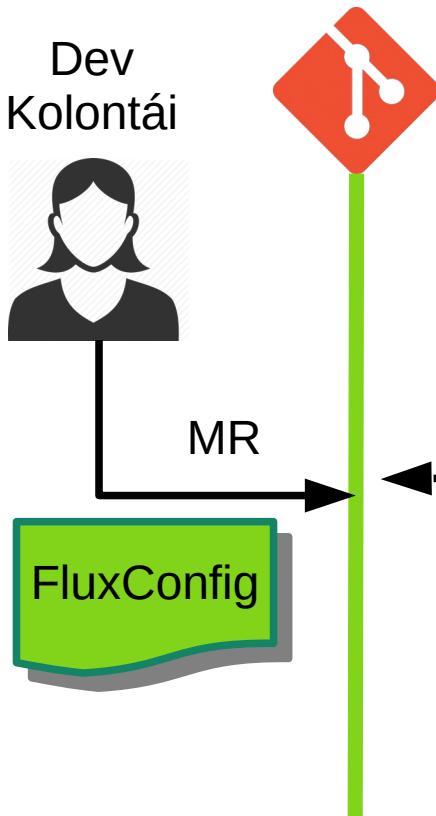
NS Kolontái

Helm  
Release

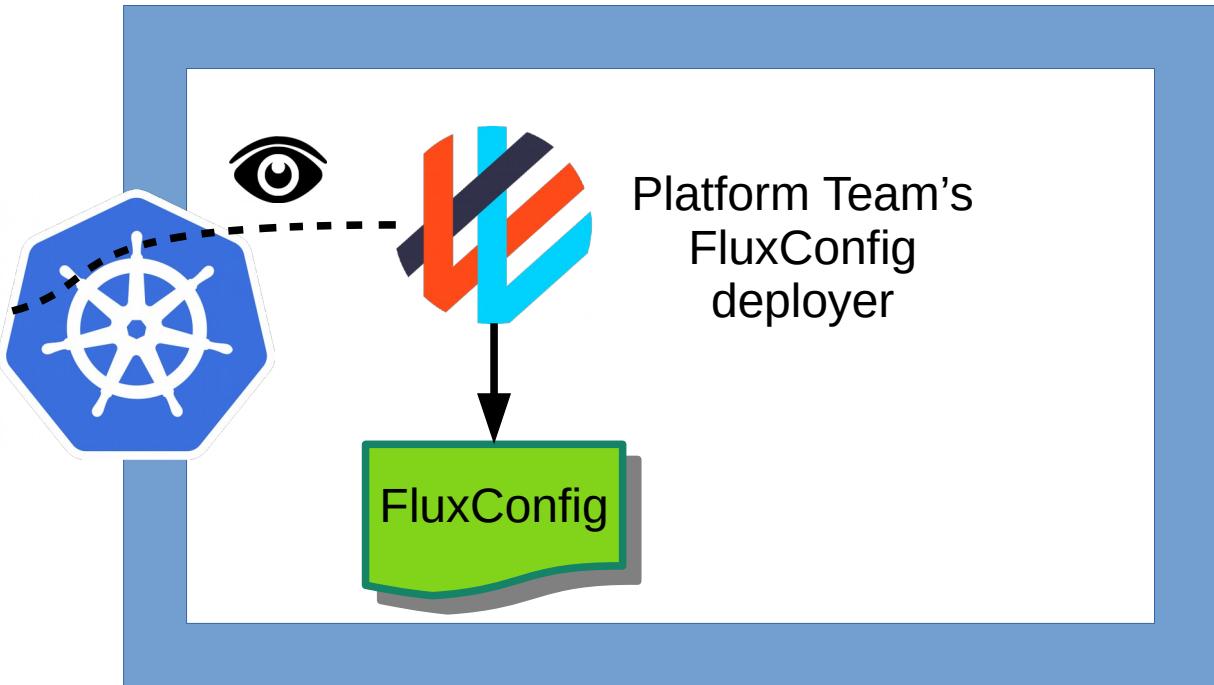


Platform Team's  
FluxConfig repo

Dev  
Kolontái



# However, this way...



master



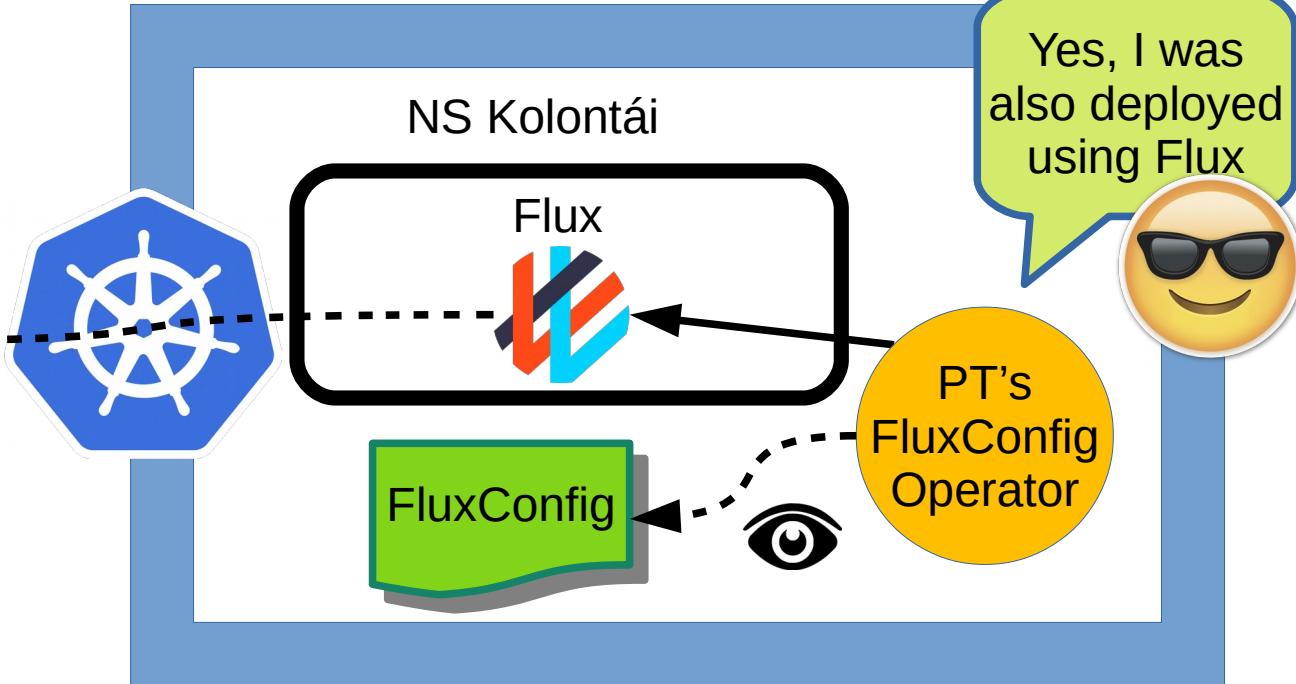
# However, this way...

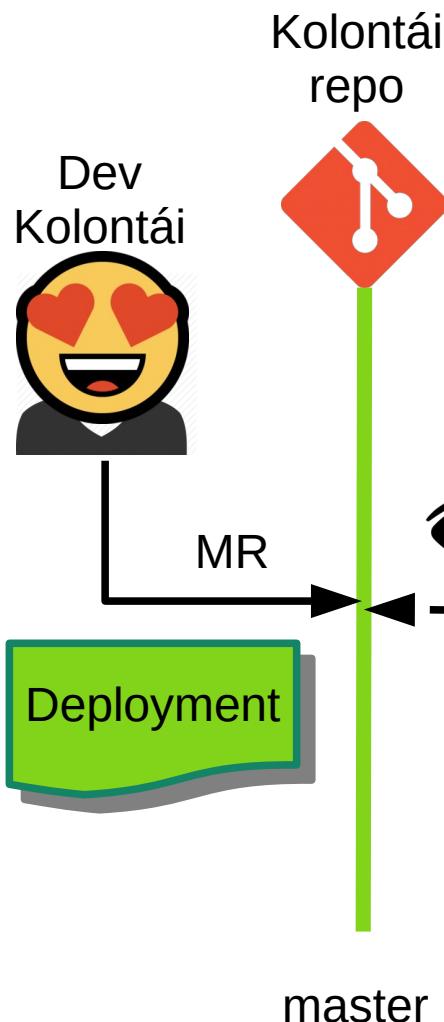
Kolontái  
repo

Dev  
Kolontái



master





# ...awesomeness ensues!



# Enter the Flux Operator!

<https://github.com/justinbarrick/flux-operator>

- **Includes a Flux CRD**
  - Teams wishing for a Flux instance deploy a resource of this CRD type
- **Two modes of operation:**
  - **Cluster mode:** Flux resources indicate in which namespace the Flux instance must be deployed
    - **When Flux definitions come from admins → Any NS**
  - **Namespace mode:** Flux instance created in the namespace where we deploy the Flux CRD
    - **When Flux definitions come from dev teams → Single NS**

# DEMO #6

Flux Operator FTW!

# Flux Operator

<https://github.com/justinbarrick/flux-operator>

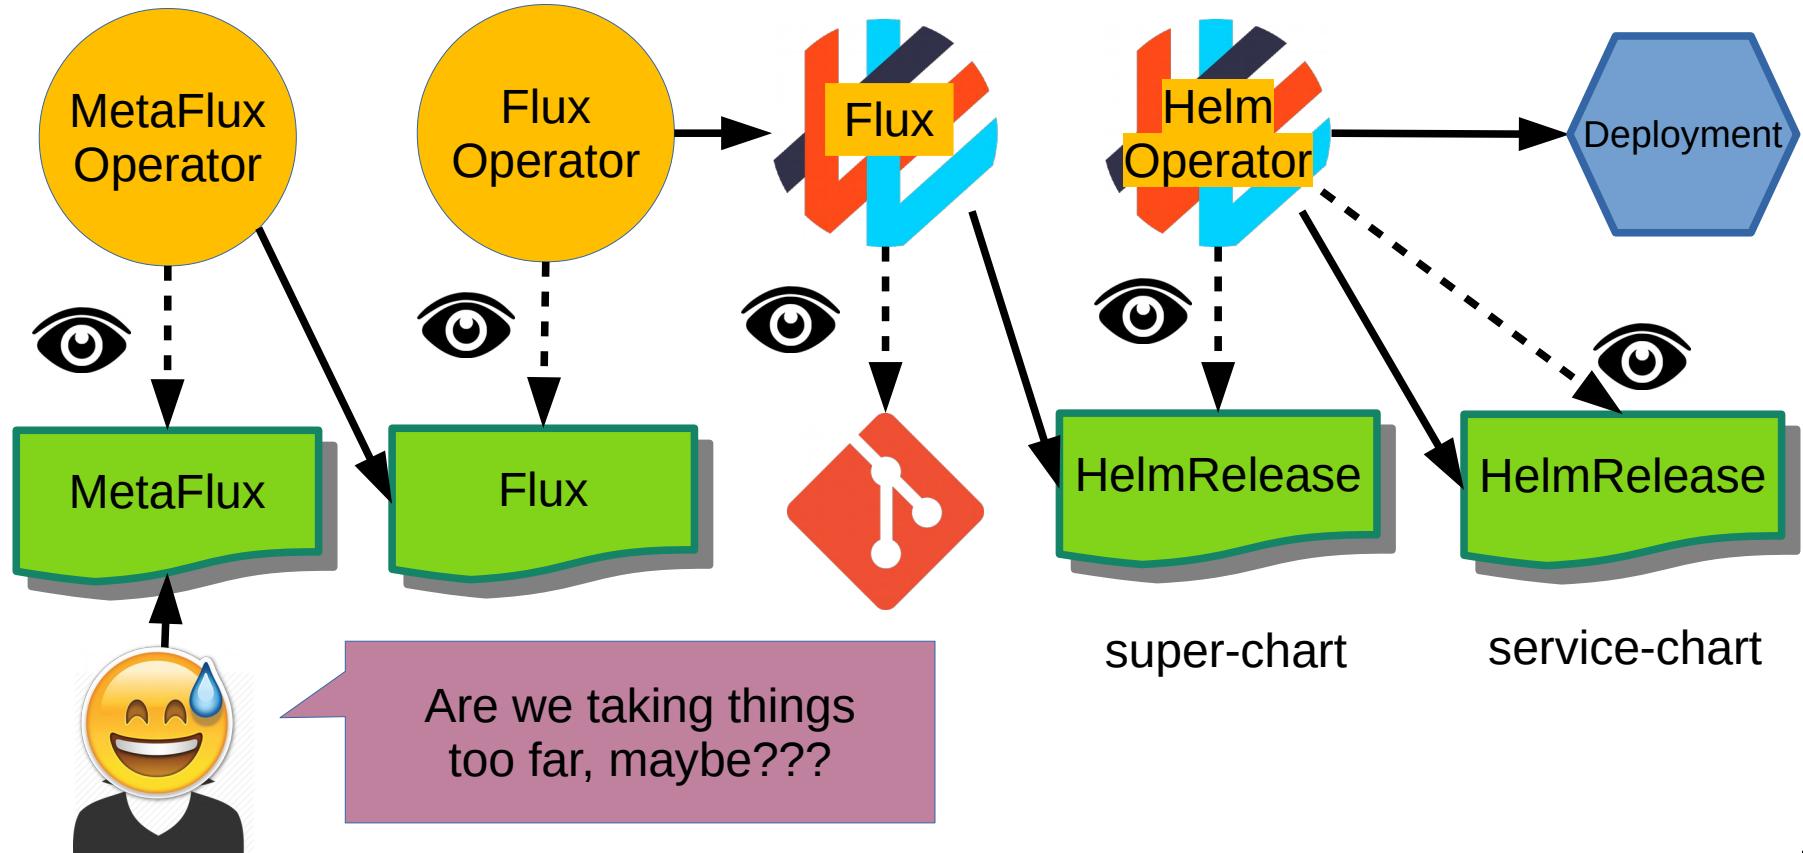
- **Drawbacks:**
  - **Very young project! (v0.0.5)**
  - **Doesn't really implement the ideal “separation of concerns” mentioned before :-(**
    - The definer of the Flux CRD resource provides **the full configuration** for the Flux instance
      - Can't define defaults for the Flux instance config in the Operator
        - e.g. Can't define default credentials for some repos centrally...
      - We can't validate the requested Flux instance config
        - e.g. Can't do white/black listing of Git repos...
    - **Therefore, we need to keep on searching...**

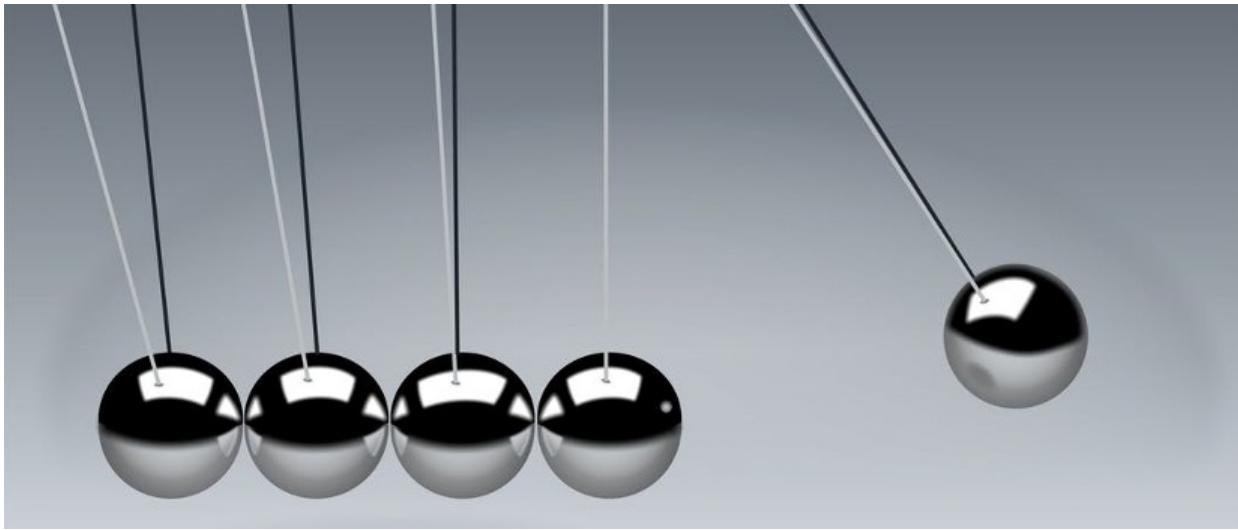


Or, you know,  
start contrib-  
uting...



# Enter our own Meta-Flux-Operator???





Platform Team's  
FluxConfig repo

Dev  
Kolontái



MR



Minimal  
FluxConfig

## Solution #1: Automation fills in the blanks

- What git repo
- What k8s cluster
- What k8s namespace
- Docker image whitelisting

Platform Team's  
FluxConfig deployer



Update  
& merge

FluxConfig

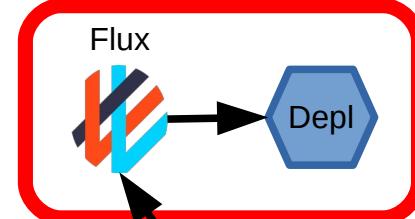
Completion  
& validation

master



FluxConfig

NS Kolontái



Flux  
Operator

Platform Team's  
FluxConfig repo

Dev  
Kolontái



MR



Minimal  
FluxConfig

## Solution #2: Automation deploys Flux (simpler)

- What git repo
- What k8s cluster
- What k8s namespace
- Docker image whitelisting



merge

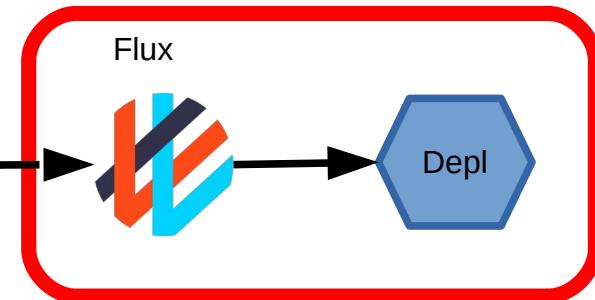
deploy Flux  
instance

Completion  
& validation

master



NS Kolontái



# But...!!!



# Can we do better?

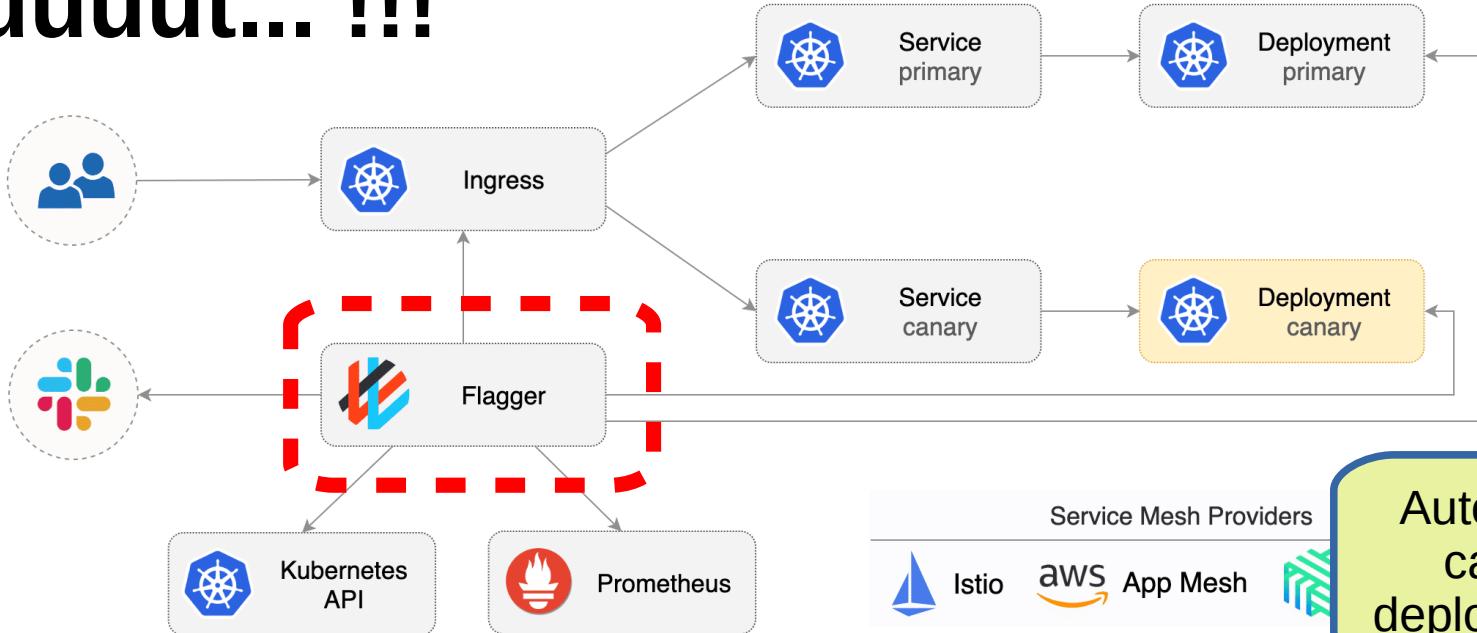




... not today.



# Buuuut... !!!

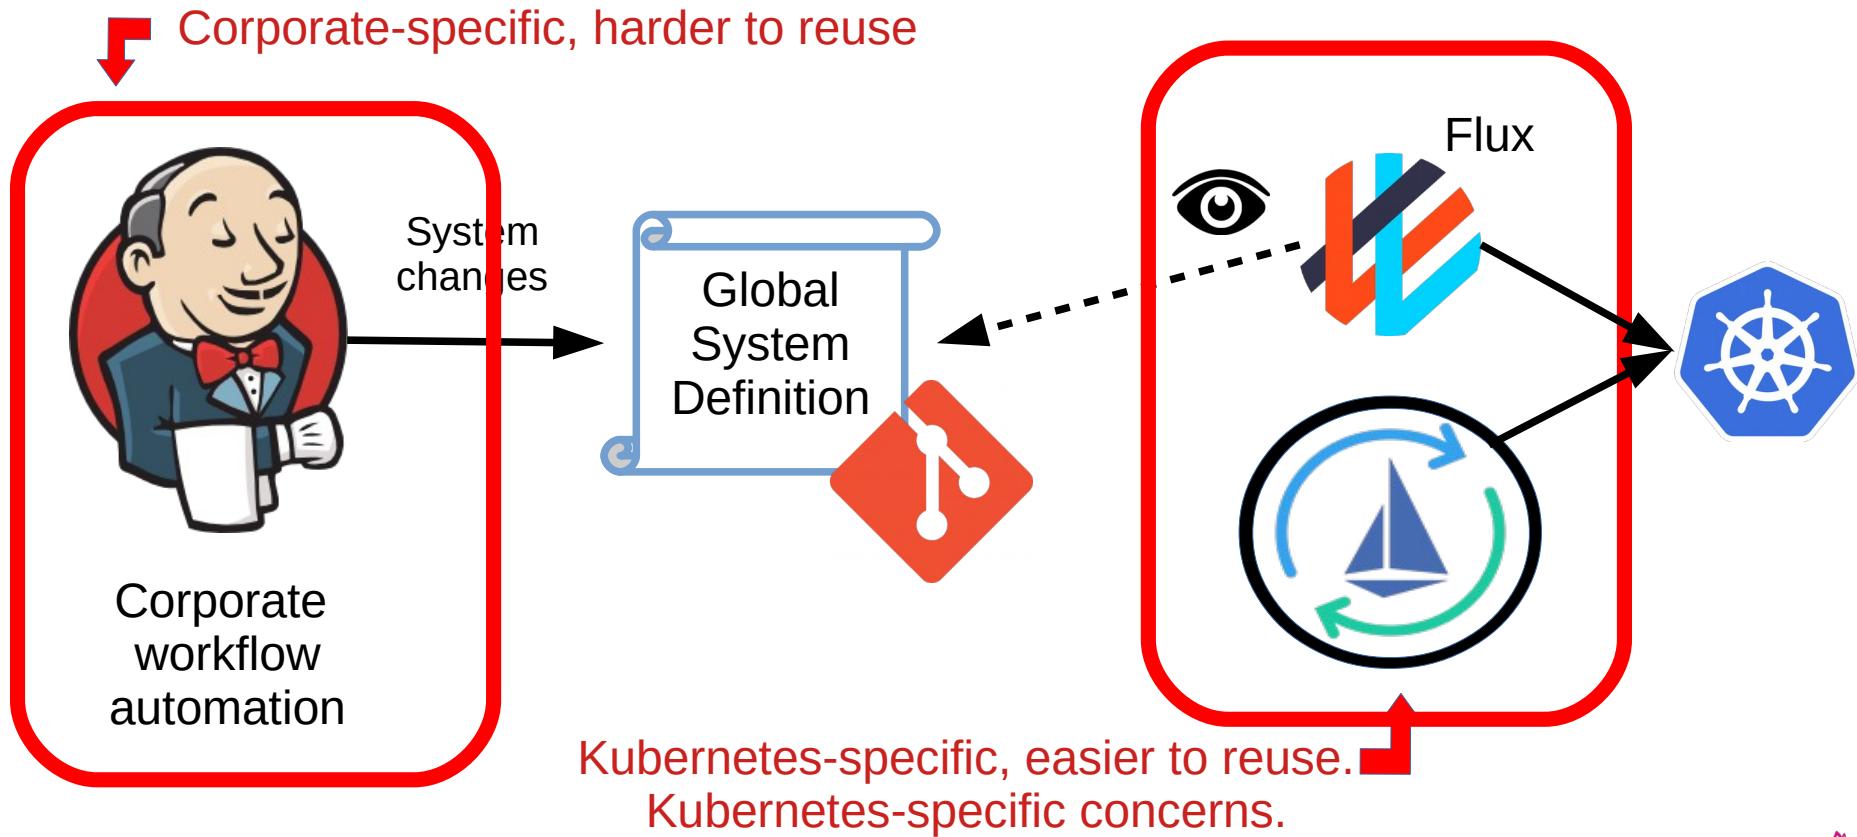


Automated  
canary  
deployments!

Let's talk some other day  
about Flagger...



# The Flagger proposal







THANK YOU